

Lecture 16:

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

The Memory System 2/3

Introduction to Computer Architecture
UC Davis EEC 170, Fall 2019

Locality Properties of Patterns

- **Two locality properties of memory references:**
 - ***Temporal Locality:*** If a location is referenced it is likely to be referenced again in the near future.
Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs
 - ***Spatial Locality:*** If a location is referenced it is likely that locations near it will be referenced in the near future.

Caches

- Caches exploit both properties of patterns.
 - Exploit *temporal* locality by remembering the contents of recently accessed locations.
Assignment Project Exam Help
 - Exploit *spatial* locality by remembering blocks of contents of recently accessed locations.
<https://tutorcs.com>
WeChat: cstutorcs

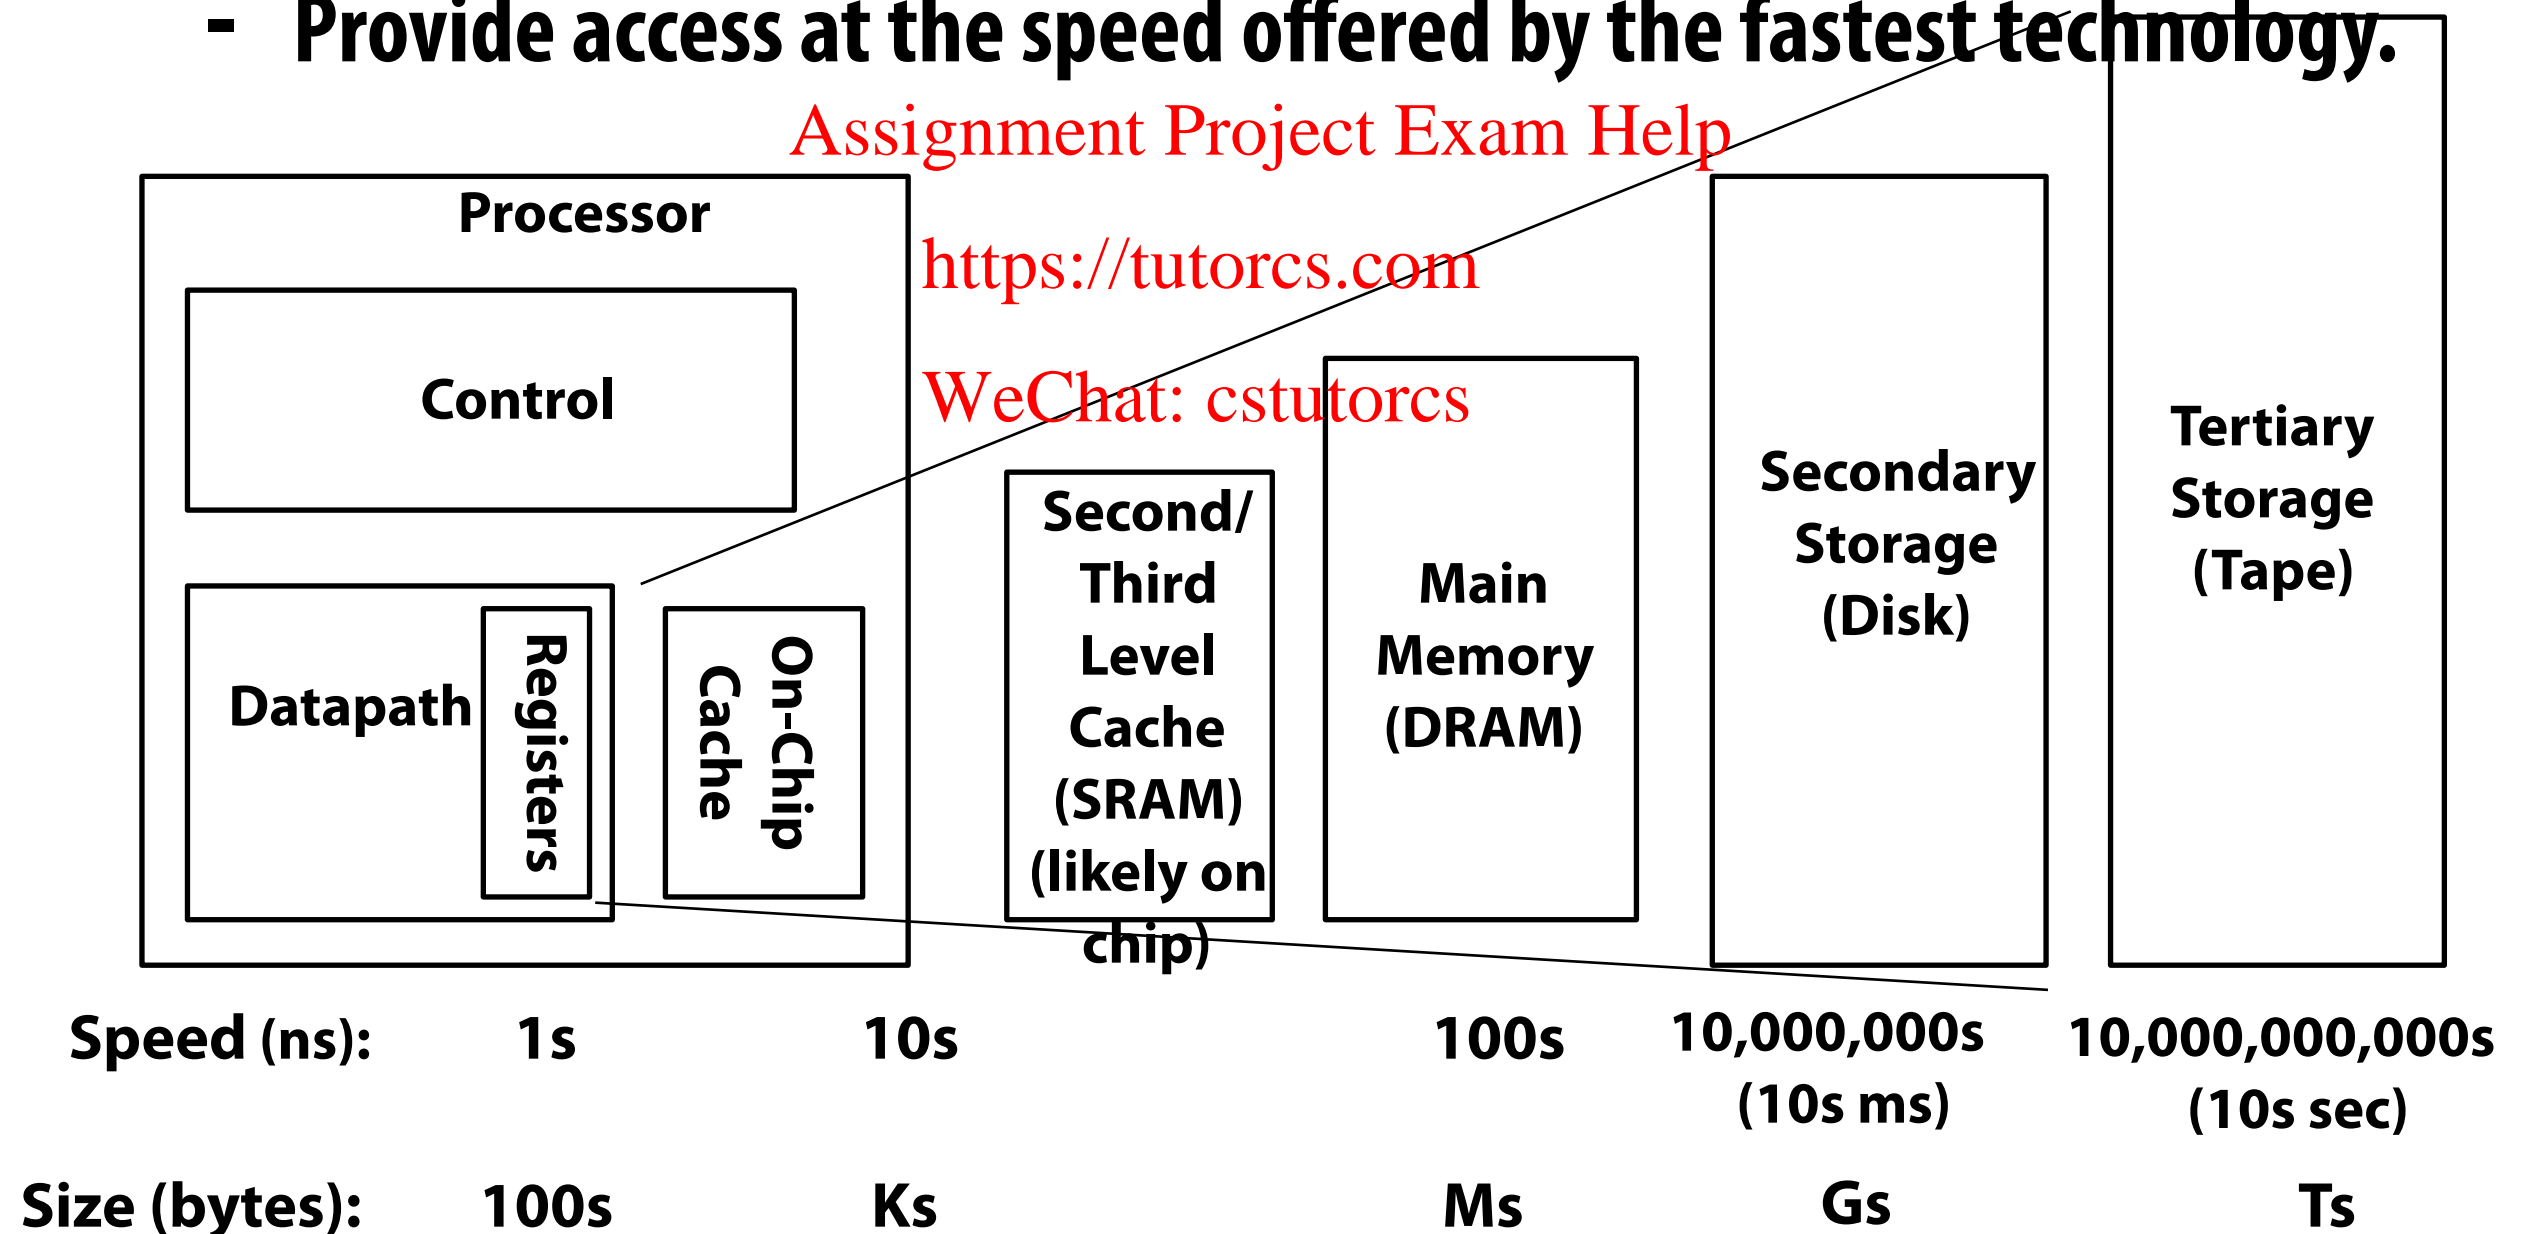
Memory Hierarchy of a Modern Computer System

- By taking advantage of the principle of *locality*:
 - Present the user with as much memory as is available in the cheapest technology.
 - Provide access at the speed offered by the fastest technology.

Assignment Project Exam Help


<https://tutorcs.com>

WeChat: cstutorcs




Direct Mapped Cache

- Location determined by address
- Direct mapped: only one choice
 - (Block address) modulo (#Blocks in cache)



- *#Blocks is a power of 2*
- *Use low-order address bits*

Address Subdivision




Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Example: Larger Block Size

- **64 blocks, 16 bytes/block**
 - **16 bytes/block means 4 bits of offset ($2^4 = 16$)**
 - **To what block number does address 1200 map?**
- **Block address = $\lfloor 1200/16 \rfloor = 75$**
 - **I am the 75th 16-byte block in the address space**
- **Block number = $75 \bmod 64 = 11$**



Block Size Considerations

- Larger blocks should reduce miss rate
 - Due to spatial locality
- But in a fixed-sized cache
 - Larger blocks \Rightarrow fewer of them
 - More competition \Rightarrow increased miss rate
 - Larger blocks \Rightarrow pollution
- Larger miss penalty
 - Can override benefit of reduced miss rate
 - Early restart and critical-word-first can help

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Cache Misses

- On cache hit, CPU proceeds normally
 - This is (hopefully) the common case
 - This is one or a few cycles at most
- On cache miss
 - Stall the CPU pipeline
 - Fetch block from next level of hierarchy
 - Instruction cache miss
 - Restart instruction fetch
 - Data cache miss
 - Complete data access

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Example: Intrinsicity FastMATH

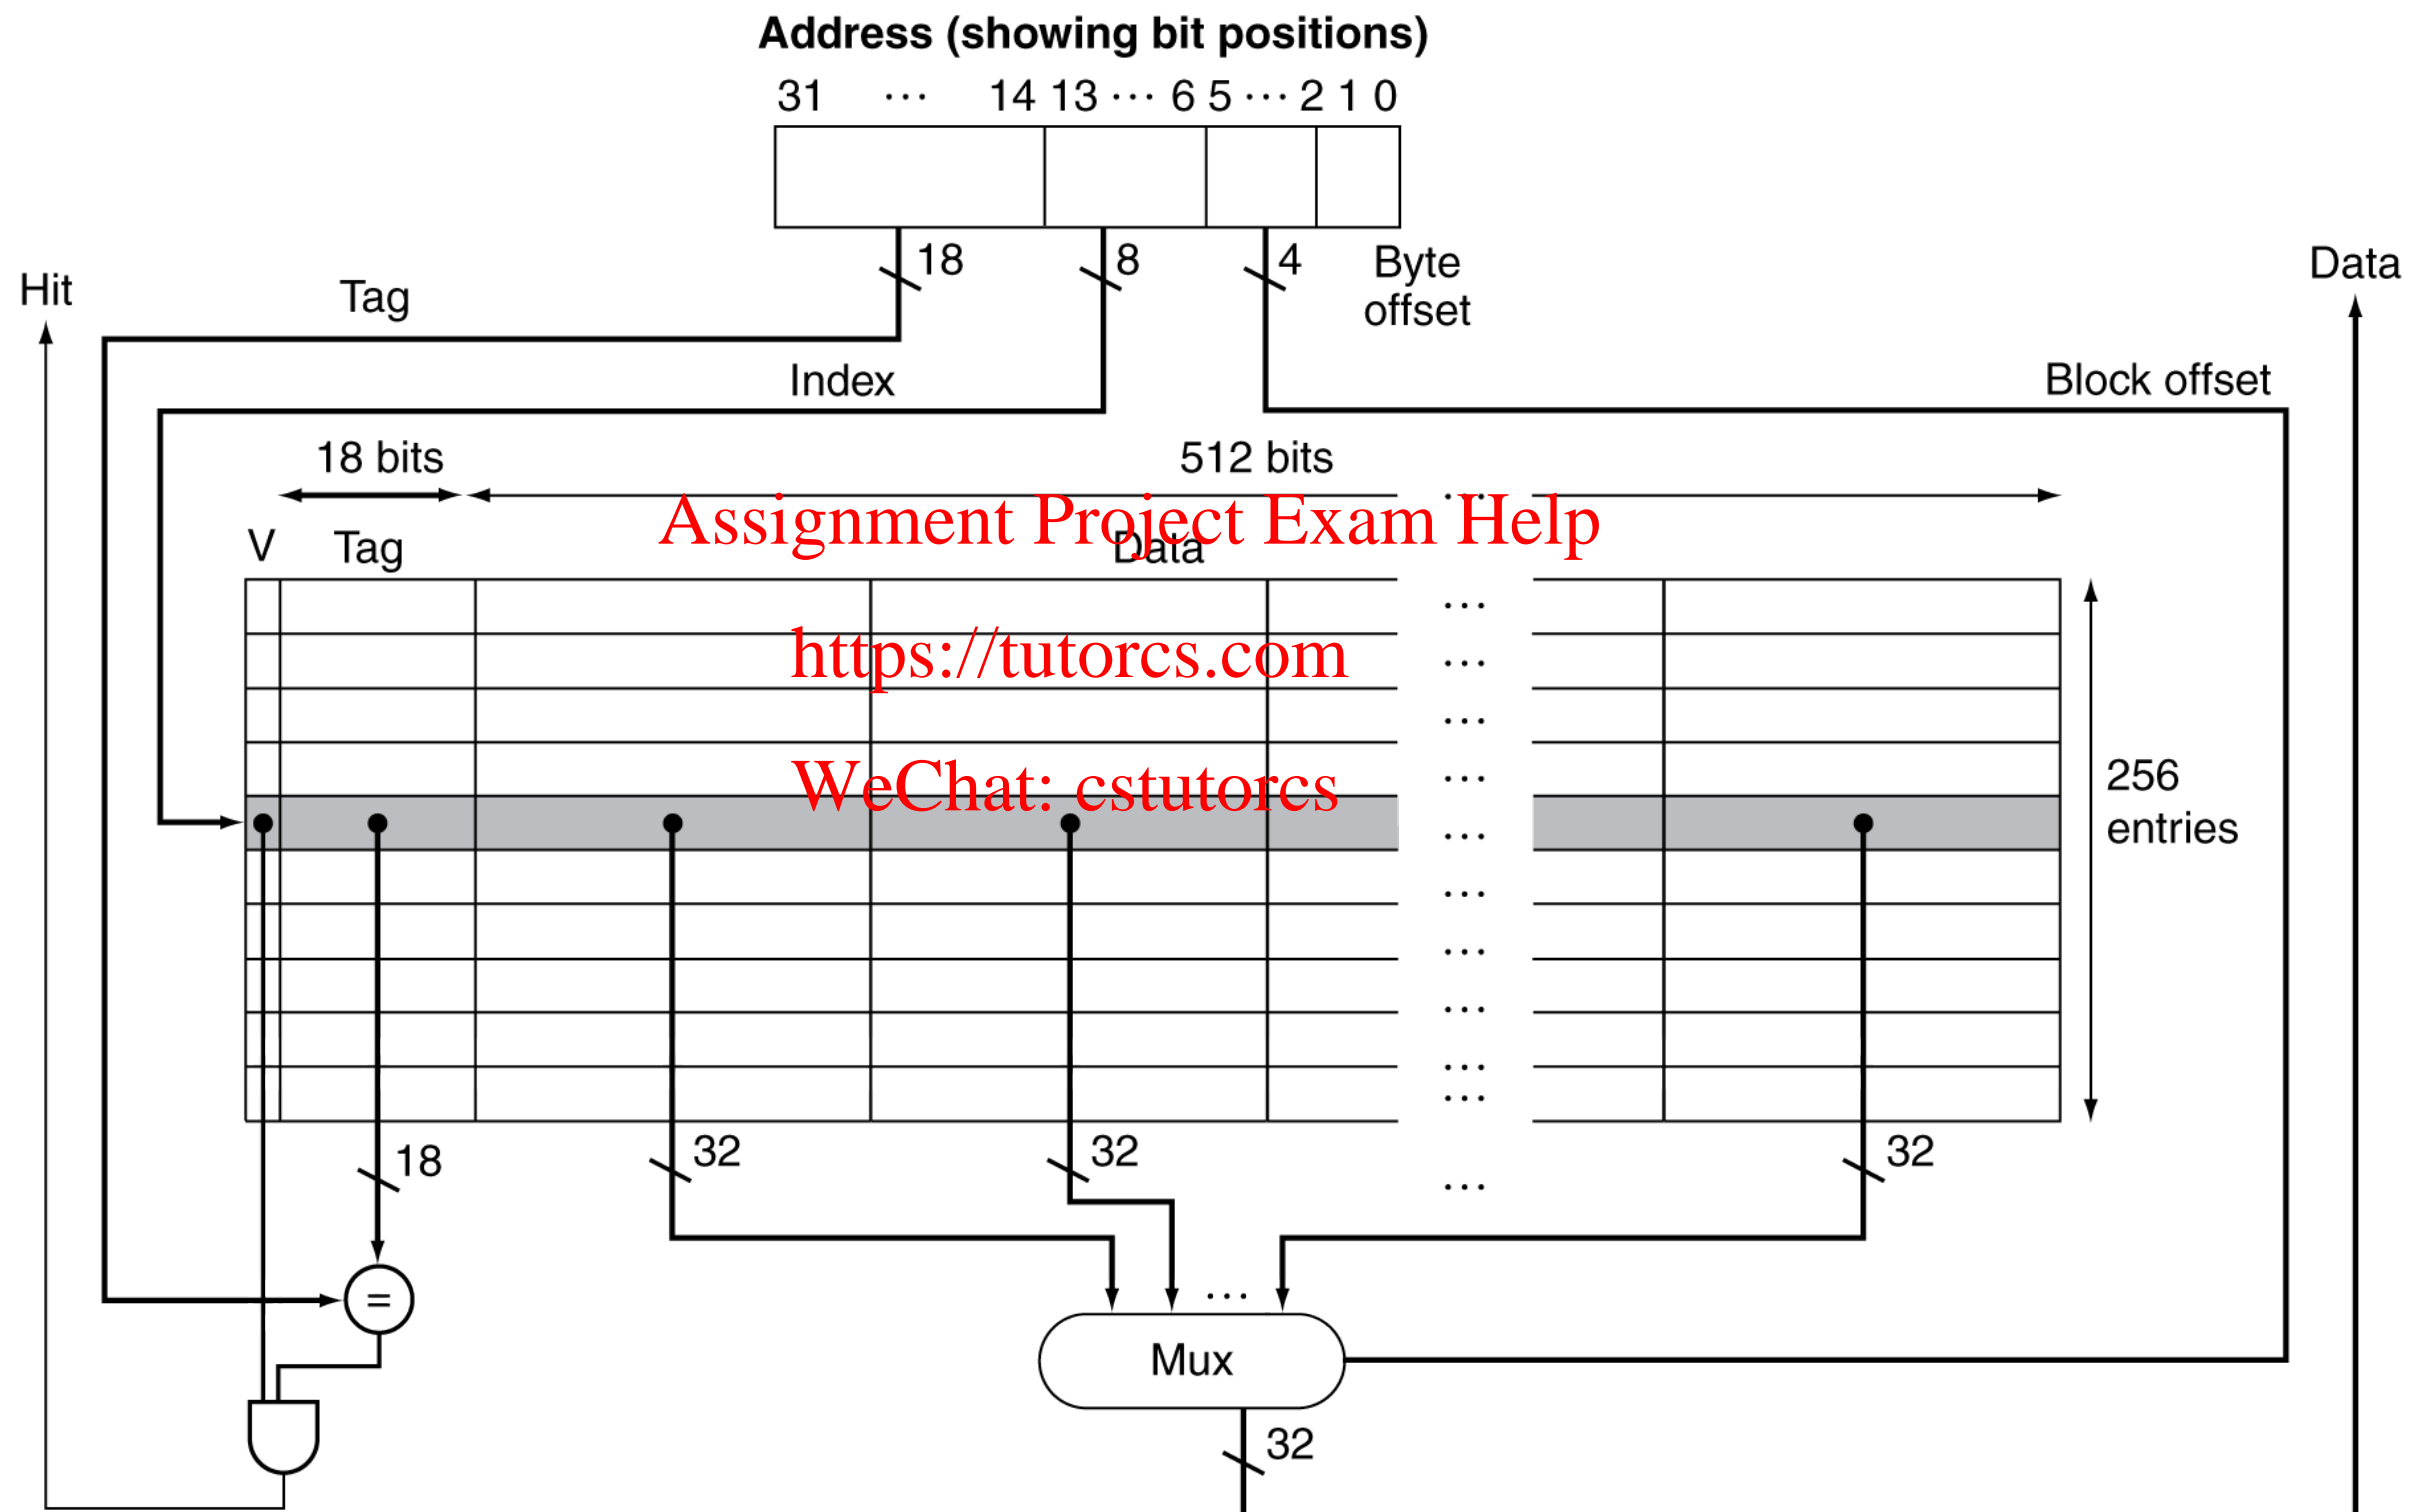
- Embedded MIPS processor
 - 12-stage pipeline
 - Instruction and data access on each cycle
- Split cache: separate I-cache and D-cache
 - Each 16KB: 256 blocks \times 16 words/block
 - D-cache: write-through or write-back (we will talk about this later)
- SPEC2000 miss rates
 - I-cache: 0.4%
 - D-cache: 11.4%
 - Weighted average: 3.2%

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Example: Intrinsity FastMATH



Associative Caches

■ Fully associative

- Allow a given block to go in any cache entry
- Requires all entries to be searched at once
- Comparator per entry (expensive)

Assignment Project Exam Help

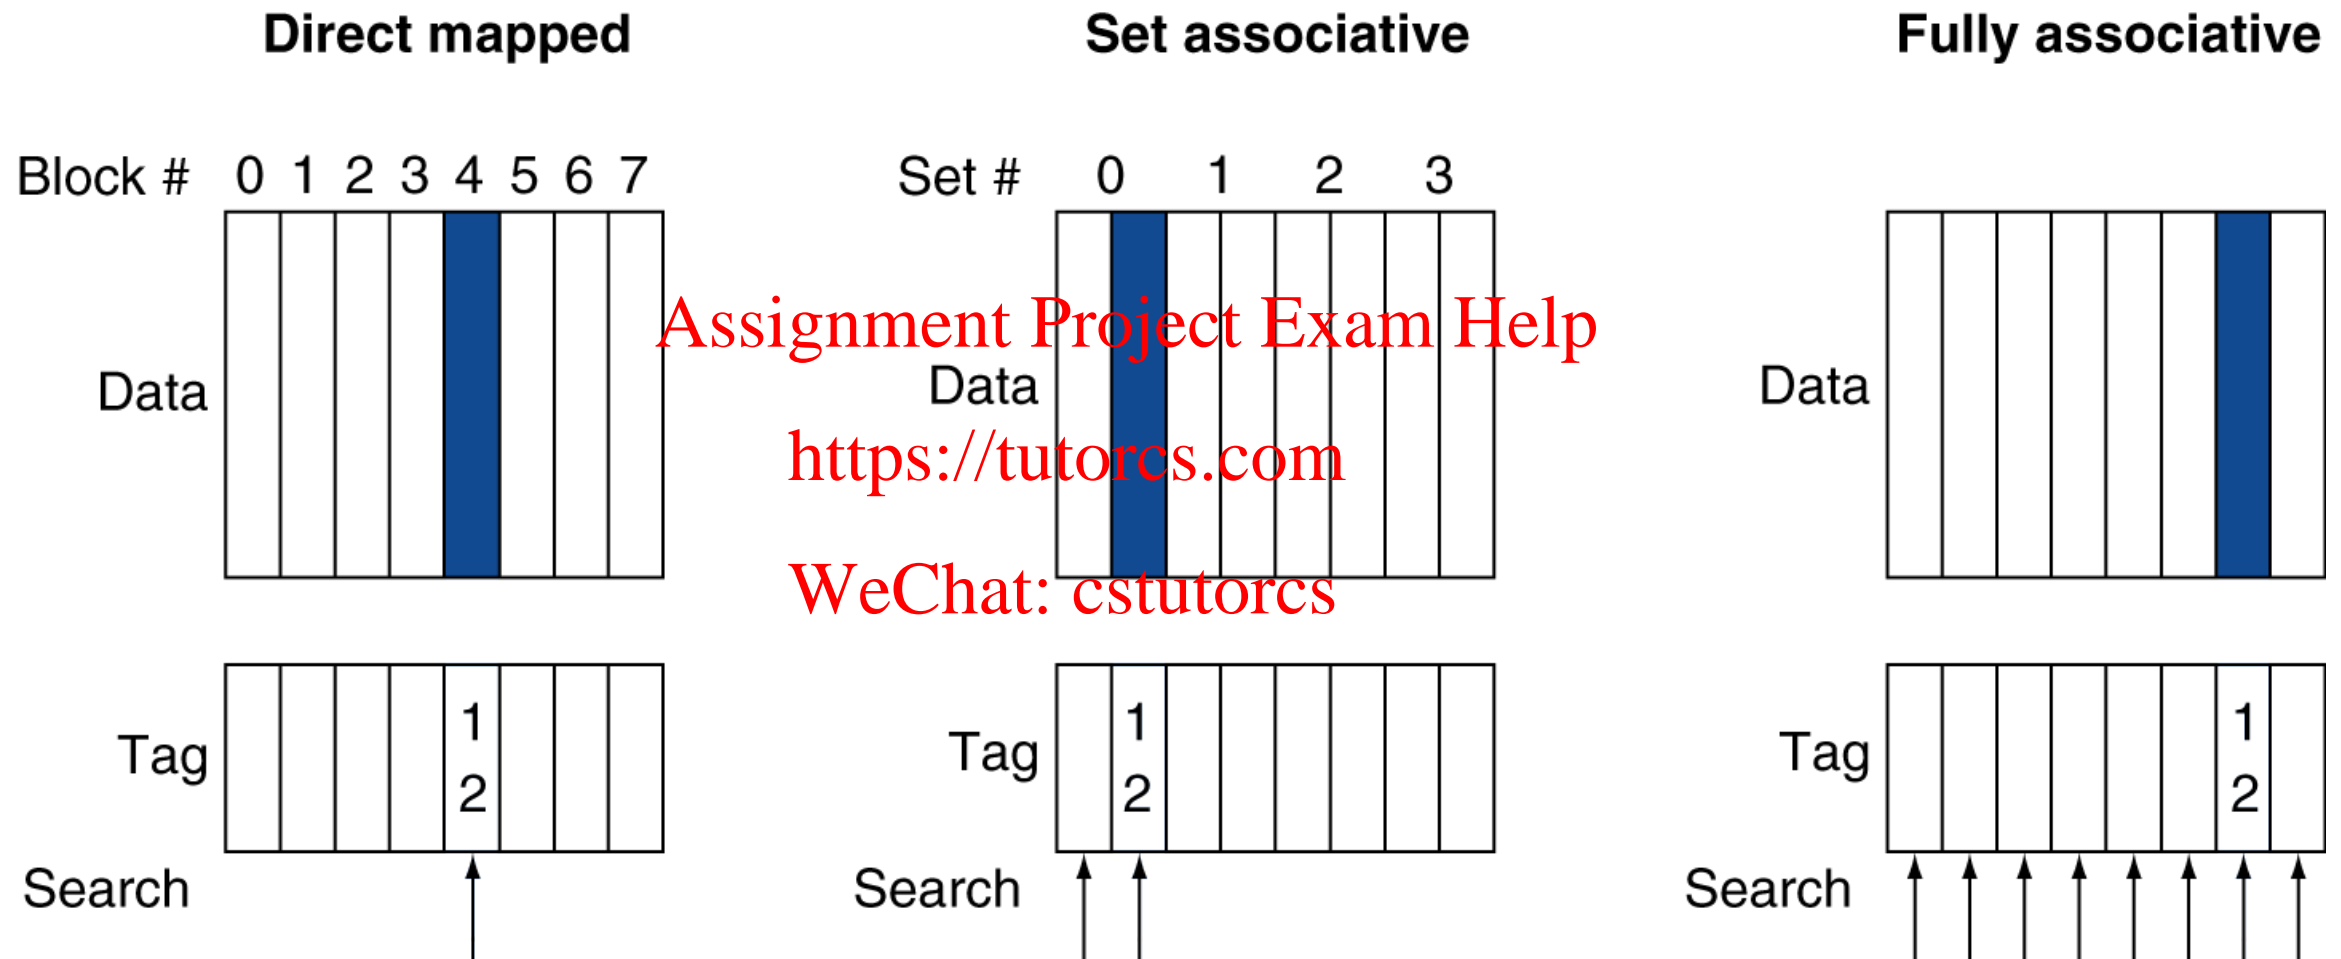
■ n -way set associative

<https://tutorcs.com>

- Each set contains n entries
 - 1-way set associative == direct mapped
- Block number determines which set
 - (Block number) modulo (#Sets in cache)
- Search all entries in a given set at once
- n comparators (less expensive)

WeChat: cstutorcs

Associative Cache Example



Spectrum of Associativity

■ For a cache with 8 entries

One-way set associative (direct mapped)

| Block | Tag | Data |
|-------|-----|------|
| 0 | | |
| 1 | | |
| 2 | | |
| 3 | | |
| 4 | | |
| 5 | | |
| 6 | | |
| 7 | | |

Two-way set associative

| Set | Tag | Data | Tag | Data |
|-----|-----|------|-----|------|
| 0 | | | | |
| 1 | | | | |
| 2 | | | | |
| 3 | | | | |

Four-way set associative

| Set | Tag | Data | Tag | Data | Tag | Data | Tag | Data |
|-----|-----|------|-----|------|-----|------|-----|------|
| 0 | | | | | | | | |
| 1 | | | | | | | | |

Eight-way set associative (fully associative)

| Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data |
|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|
| | | | | | | | | | | | | | | | |

Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, fully associative
 - Block access sequence: 0, 8, 0, 6, 8

Assignment Project Exam Help

- Direct mapped

<https://tutorcs.com>

WeChat: cstutorcs

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|---|--------|---|
| | | | 0 | 1 | 2 | 3 |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[8] | | | |
| 0 | 0 | miss | Mem[0] | | | |
| 6 | 2 | miss | Mem[0] | | Mem[6] | |
| 8 | 0 | miss | Mem[8] | | Mem[6] | |

Associativity Example

- Block access sequence: 0, 8, 0, 6, 8
- 2-way set associative

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|--------|-------|--|
| | | | Set 0 | | Set 1 | |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[0] | Mem[8] | | |
| 0 | 0 | hit | Mem[0] | Mem[8] | | |
| 6 | 0 | miss | Mem[0] | Mem[6] | | |
| 8 | 0 | miss | Mem[8] | Mem[6] | | |

- Fully associative

| Block address | | Hit/miss | Cache content after access | | | |
|---------------|--|----------|----------------------------|--------|--------|--|
| 0 | | miss | Mem[0] | | | |
| 8 | | miss | Mem[0] | Mem[8] | | |
| 0 | | hit | Mem[0] | Mem[8] | | |
| 6 | | miss | Mem[0] | Mem[8] | Mem[6] | |
| 8 | | hit | Mem[0] | Mem[8] | Mem[6] | |

How Much Associativity


- Increased associativity decreases miss rate
 - But with diminishing returns
- Simulation of a system with 64 KB D-cache, 16-word blocks, SPEC2000
 - 1-way: 10.3%
 - 2-way: 8.6%
 - 4-way: 8.3%
 - 8-way: 8.1%

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Set Associative Cache Organization



Replacement Policy

- **Direct mapped: no choice**
- **Set associative**
 - **Prefer non-valid entry, if there is one**
 - **Otherwise, choose among entries in the set**
- **Least-recently used (LRU)**
 - **Choose the one unused for the longest time**
 - **Simple for 2-way, manageable for 4-way, too hard beyond that**
- **Random**
 - **Gives approximately the same performance as LRU for high associativity**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Measuring Cache Performance

■ Components of CPU time

- Program execution cycles
 - Includes cache hit time

- Memory stall cycles

- Mainly from cache misses

Assignment Project Exam Help

<https://tutorcs.com>

■ With simplifying assumptions:

Memory stall cycles

$$= \frac{\text{Memory accesses}}{\text{Program}} \times \text{Miss rate} \times \text{Miss penalty}$$

$$= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Misses}}{\text{Instruction}} \times \text{Miss penalty}$$

Cache Performance Example

■ Given

- I-cache miss rate = 2%
 - D-cache miss rate = 4%
 - Miss penalty = 100 cycles
 - Base CPI (ideal cache) = 2
 - Load & stores are 36% of instructions
- Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

■ Miss cycles per instruction

- I-cache: $0.02 \times 100 = 2$
- D-cache: $0.36 \times 0.04 \times 100 = 1.44$

■ Actual CPI = $2 + 2 + 1.44 = 5.44$

- Ideal CPU is $5.44/2 = 2.72$ times faster

Average Access Time

- Hit time is also important for performance
- Average memory access time (AMAT)
 - $AMAT = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$
- Example
 - CPU with 1 ns clock, hit time = 1 cycle, miss penalty = 20 cycles, l-cache miss rate = 5%
 - $AMAT = 1 + 0.05 \times 20 = 2\text{ns}$
 - 2 cycles per instruction

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Performance Summary

- When CPU performance increased
 - Miss penalty becomes more significant
- Decreasing base CPI
 - Greater proportion of time spent on memory stalls
- Increasing clock rate <https://tutorcs.com>
 - Memory stalls account for more CPU cycles
- Can't neglect cache behavior when evaluating system performance
 - Consequence: Increased attention to cache; more levels of cache

Multilevel Caches

- **Primary cache attached to CPU**
 - **Small, but fast**
- **Level-2 cache services misses from primary cache**
 - **Larger, slower, but still faster than main memory**
- **Main memory services L2 cache misses**
- **Some high-end systems include L3 cache**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: estuorcs

Multilevel Cache Example

■ Given

- CPU base CPI = 1, clock rate = 4GHz
- Miss rate/instruction = 2%
- Main memory access time = 100 ns

■ With just primary cache

- Miss penalty = $100 \text{ ns} / 0.25 \text{ ns} = 400$ cycles
- Effective CPI = $1 + 0.02 \times 400 = 9$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Example (cont.)

- Now add L2 cache
 - Access time = 5 ns
 - Global miss rate to main memory = 0.5%
 - *Be careful on miss rates: miss rate on L2 alone, or global?*
- Primary miss with L2 hit
 - Penalty = $5 \text{ ns} / 0.25 \text{ ns} = 20 \text{ cycles}$
- Primary miss with L2 miss
 - Extra penalty = 400 cycles
- $\text{CPI} = 1 + 0.02 \times 20 + 0.005 \times 400 = 3.4$
- Performance ratio = $9 / 3.4 = 2.6$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Multilevel Cache Considerations

■ Primary cache

- Focus on minimal hit time

■ L2 cache

- Focus on low miss rate to avoid main memory access
- Hit time has less overall impact

■ Results

- L1 cache usually small enough to allow single-cycle access
- L1 block size smaller than L2 block size

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Main Memory Supporting Caches

- **Use DRAMs for main memory**
 - **Fixed width (e.g., 1 word)**
 - **Connected by fixed-width clocked bus**
 - **Bus clock is typically slower than CPU clock**
- **Example cache block read**
 - **1 bus cycle for address transfer**
 - **15 bus cycles per DRAM access**
 - **1 bus cycle per data transfer**
- **For 4-word block, 1-word-wide DRAM**
 - **Miss penalty = $1 + 4 \times 15 + 4 \times 1 = 65$ bus cycles**
 - **Bandwidth = $16 \text{ bytes} / 65 \text{ cycles} = 0.25 \text{ B/cycle}$**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Summary

■ Two Different Types of Locality:

- **Temporal Locality (Locality in Time):** If an item is referenced, it will tend to be referenced again soon.
- **Spatial Locality (Locality in Space):** If an item is referenced, items whose addresses are close by tend to be referenced soon.

Assignment Project Exam Help

■ By taking advantage of the principle of locality:

<https://tutorcs.com>

- Present the user with as much memory as is available in the cheapest technology.
- Provide access at the speed offered by the fastest technology.

WeChat: cstutorcs

■ DRAM is slow but cheap and dense:

- Good choice for presenting the user with a BIG memory system

■ SRAM is fast but expensive and not very dense:

- Good choice for providing the user FAST access time.

Interactions with Advanced CPUs

- **Out-of-order CPUs can execute instructions during cache miss**
 - **Pending store stays in load/store unit**
 - **Dependent instructions wait in reservation stations**
 - **Independent instructions continue**
- **Effect of miss depends on program data flow**
 - **Much harder to analyze**
 - **Use system simulation**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Interactions with Software




■ Misses depend on memory access patterns

- Algorithm behavior
- Compiler optimization for memory access

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs




Software Optimization via Blocking

- Goal: maximize accesses to data before it is replaced

- Consider inner loops of DGEMM:

```
for (int i = 0; i < m; i++) {  
    for (int j = 0; j < n; ++j) {  
        double cij = C[i+j*n];  
        for (int kk = 0; kk < k; kk++)  
            cij += A[i*k+kk] * B[j+kk*n];  
        C[j+i*n] = cij;  
    }  
}
```

*I did this indexing
by hand and it's
probably wrong*



DGEMM Access Pattern

■ C, A, and B arrays




older accesses

new accesses

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



note you can transpose! might be helpful

Cache Blocked DGEMM


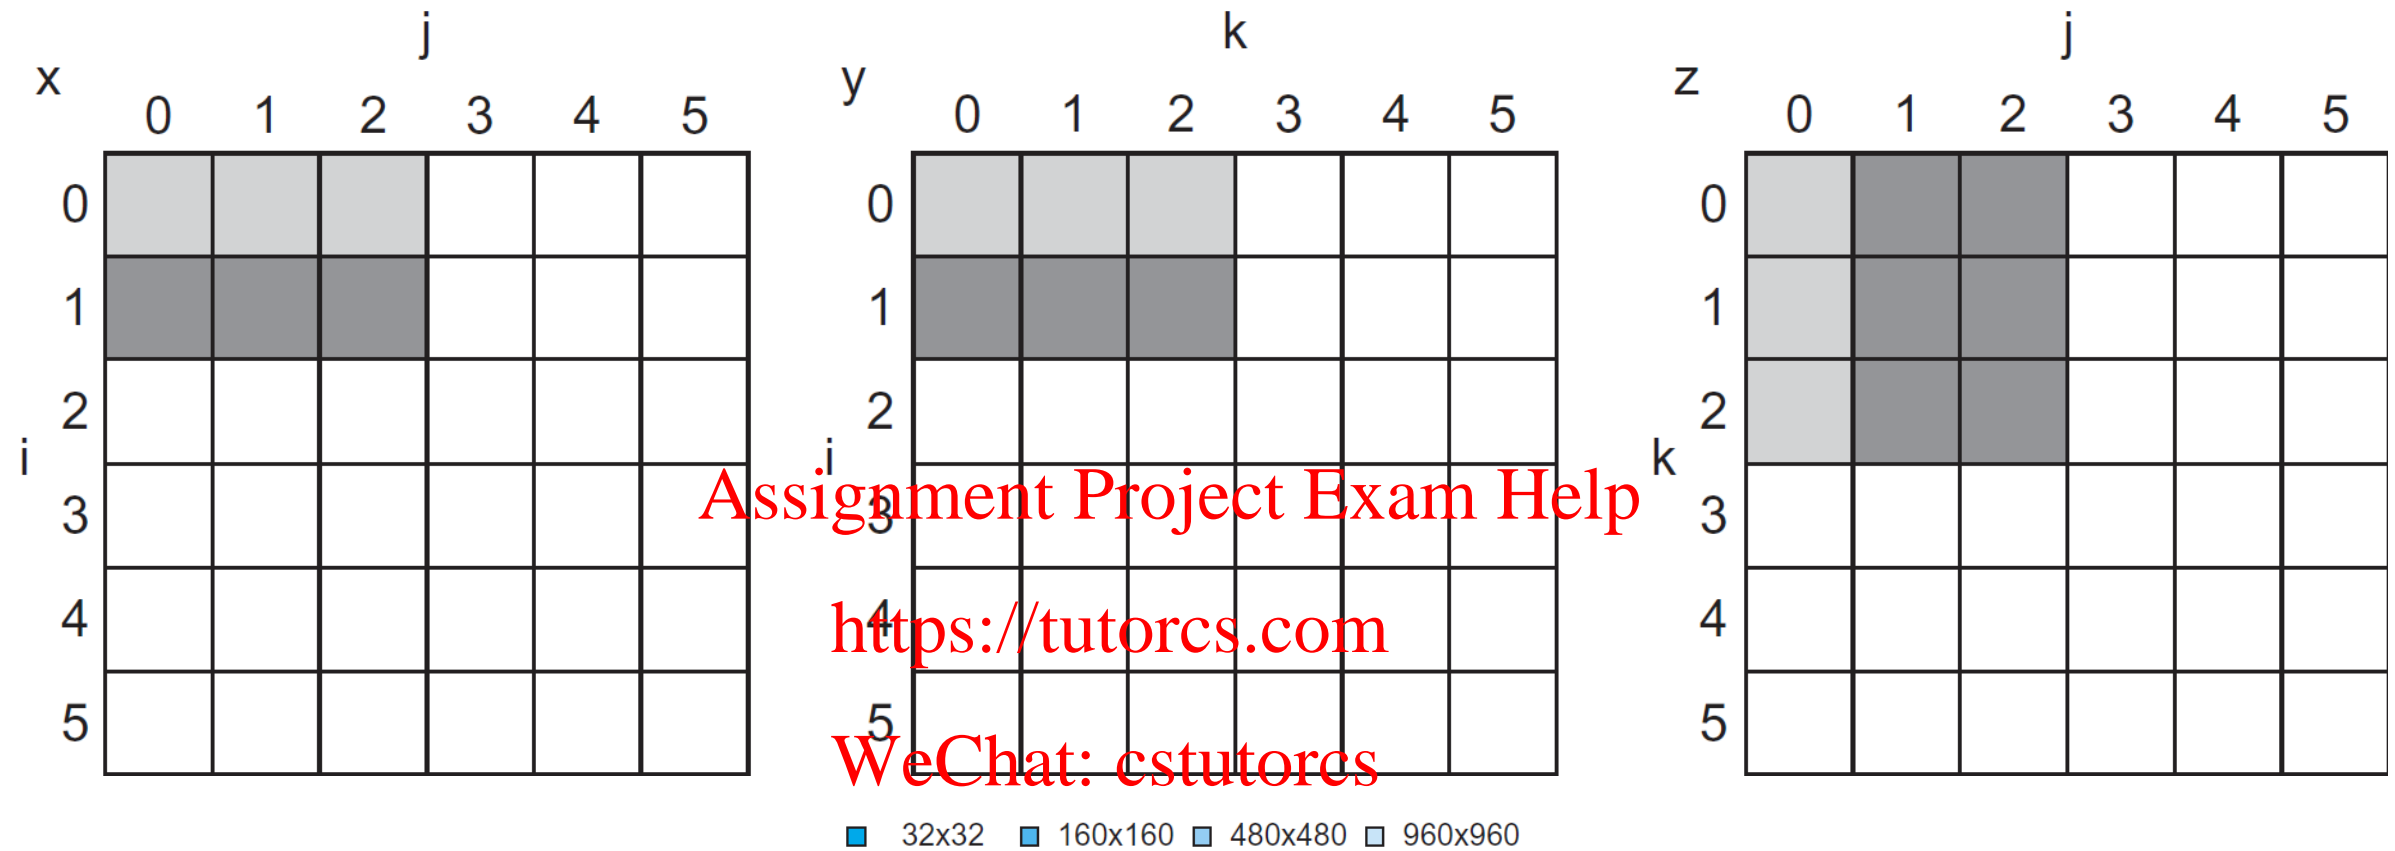
```
1 #define BLOCKSIZE 32
2 void do_block (int n, int si, int sj, int sk, double *A, double
3 *B, double *C)
4 {
5     for (int i = si; i < si+BLOCKSIZE; ++i)
6         for (int j = sj; j < sj+BLOCKSIZE; ++j)
7             {
8                 double cij = C[i+j*n]; /* cij = C[i][j] */
9                 for( int k = sk; k < sk+BLOCKSIZE; k++ )
10                     cij += A[i+k*n] * B[k+j*n]; /* cij += A[i][k] * B[k][j] */
11                 C[i+j*n] = cij; /* C[i][j] = cij */
12             }
13 }
14 void dgemm (int n, double* A, double* B, double* C)
15 {
16     for ( int sj = 0; sj < n; sj += BLOCKSIZE )
17         for ( int si = 0; si < n; si += BLOCKSIZE )
18             for ( int sk = 0; sk < n; sk += BLOCKSIZE )
19                 do_block(n, si, sj, sk, A, B, C);
20 }
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Blocked DGEMM Access Pattern



Four Questions for Caches and Memory Hierarchy


- **Q1: Where can a block be placed in the upper level? (Block placement)**
- **Q2: How is a block found if it is in the upper level? (Block identification)**
- **Q3: Which block should be replaced on a miss? (Block replacement)**
Assignment Project Exam Help
<https://tutorcs.com>
- **Q4: What happens on a write? (Write strategy)**
WeChat: cstutorcs

Q1: Where can a block be placed in the upper level?


■ Block 12 placed in 8 block cache:

- Fully associative, direct mapped, 2-way set associative
- S.A. Mapping = Block Number Modulo Number Sets


Fully associative:
block 12 can go
anywhere




Block-frame address



Direct mapped:
block 12 can go only
into block 4 (12 mod
8)



Set associative:
block 12 can go
anywhere in set 0
(12 mod 4)



Q2: How is a block found if it is in the upper level?



- Direct indexing (using index and block offset), tag compares, or combination
- Increasing associativity shrinks index, expands tag

Q3: Which block should be replaced on a miss?

- Easy for Direct Mapped
- Set Associative or Fully Associative:
 - Random
 - LRU (Least Recently Used)

Assignment Project Exam Help

<https://tutorcs.com>

| Associativity | 2 way | 2 way | 4 way | 4 way | 8 way | 8 way |
|---------------|-------|--------|-------|--------|-------|--------|
| Size | LRU | Random | LRU | Random | LRU | Random |
| 16 KB | 5.2% | 5.7% | 4.7% | 5.3% | 4.4% | 5.0% |
| 64 KB | 1.9% | 2.0% | 1.5% | 1.7% | 1.4% | 1.5% |
| 256 KB | 1.15% | 1.17% | 1.13% | 1.13% | 1.12% | 1.12% |

Random Replacement (Don McLane, UW)

- Build a single Pseudorandom Number generator for the WHOLE cache. On a miss, roll the dice and throw out a cache line at random.
- Updates only on misses.
- How do you build a random number generator (it's easier than you might think).

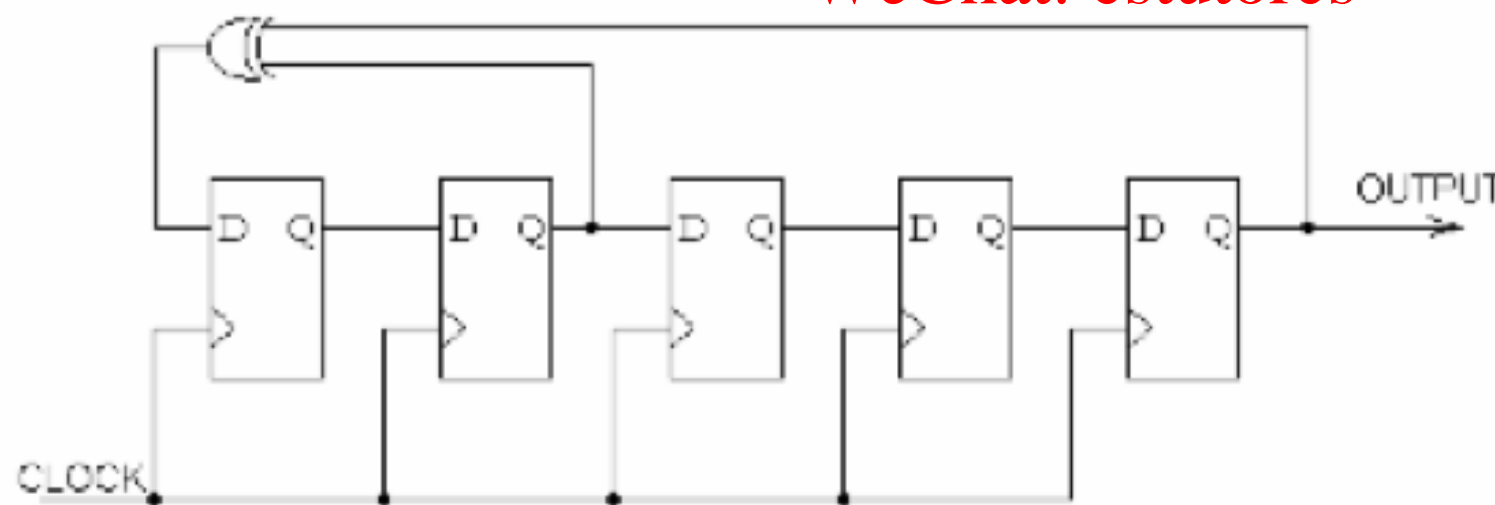


Overhead is
 $O(\log_2 N)$
 bits/cache!

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs




Pseudorandom Linear Feedback Shift Register

Counting Sequence

| | | | |
|-------|------|-------|------|
| 11111 | 0x1F | 01000 | 0x08 |
| 01111 | 0x0F | 10100 | 0x14 |
| 00111 | 0x07 | 01010 | 0x0A |
| 10011 | 0x13 | 10101 | 0x15 |
| 11001 | 0x19 | 11010 | 0x1A |
| 01100 | 0x0C | 11101 | 0x1D |
| 10110 | 0x16 | 01110 | 0x0E |
| 01011 | 0x0B | 10111 | 0x17 |
| 00101 | 0x05 | 11011 | 0x1B |
| 10010 | 0x12 | 01101 | 0x0D |
| 01001 | 0x09 | 00110 | 0x06 |
| 00100 | 0x04 | 00011 | 0x03 |
| 00010 | 0x02 | 10001 | 0x11 |
| 00001 | 0x01 | 11000 | 0x18 |
| 10000 | 0x10 | 11100 | 0x1C |
| | | 11110 | 0x1E |


3Cs Absolute Miss Rate (SPEC92)



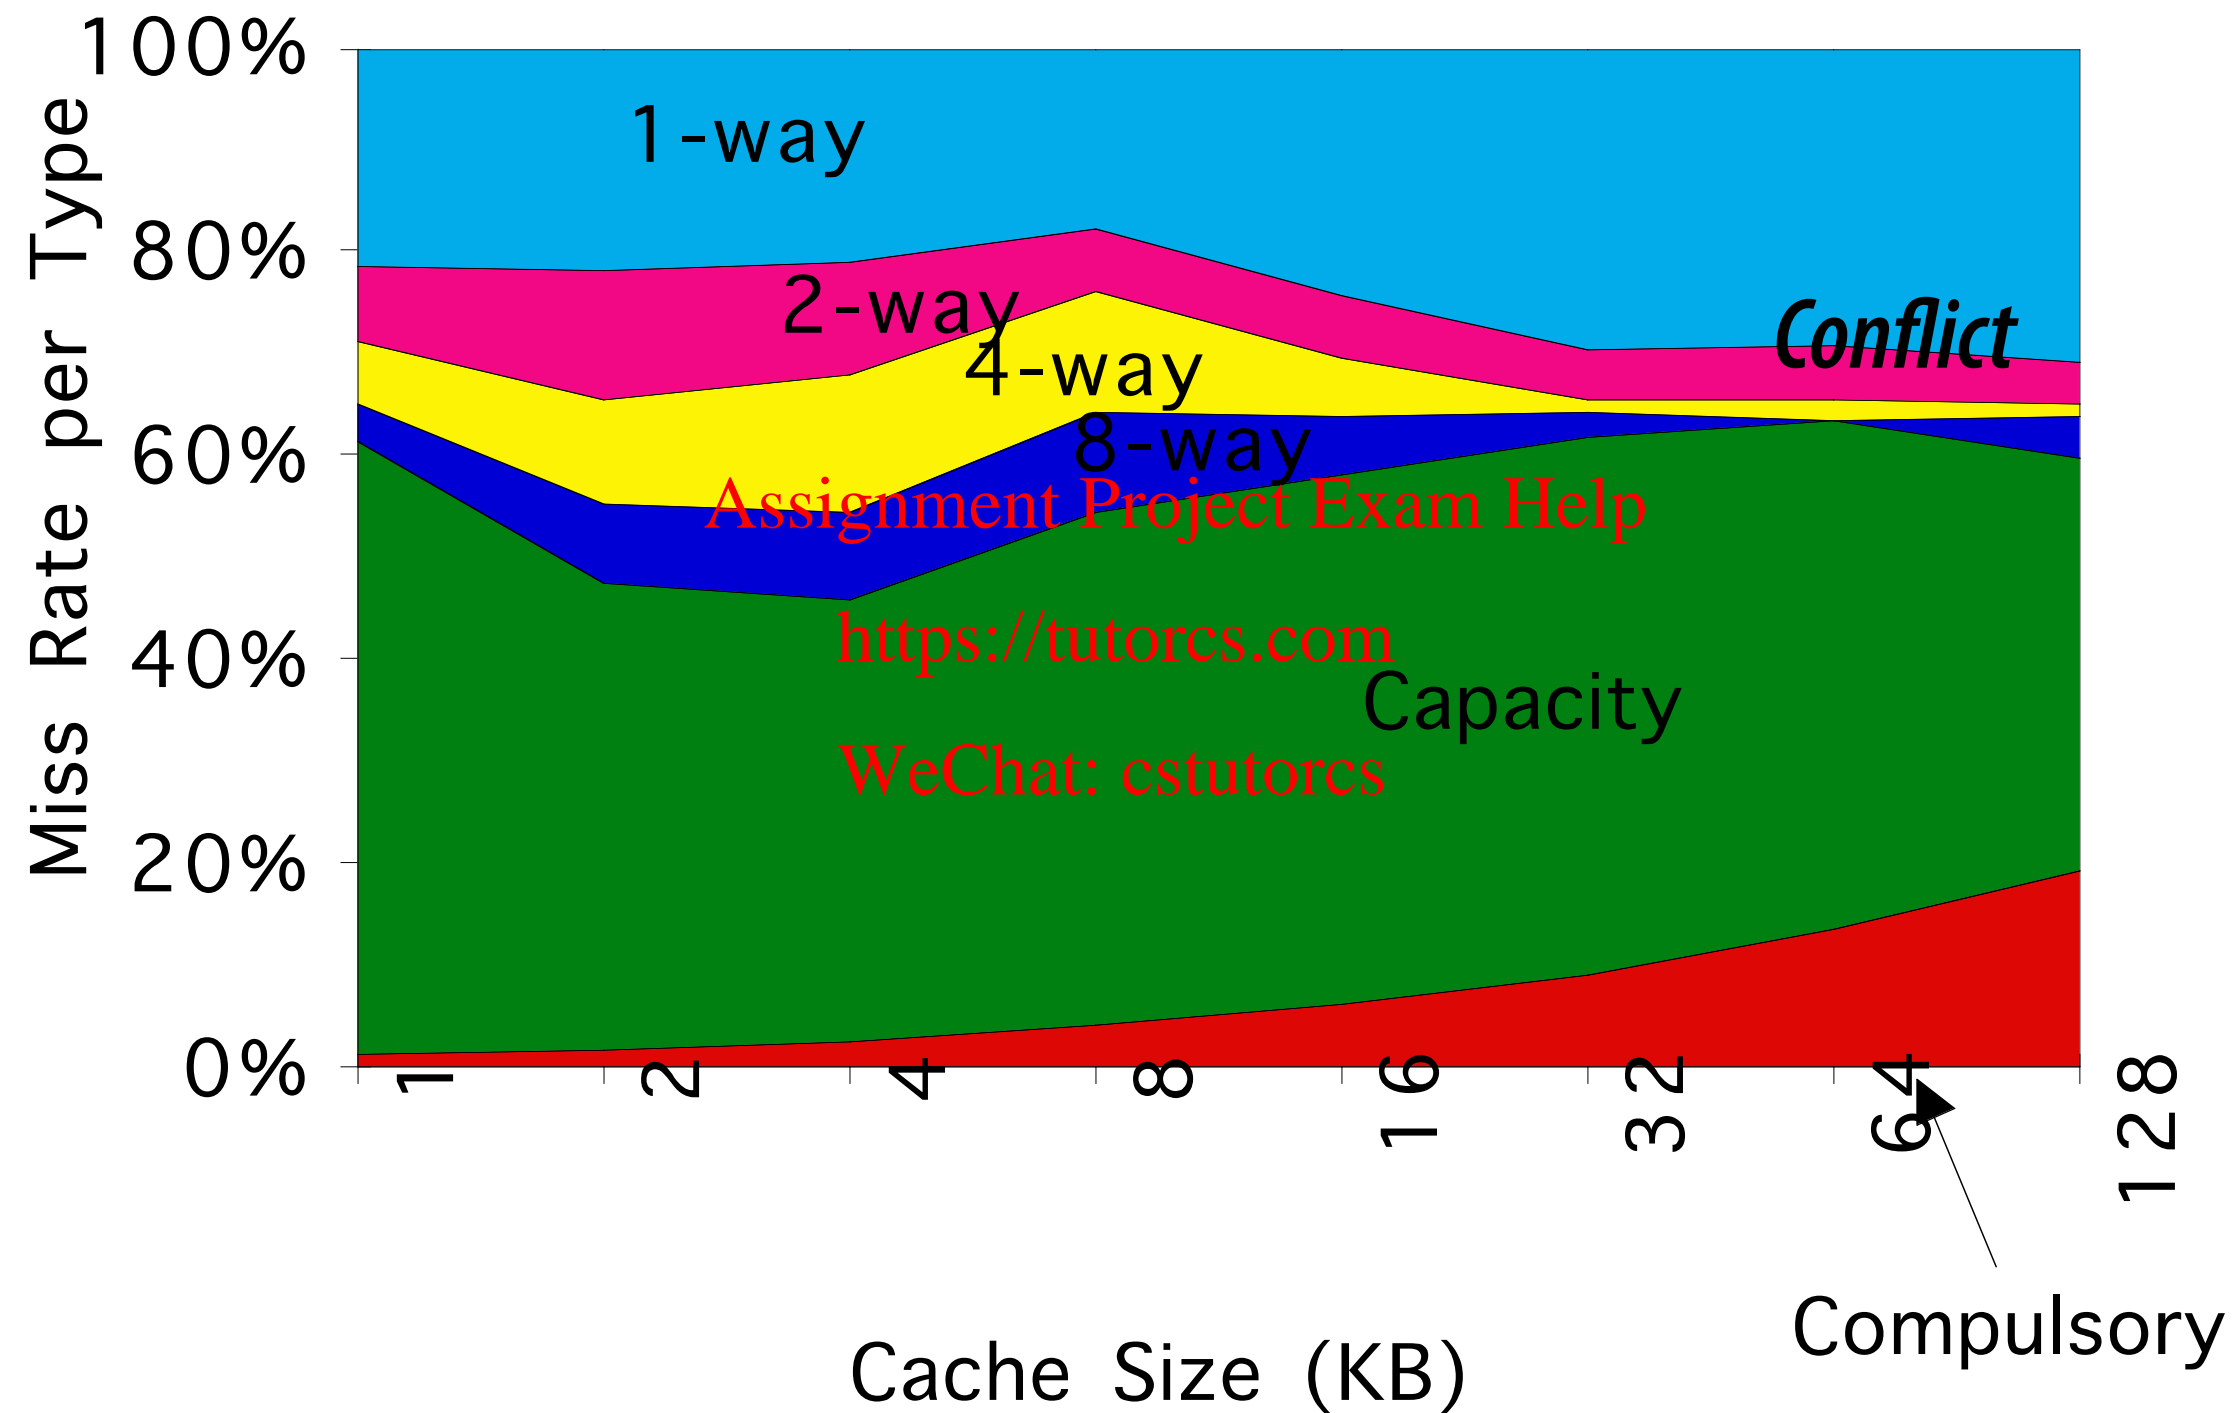
2:1 Cache Rule

miss rate 1-way associative cache size X

■ *= miss rate 2-way associative cache size $X/2$*



3Cs Relative Miss Rate



Q4: What happens on a write?


- What happens on a write miss?
 - Don't have to read data or check tags - just write!
- Writes are more complex than reads:
 - Desirable: Cache is always a subset of memory ("consistent")
 - So when we write to the cache, we also should write to the memory as well ("write through")
 - If we don't do this: "write back"
 - Write-through could lead to performance problems - writing to memory could mean performance is proportional to memory time instead of cache time
 - Solution: "write buffer"

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Write Buffer for Write Through



- A Write Buffer is needed between the Cache and Memory
 - Processor: writes data into the cache and the write buffer
 - Memory controller: write contents of the buffer to memory
- Write buffer is just a FIFO:
 - Typical number of entries: 4
 - Must handle bursts of writes
 - Works fine if: Store frequency (w.r.t. time) $\ll 1 / \text{DRAM write cycle}$

Q4: What happens on a write?

- **Write through**—The information is written to both the block in the cache and to the block in the lower-level memory.
 - On a write miss, do we allocate space in the cache? Typically no, since all writes have to go to main memory anyway. This is called *write no-allocate*.
- **Write back**—The information is written only to the block in the cache. The modified cache block is written to main memory only when it is replaced.
 - On a write miss, do we allocate space in the cache? Typically yes, since we hope future writes might also be captured by the cache. This is called *write allocate*.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Q4: What happens on a write?

- **Write through**—The information is written to both the block in the cache and to the block in the lower-level memory.
- **Write back**—The information is written only to the block in the cache. The modified cache block is written to main memory only when it is replaced.
[Assignment Project Exam Help](https://tutorcs.com)
 - is block clean or dirty?
- **Pros and Cons of each?** [WeChat: cstutorcs](#)
 - **WT**: read misses cannot result in writes, also simpler
 - **WB**: no writes of repeated writes (so less bandwidth), more complex (dirty bit)
- **WT always combined with write buffers so that don't wait for lower level memory**