

Lecture 10:

Implementing a Processor 2/5

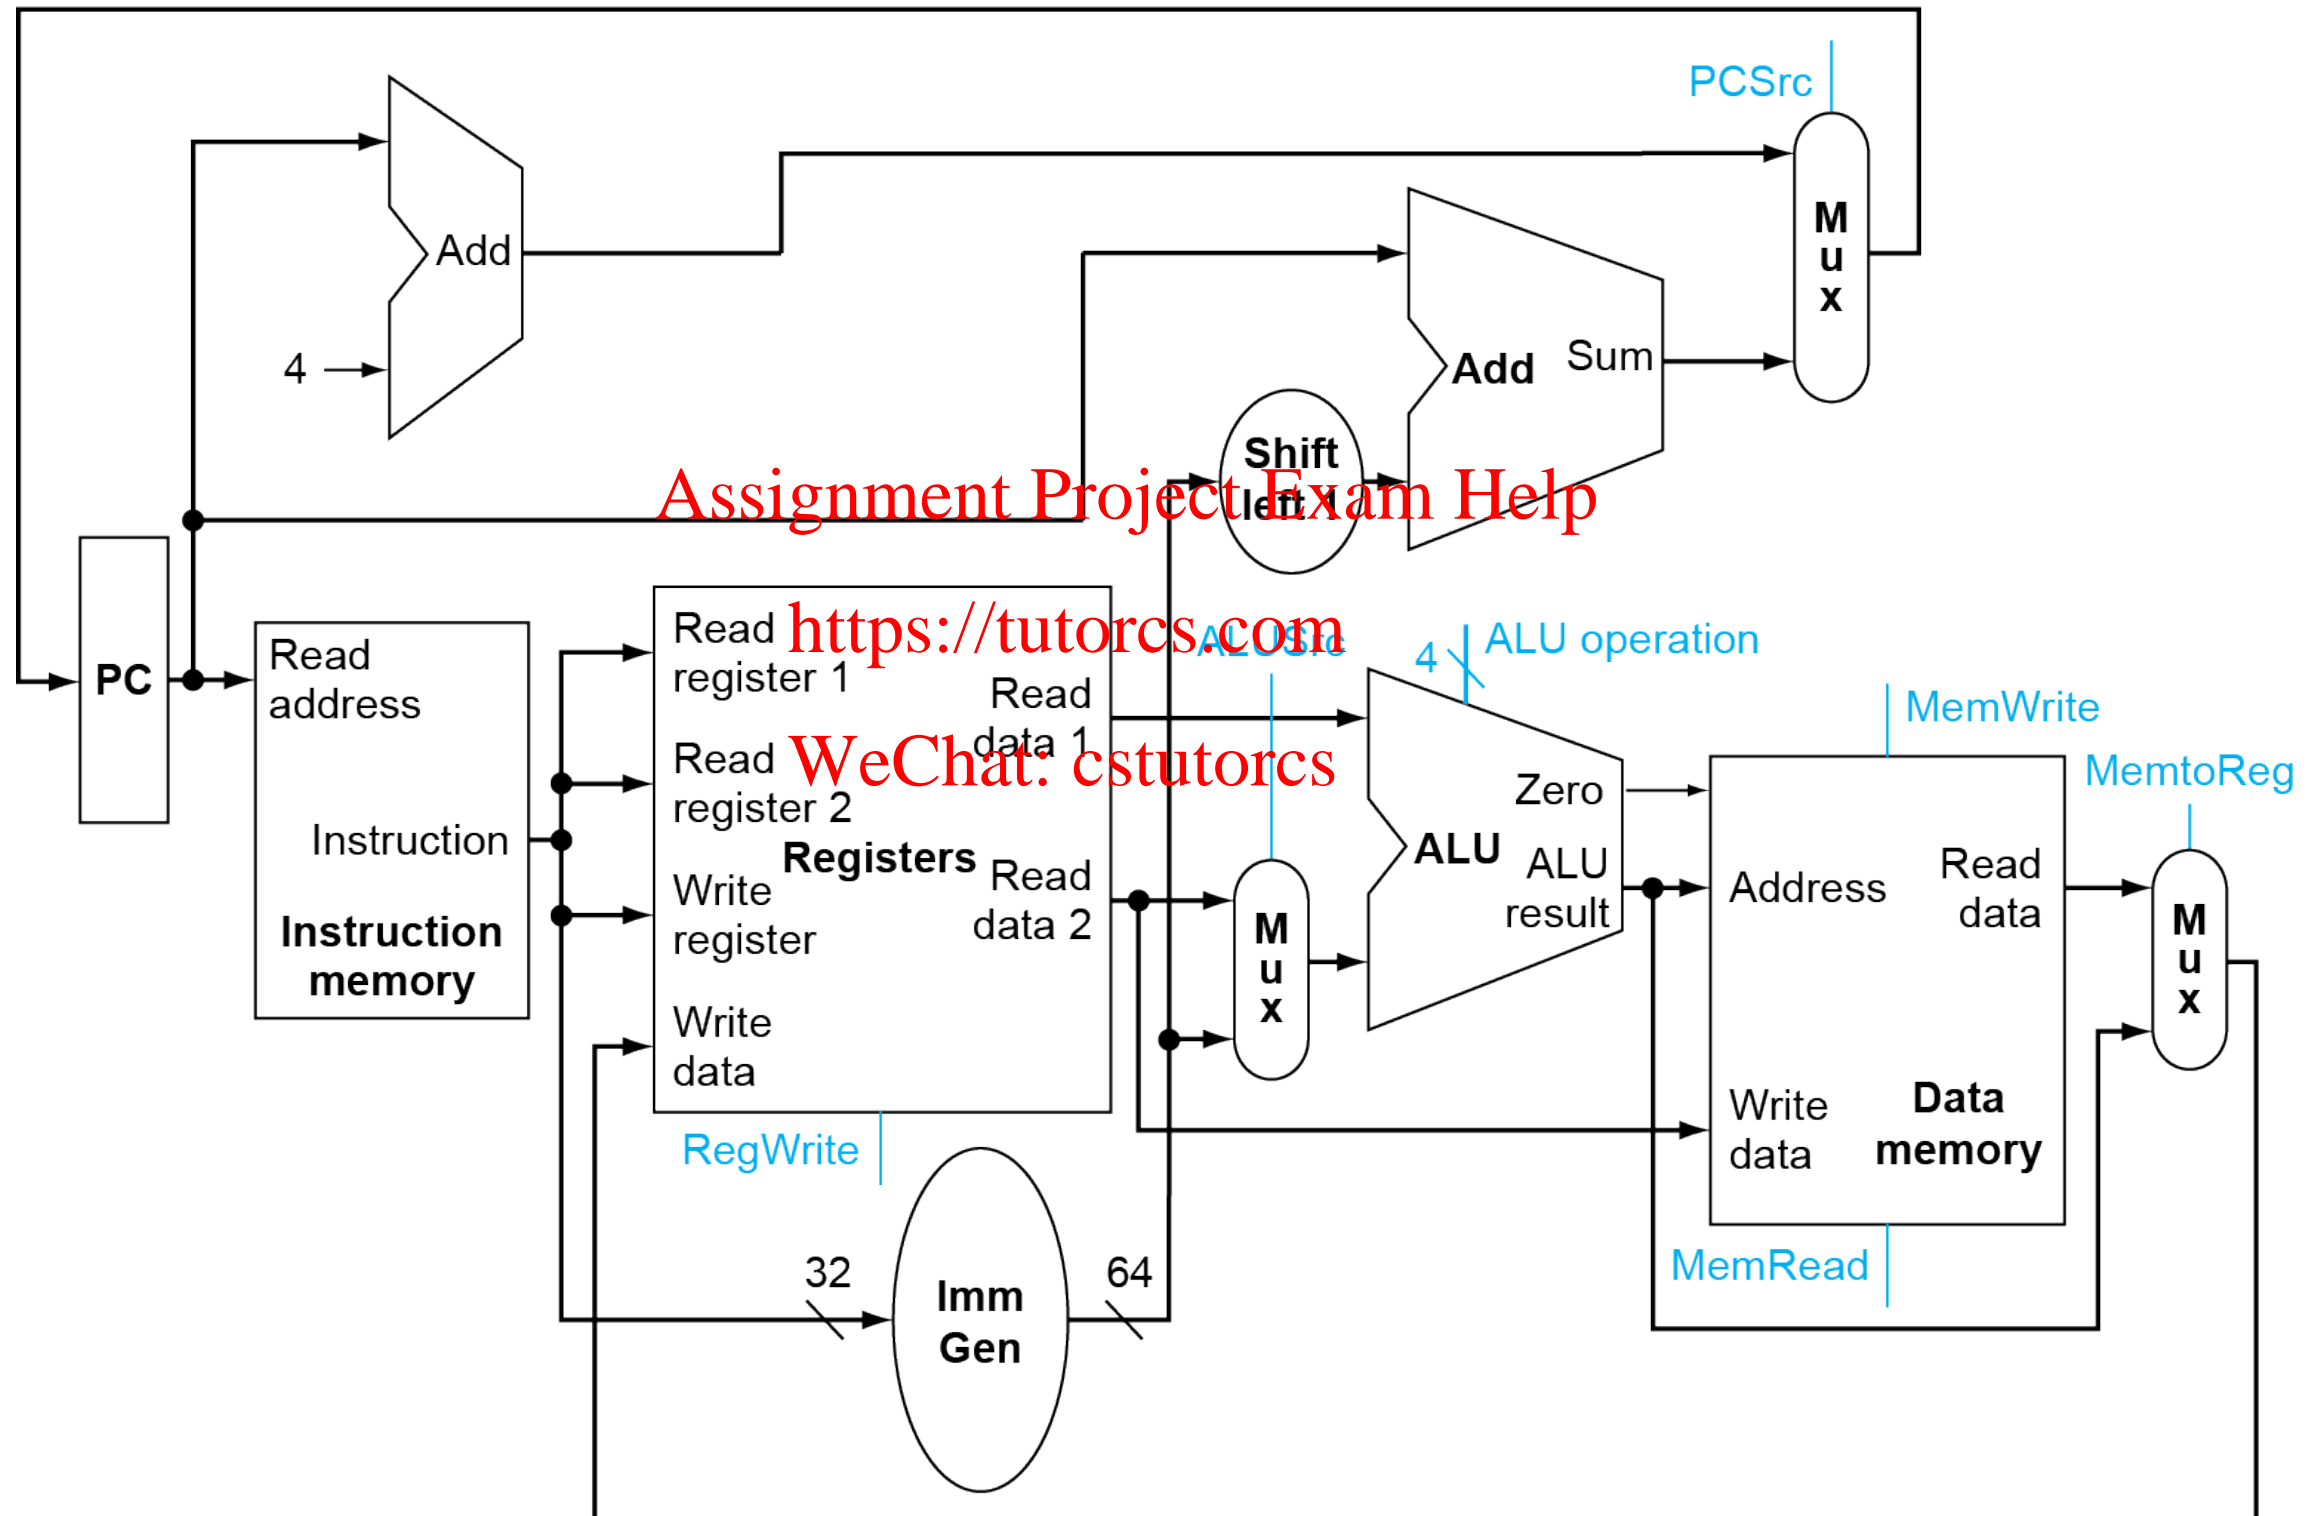
Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

**Introduction to Computer Architecture
UC Davis EEC 170, Fall 2019**


Introduction

- **CPU performance factors**
 - **Instruction count**
 - **Determined by ISA and compiler**
 - **CPI and Cycle time**
 - **Determined by CPU hardware**
- **We will examine two RISC implementations**
 - **A simplified version**
 - **A more realistic pipelined version**
- **Simple subset, shows most aspects**
 - **Memory reference: ld, sd**
 - **Arithmetic/logical: add, sub, and, or**
 - **Control transfer: beq**

Full Datapath



Datapath With Control



Implementing Logic

- Possibly faster to factor logic into multiple stages
- Recall opcode (6 bits) + funct7 (7 bits) + funct3 (3 bits)
 - Likely a 16-bit function will be very complex!
- Instead, can we do this in stages? Decode the opcode separately from the other bits.
 - Choice of encodings for opcodes can make this much easier

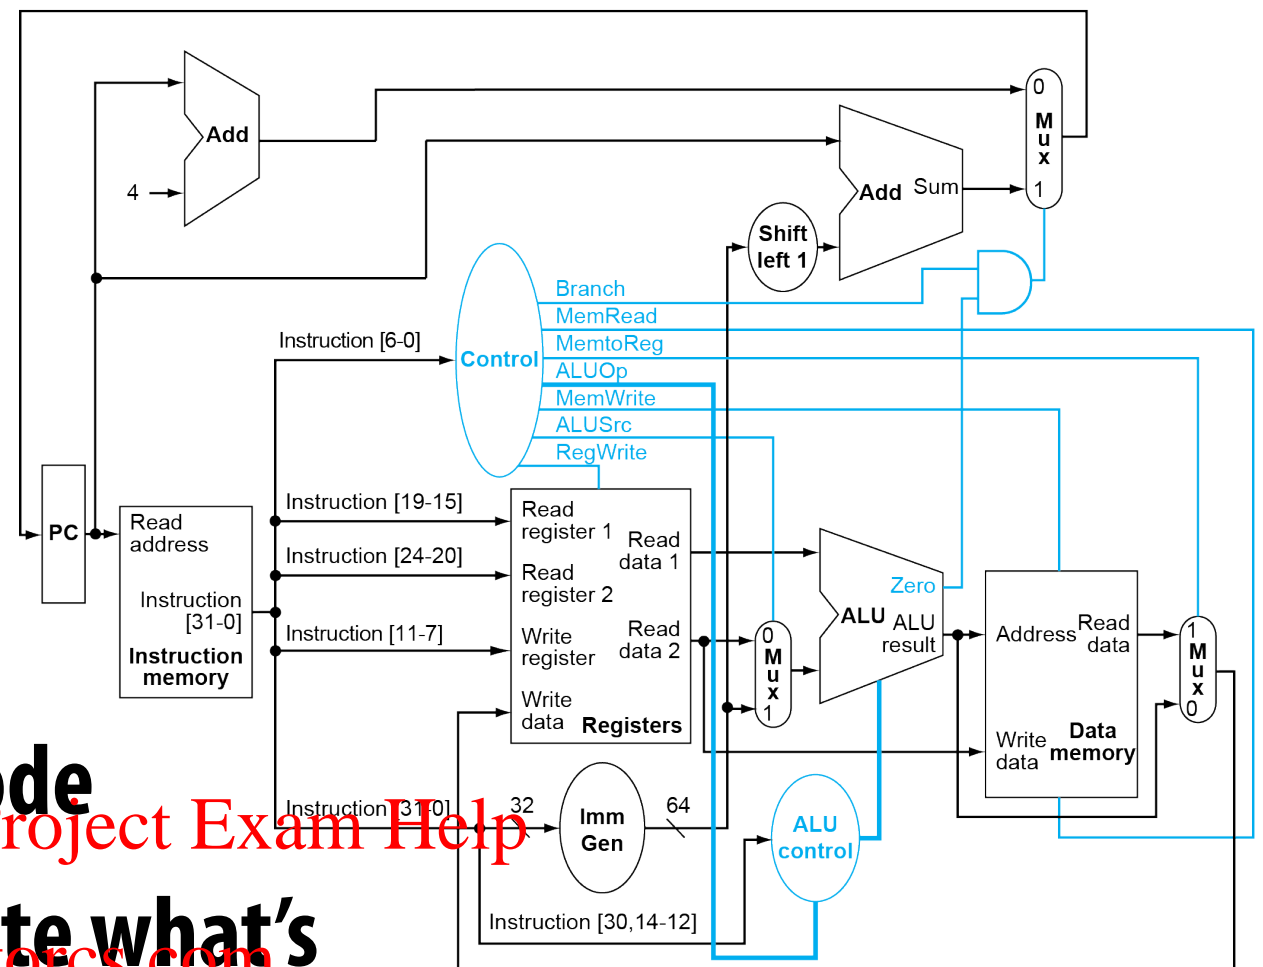
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

ALU Control

- ALU used for
 - Load/Store: F = add
 - Branch: F = subtract
 - R-type: F depends on opcode
- We want to eventually generate what's below. How? Next slide.



ALU control	Function
0000	AND
0001	OR
0010	add
0110	subtract

ALU Control

- Assume 2-bit ALUOp derived from opcode (bits [6:0])
 - Then combinational logic derives ALU control
- Goal here: Set “ALU Control” as a function of ALUOp + funct7/3

opcode	ALUOp	Operation	funct7/3 fields	ALU function	ALU control
ld	00	load word	XXXXXXXXXXXX	add	0010
sd	00	store dword	XXXXXXXXXXXX	add	0010
beq	01	branch on equal	XXXXXXXXXXXX	subtract	0110
R-type	10	add	0000000/000	add	0010
		subtract	0100000/000	subtract	0110
		AND	0000000/111	AND	0000
		OR	0000000/110	OR	0001

The Main Control Unit

■ Control signals derived from instruction

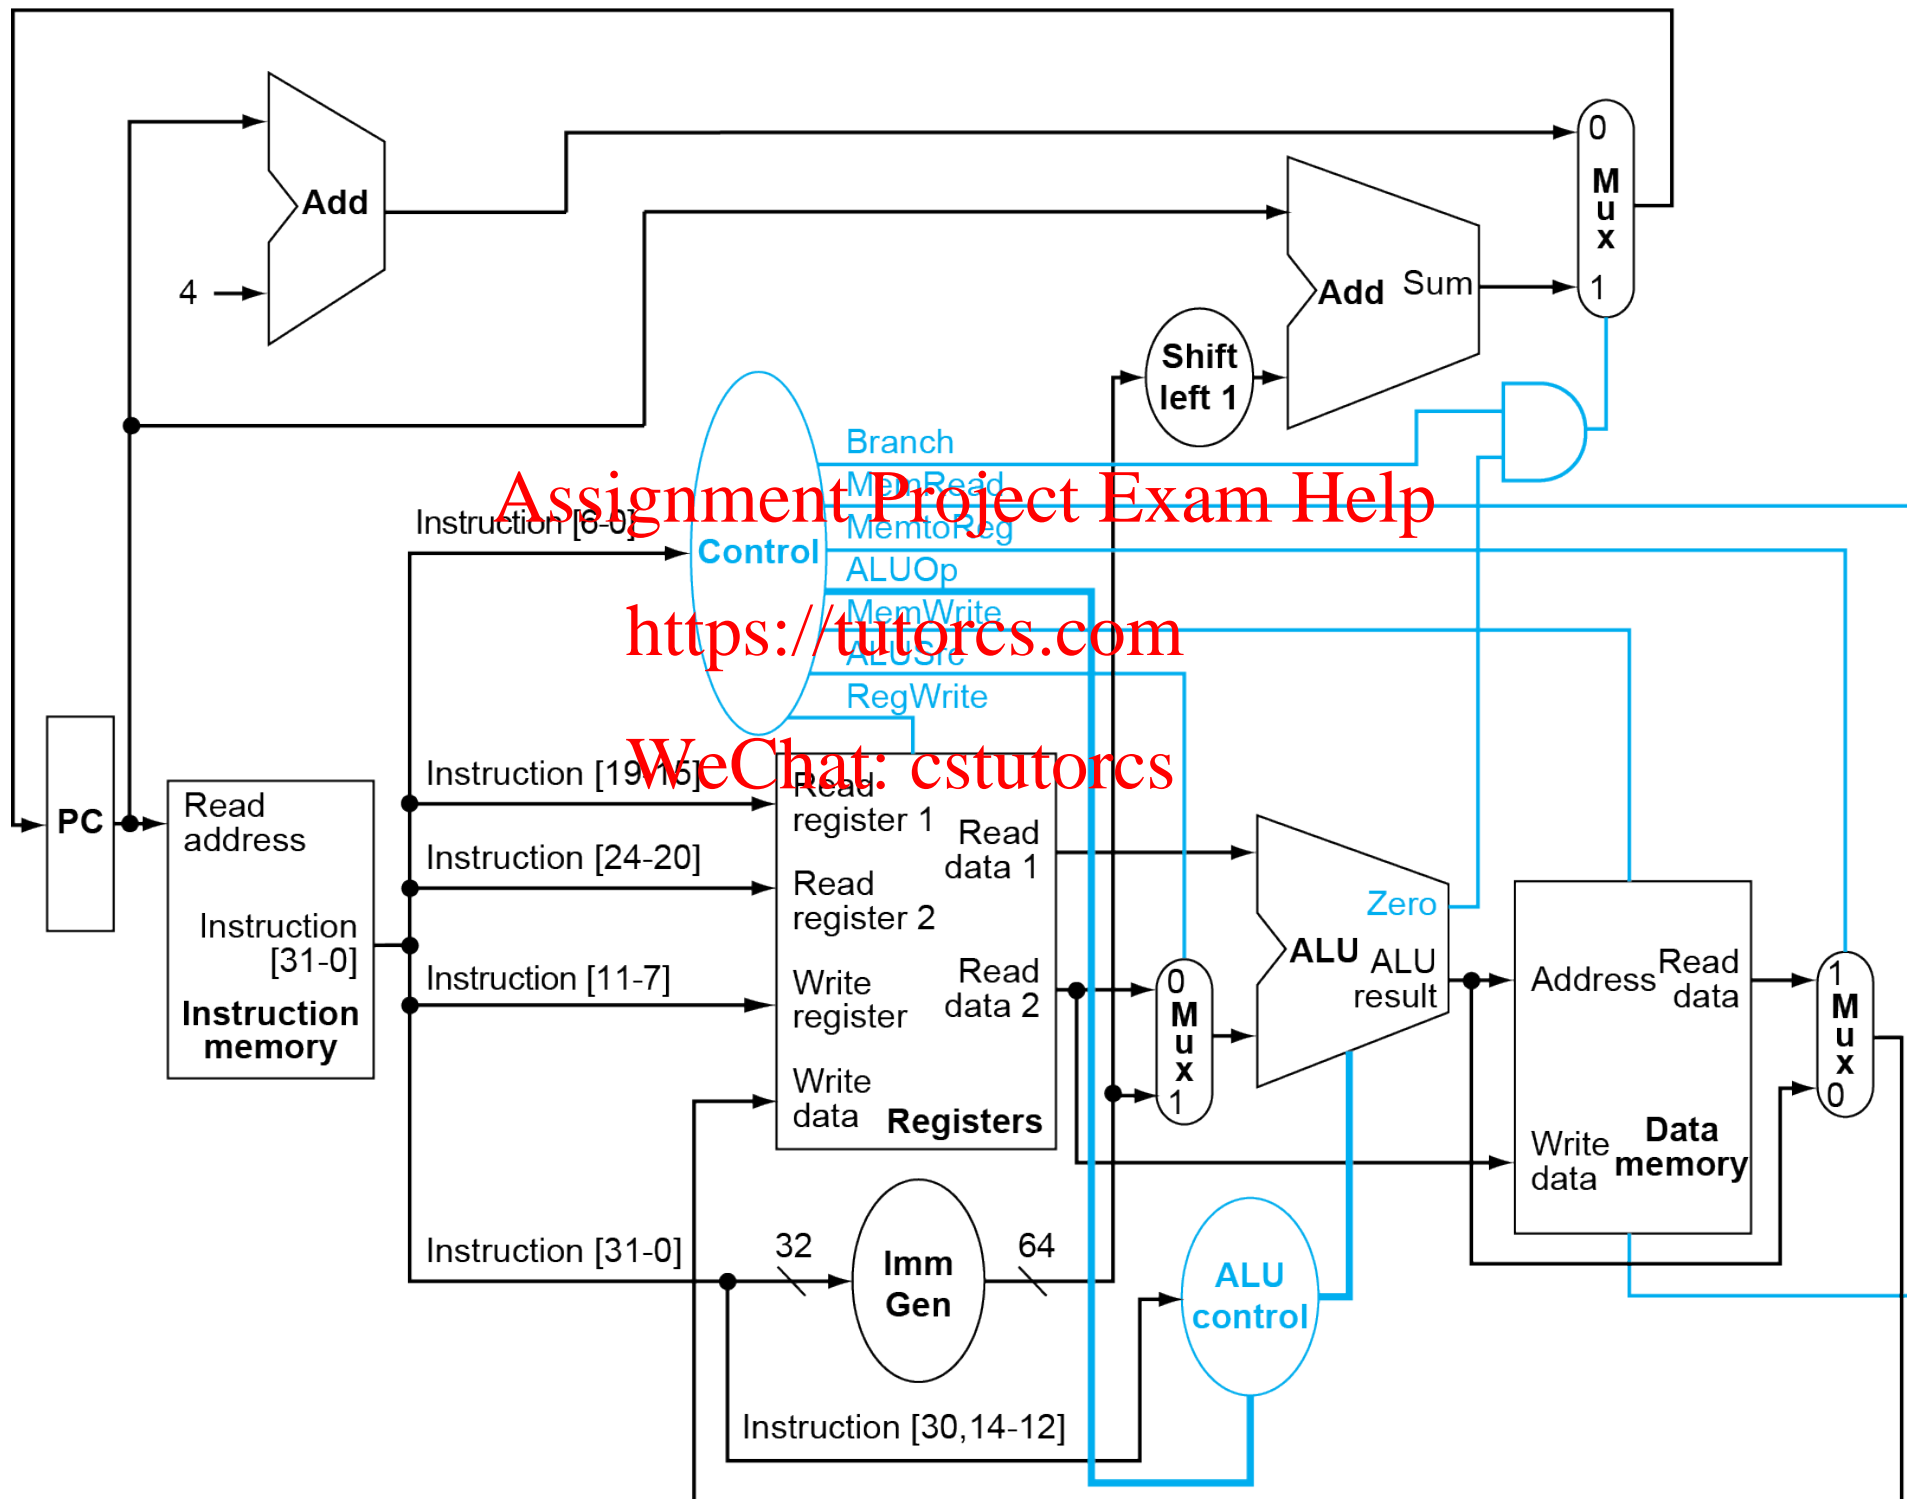
opcode	ALUOp	ALU control	Function
ld	00	0000	AND
sd	00	0001	OR
beq	01	0010	add
R-type	10	0110	subtract

Name (Bit position)	31:25	24:20	19:15	14:12	11:7	6:0
(a) R-type	funct7	rs2	rs1	funct3	rd	opcode
(b) I-type	immediate[11:0]	rs2	rs1	funct3	rd	opcode
(c) S-type	immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode
(d) SB-type	immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode

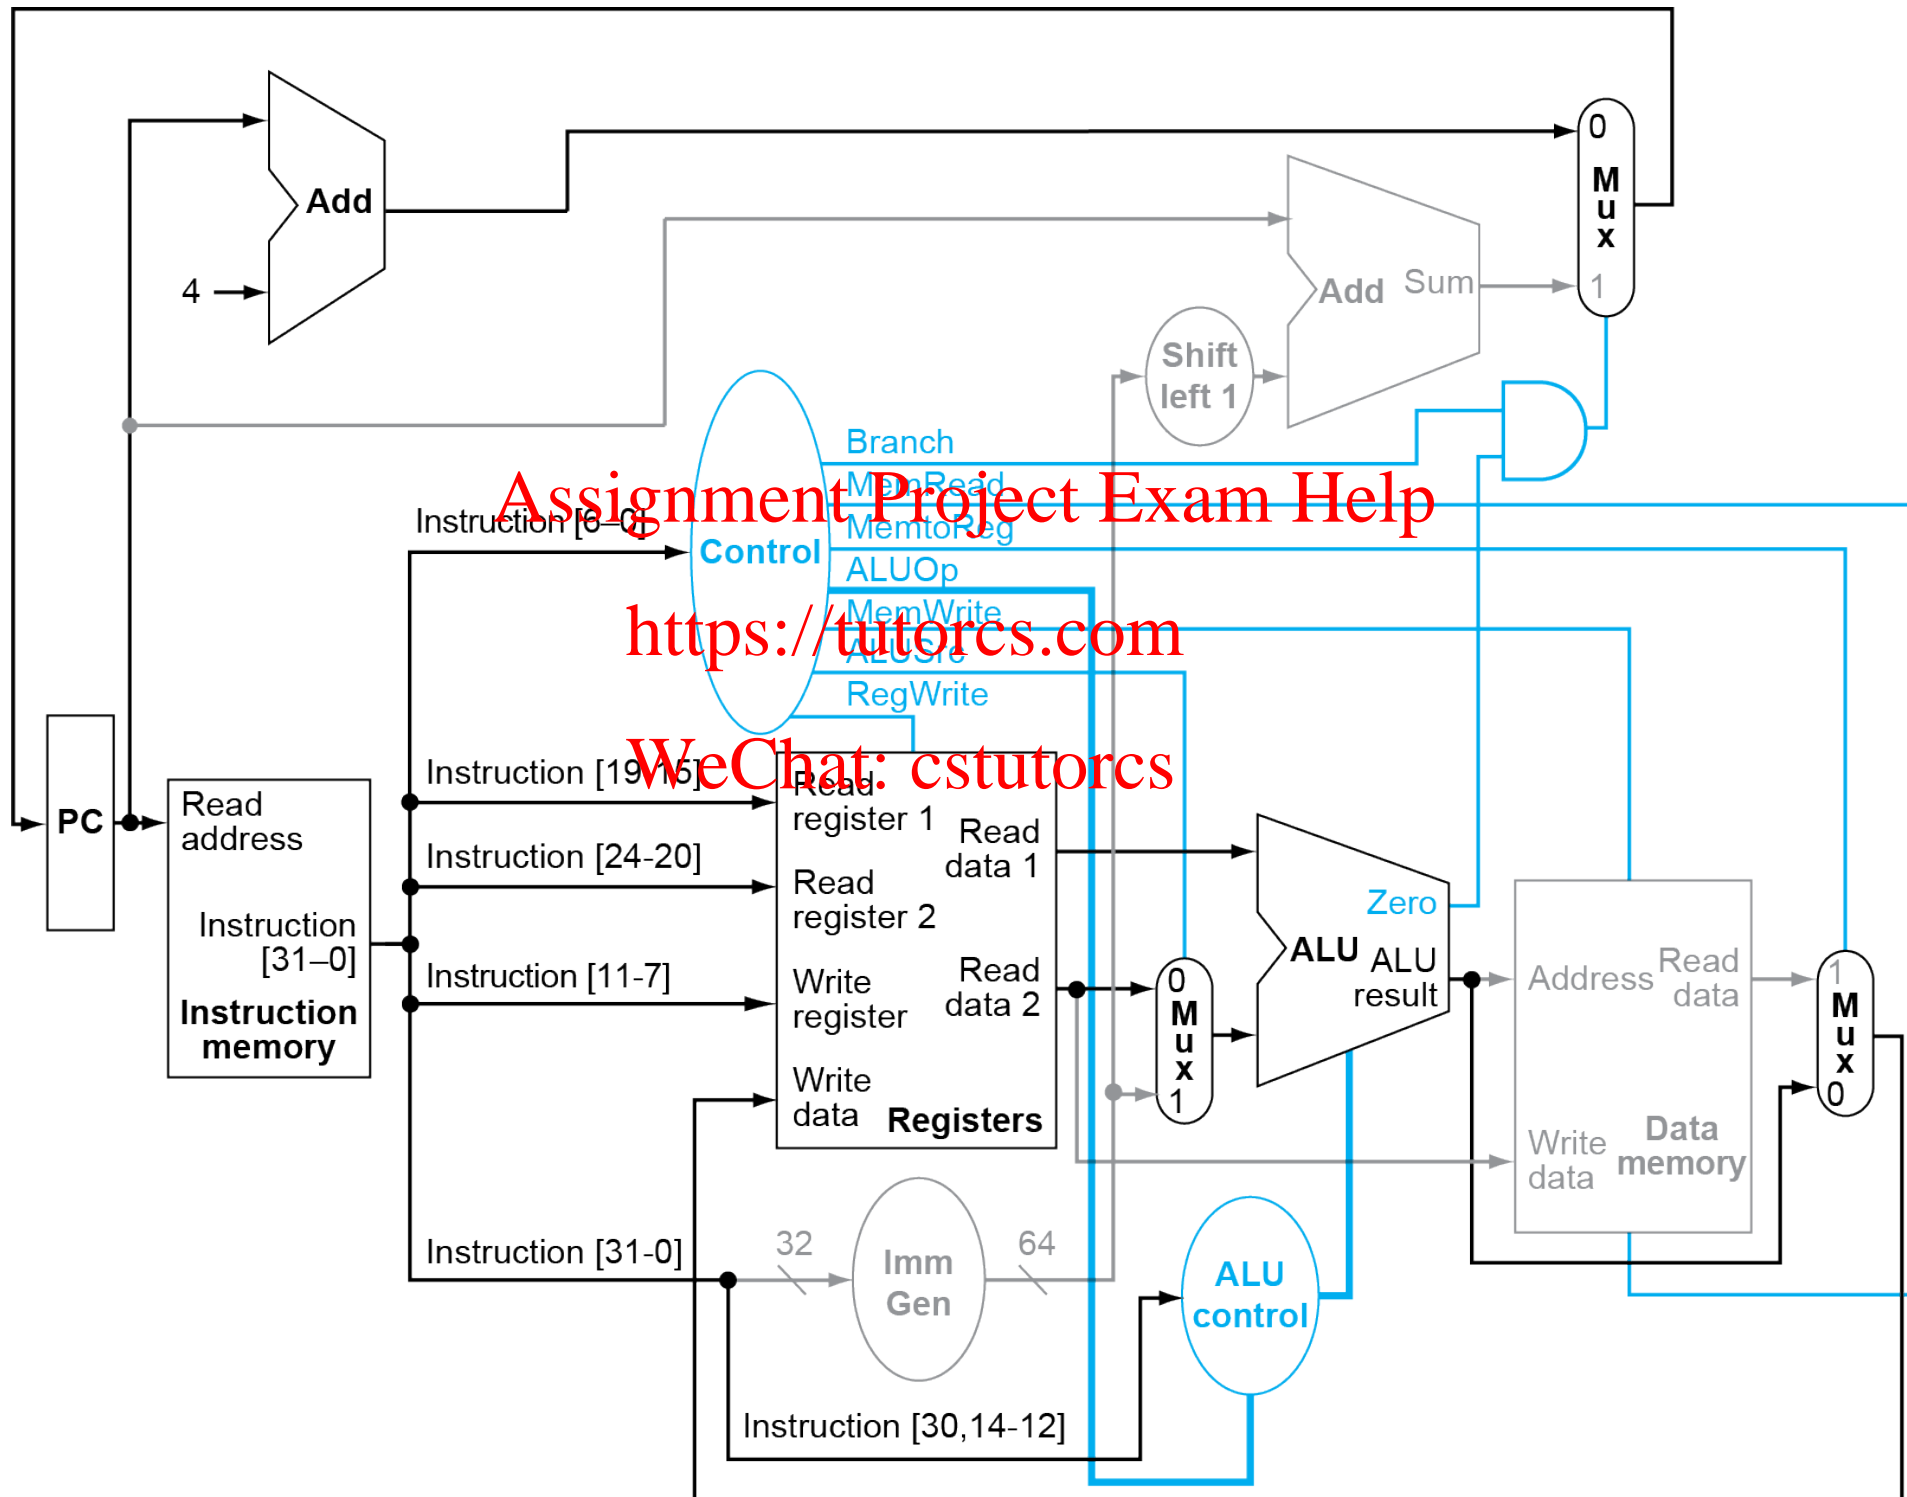
ALUOp		Funct7 field							Funct3 field			Operation
ALUOp1	ALUOp0	I[31]	I[30]	I[29]	I[28]	I[27]	I[26]	I[25]	I[14]	I[13]	I[12]	
0	0	X	X	X	X	X	X	X	X	X	X	0010
X	1	X	X	X	X	X	X	X	X	X	X	0110
1	X	0	0	0	0	0	0	0	0	0	0	0010
1	X	0	1	0	0	0	0	0	0	0	0	0110
1	X	0	0	0	0	0	0	0	1	1	1	0000
1	X	0	0	0	0	0	0	0	1	1	0	0001

these four bits are the only ones that matter for generating "operation = ALU control"

Datapath With Control



R-Type Instruction




Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Load Instruction

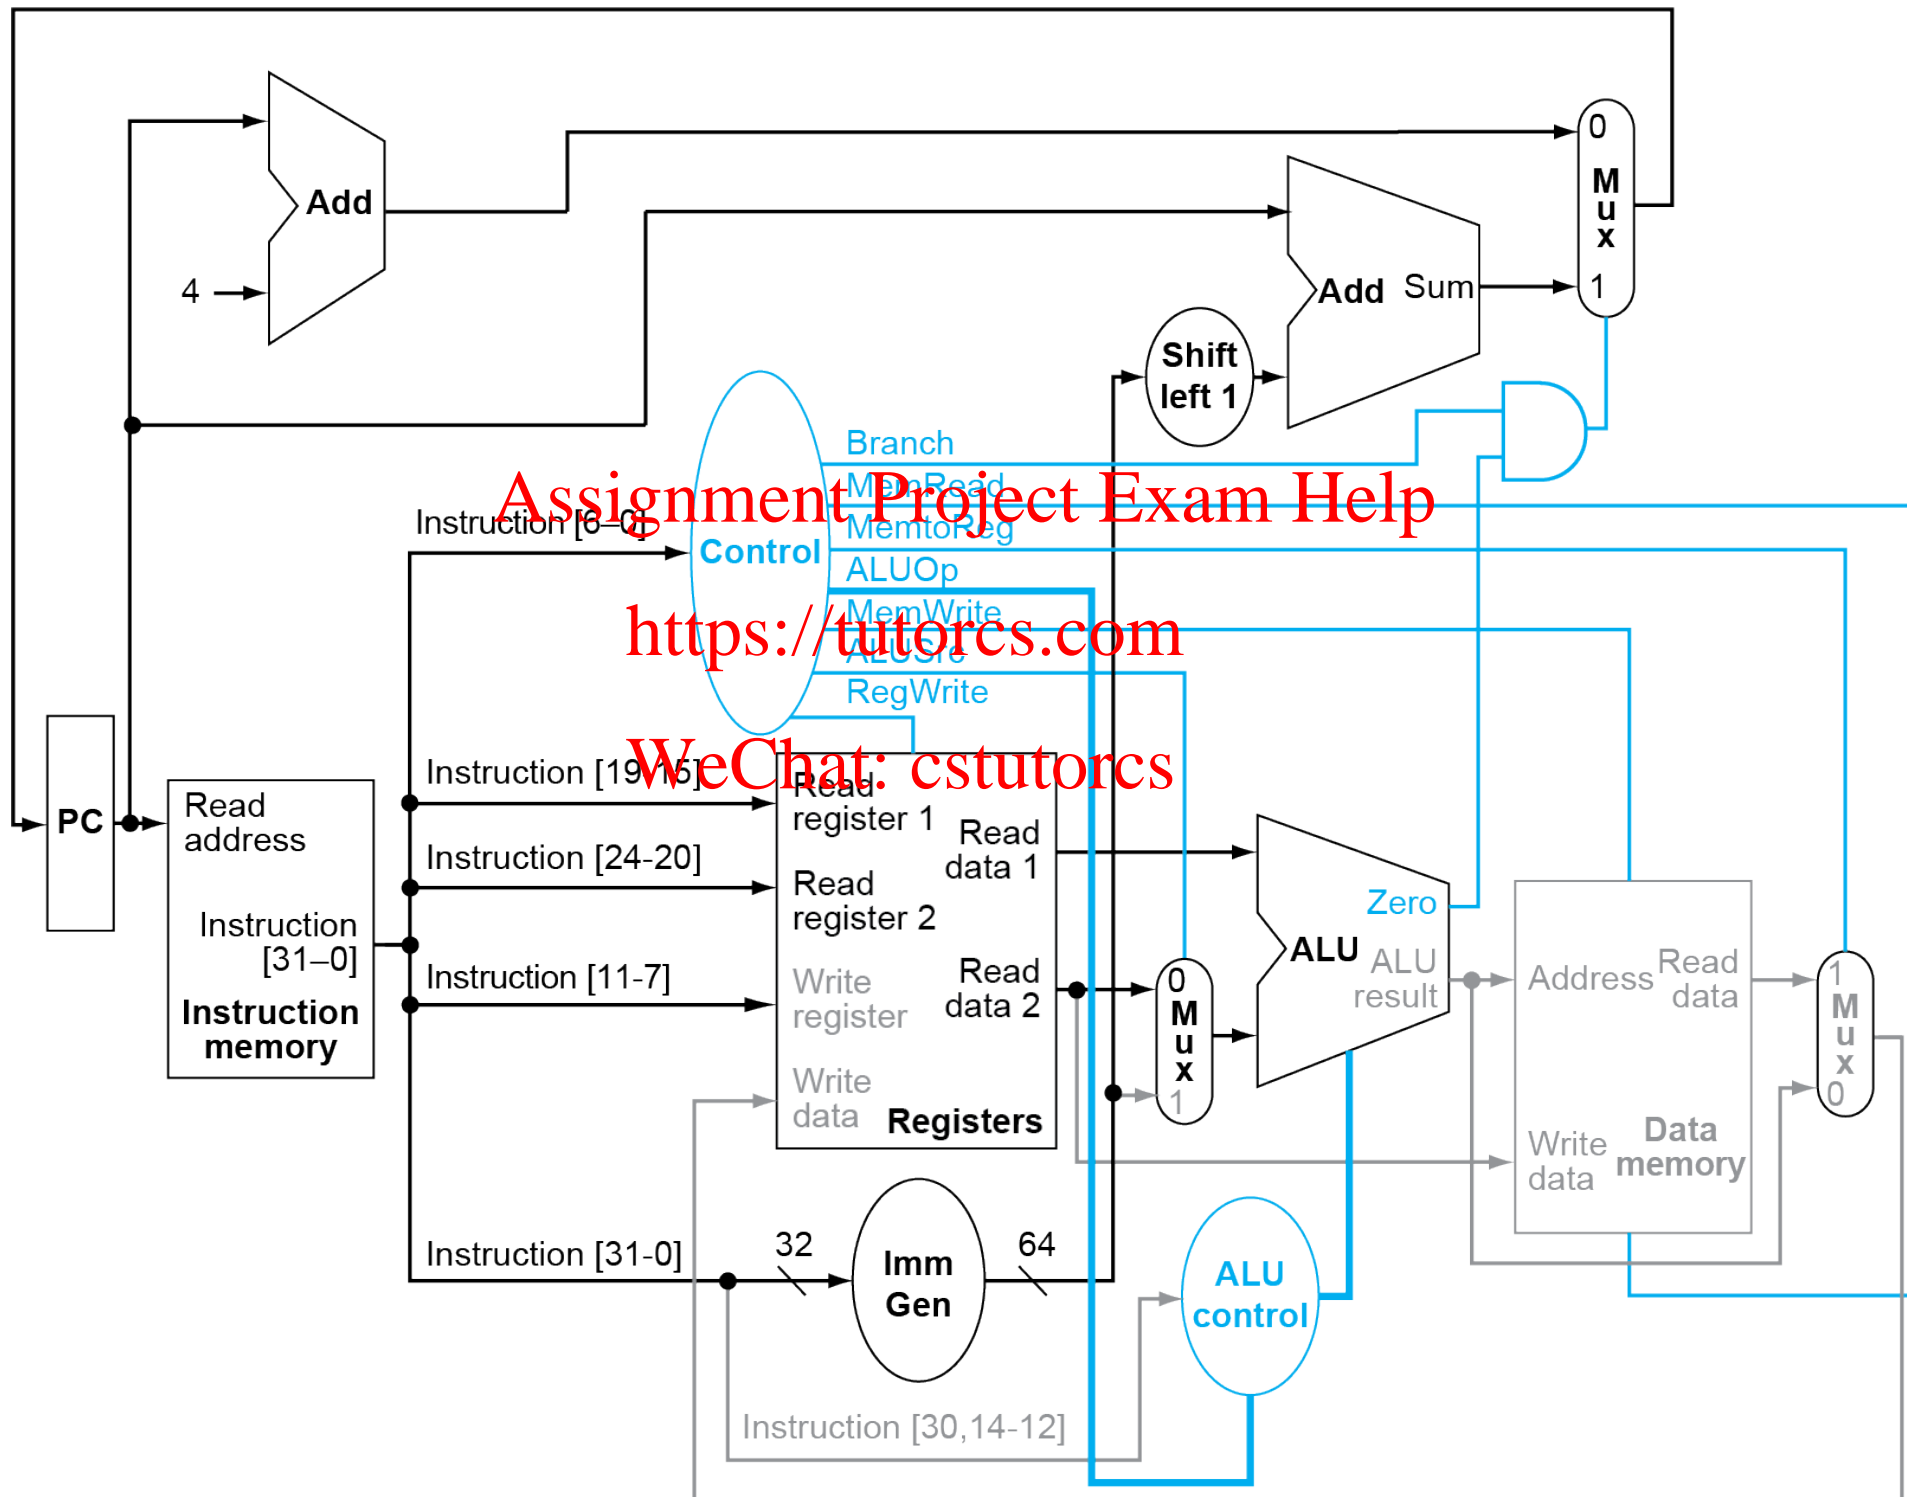


Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

BEQ Instruction



Building an ISA

- **Should we add this other instruction?**
- **How often would we use it?**
- **What are the changes to the {datapath, control}?**
 - **How complicated is the additional logic?**
 - **Do we have the design time to add it?**
 - **Would it affect the critical path?**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Good Exam Questions

- What are the datapath changes necessary to support new instruction X?
- Given a datapath, what are the control signals necessary to perform instruction Y?
- Given a set of control signals and opcodes, what is the logic for generating control signal Z?

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Performance Issues

- Longest delay determines clock period
 - Critical path: load instruction
 - Instruction memory → register file → ALU → data memory → register file
 - But other instructions don't take this long!
- Not feasible to vary period for different instructions
 - Clock distribution, among other issues
- Violates design principle
 - Making the common case fast
- We will improve performance by *pipelining*



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Pipelining Analogy

- **Pipelined laundry: overlapping execution**
 - **Parallelism improves performance**
- **How many “stages” is this laundry process?**



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Four loads:

- **Speedup**

$$= 8 / 3.5 = 2.3$$

Non-stop:

- **Speedup**

$$= 2n / 0.5n + 1.5 \approx 4$$

$$= \text{number of stages}$$

RISC-V Pipeline

- **Five stages, one step per stage**
 1. **IF: Instruction fetch from memory**
 2. **ID: Instruction decode & register read**
 3. **EX: Execute operation or calculate address**
 4. **MEM: Access memory operand**
 5. **WB: Write result back to register**

Assignment Project Exam Help

<https://tutores.com>

WeChat: ctutores

Pipeline Performance

- Assume time for stages is
 - 100ps for register read or write
 - 200ps for other stages
- Compare pipelined datapath with single-cycle datapath

Assignment Project Exam Help

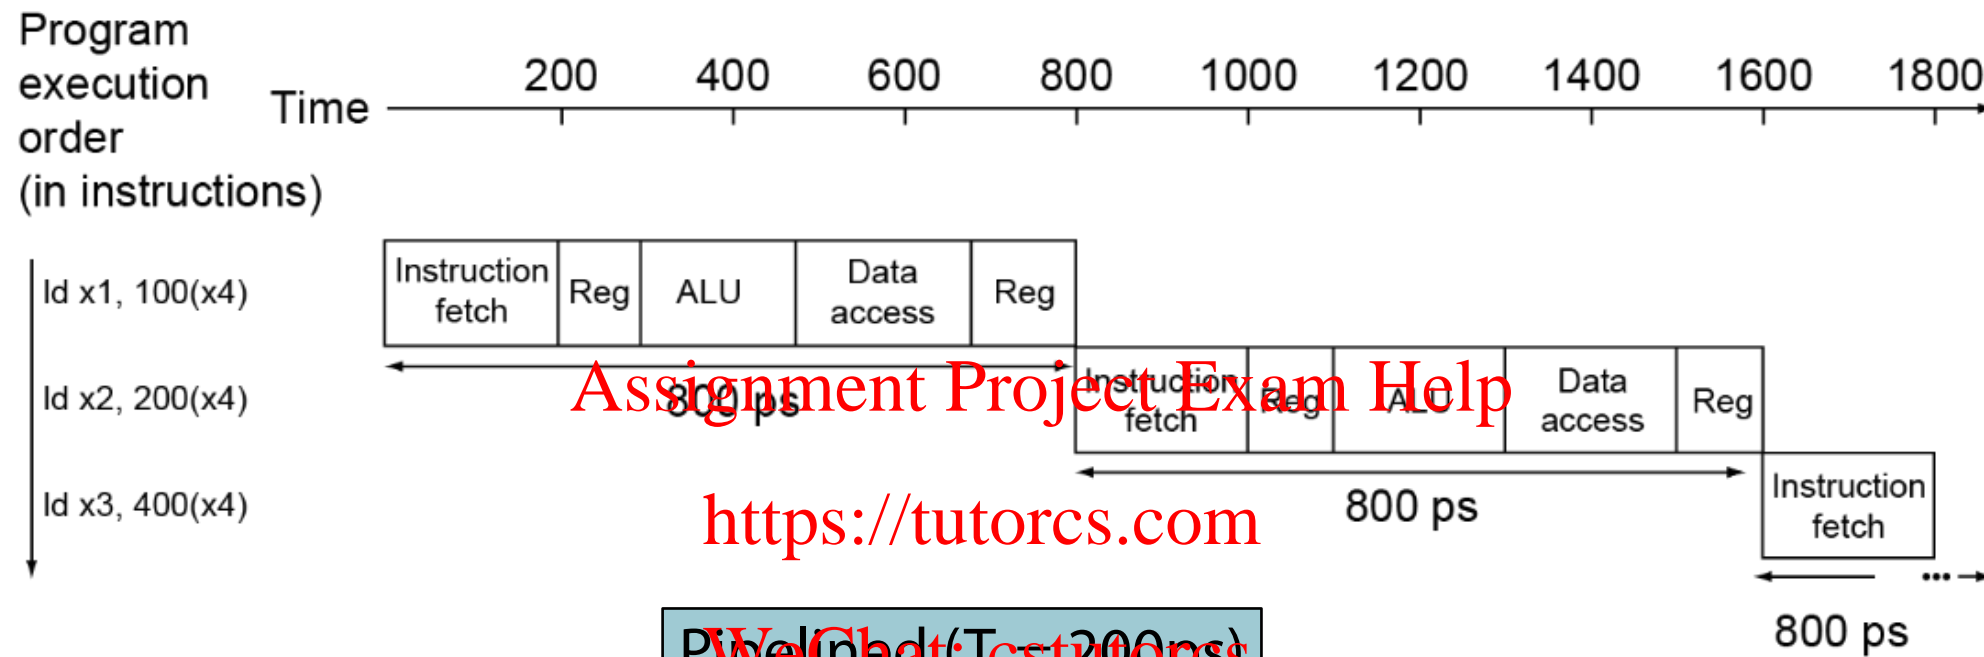
<https://tutorcs.com>

WeChat: cstutorcs

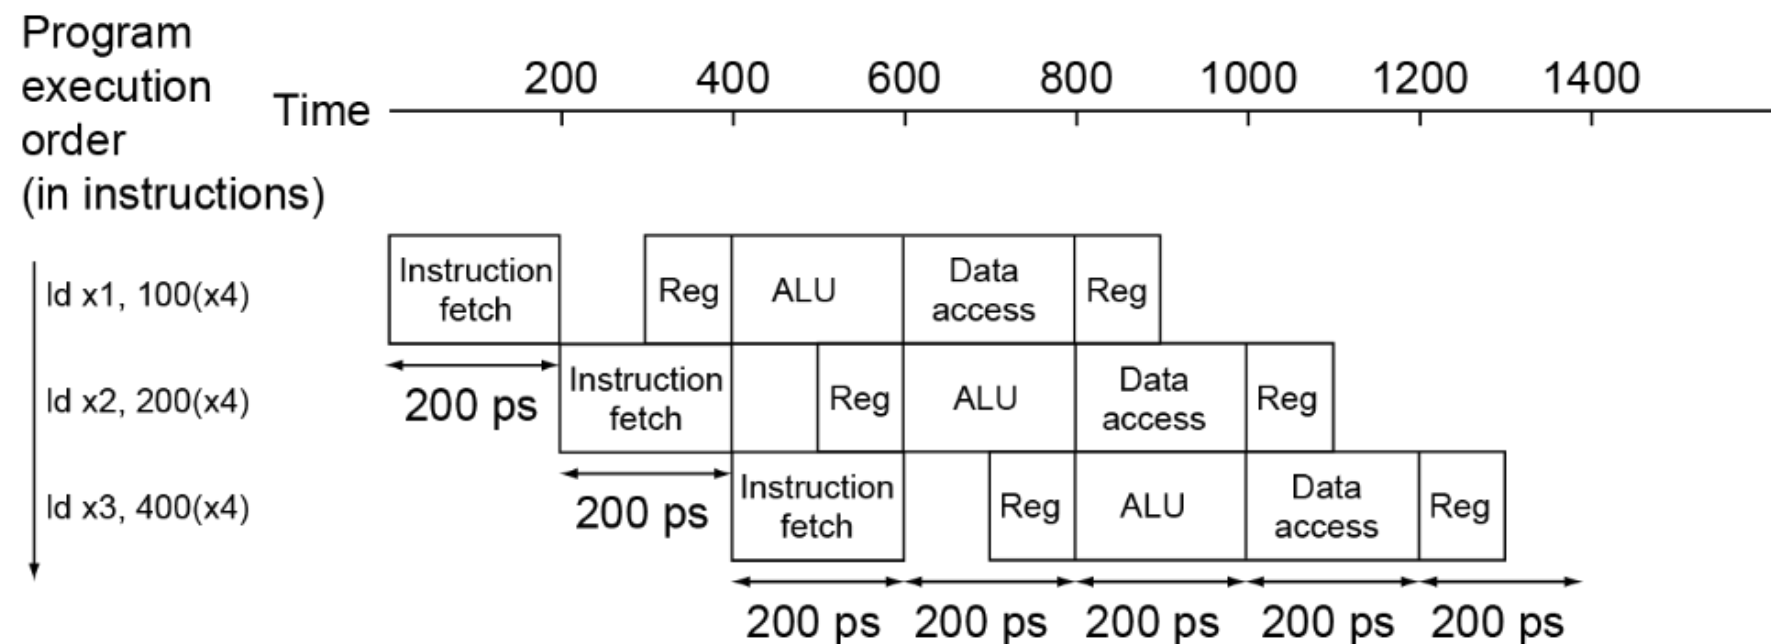
Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
ld	200ps	100 ps	200ps	200ps	100 ps	800ps
sd	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

Pipeline Performance

Single-cycle ($T_c = 800\text{ps}$)



Pipeline ($T_c = 200\text{ps}$)



Pipeline Speedup

- If all stages are “balanced”:
 - i.e., all take the same time
 - then: $\text{Time between instructions}_{\text{pipelined}} = \frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of stages}}$
- If not balanced, speedup is less
- Speedup is due to increased throughput
 - Latency (time for each instruction) does not decrease
 - *Pipelining helps throughput, not latency*

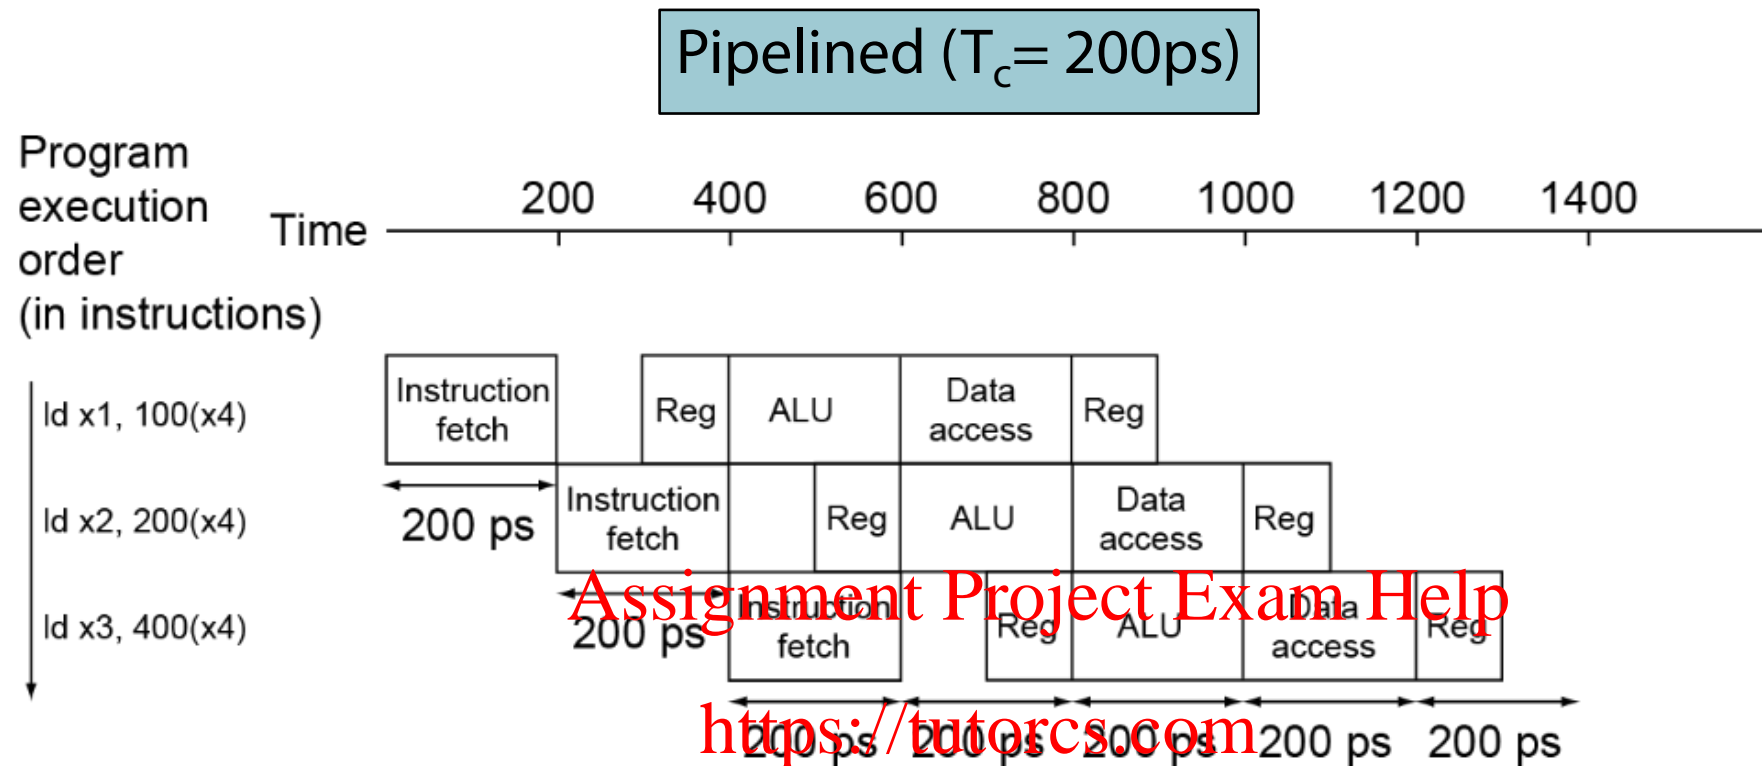
Assignment Project Exam Help

Number of stages

<https://tutorcs.com>

WeChat: cstutorcs

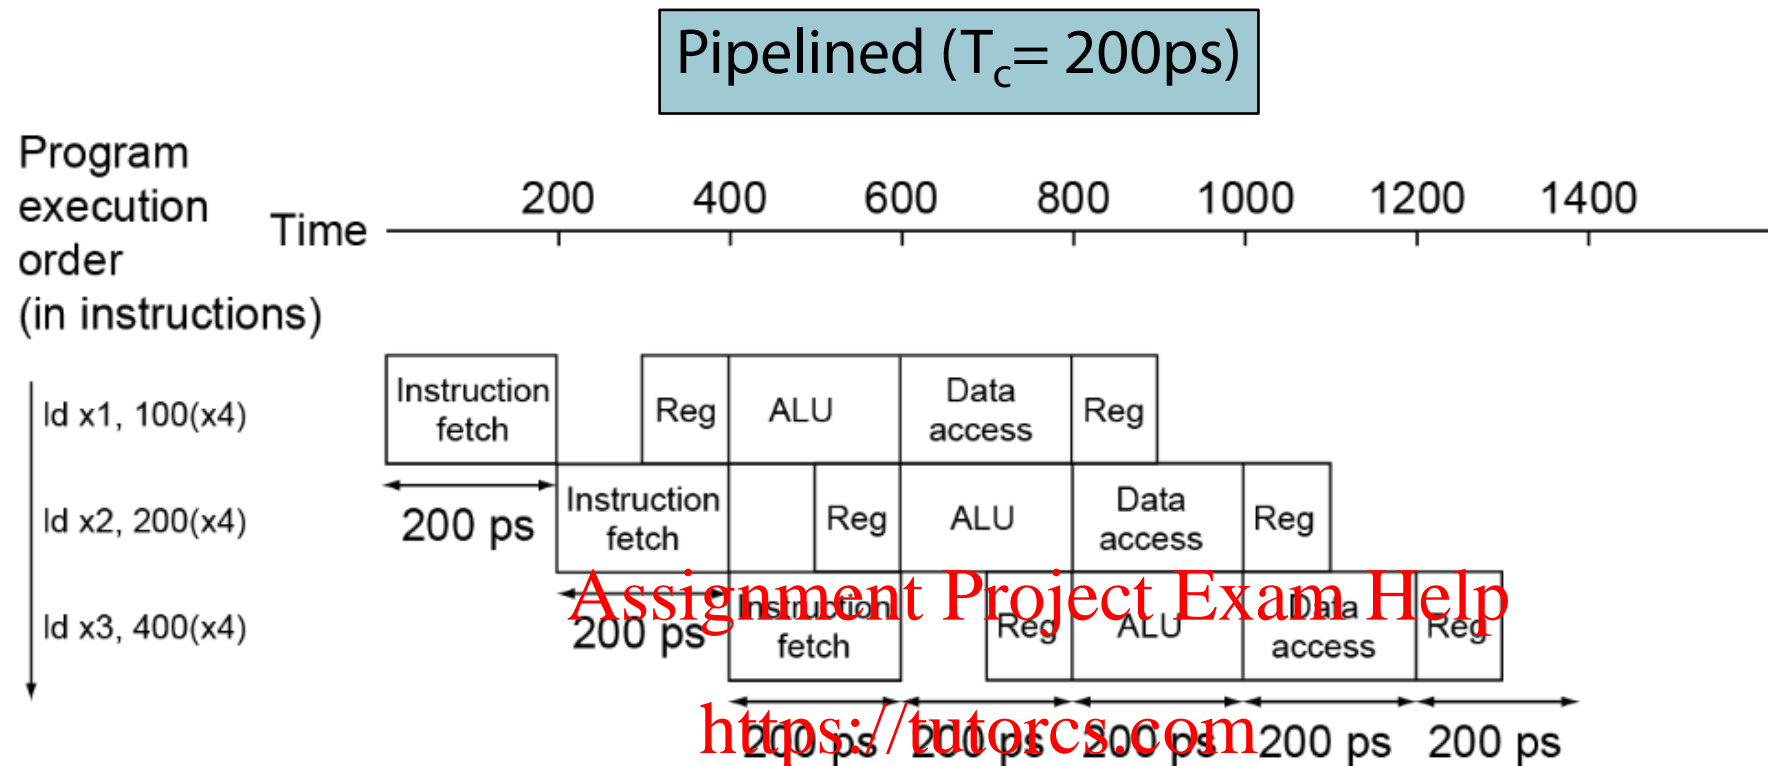
Fetching the next instruction



WeChat: cstutorcs

- **For “instruction fetch”, must:**
 - **Go to memory and get instruction**
 - **Decode instruction to figure out how long it is**
 - **Increment program counter appropriately to point to next instruction**

Fetching the next instruction



WeChat: cstutorcs

■ For “decode”, must:

- **Decode instruction to figure out what kind of instruction it is**
- **Based on what instruction it is, fetch registers from the RF**
 - **That's a lot to do! How do we make it fast?**

Pipelining and ISA Design

- **RISC-V ISA designed for pipelining**
 - **All instructions are 32 bits**
 - **This is a critical aspect of the RISC philosophy**
 - **Easier to fetch and decode in one cycle**
 - **c.f. x86: 1- to 17-byte instructions**
 - **Few and regular instruction formats**
 - **Can decode and read registers in one step**
 - **Load/store addressing**
 - **Can calculate address in 3rd stage, access memory in 4th stage**

Assignment Project Exam Help

<https://tutores.com>

WeChat: cstutores

RISC-V Registers are always in the same place!

If we read one register, it's here

If we read a second register, it's here

CORE INSTRUCTION FORMATS

	31	27	26	25	24	20	19	15	14	12	11	7	6	0	
R	funct7				rs2	rs1			funct3	rd		opcode			
I	imm[11:0]				rs1			funct3	rd		opcode				
S	imm[11:5]				rs2	rs1			funct3	imm[4:0]		opcode			
SB	imm[12 10:5]				rs2	rs1			funct3	imm[4:1 11]		opcode			
U	imm[31:12]										rd		opcode		
UJ	imm[20 10:1 11 19:12]										rd		opcode		

If we write to a register, it's here

- We don't even have to decode an instruction to start register access!

Recap: RISC-V Pipeline

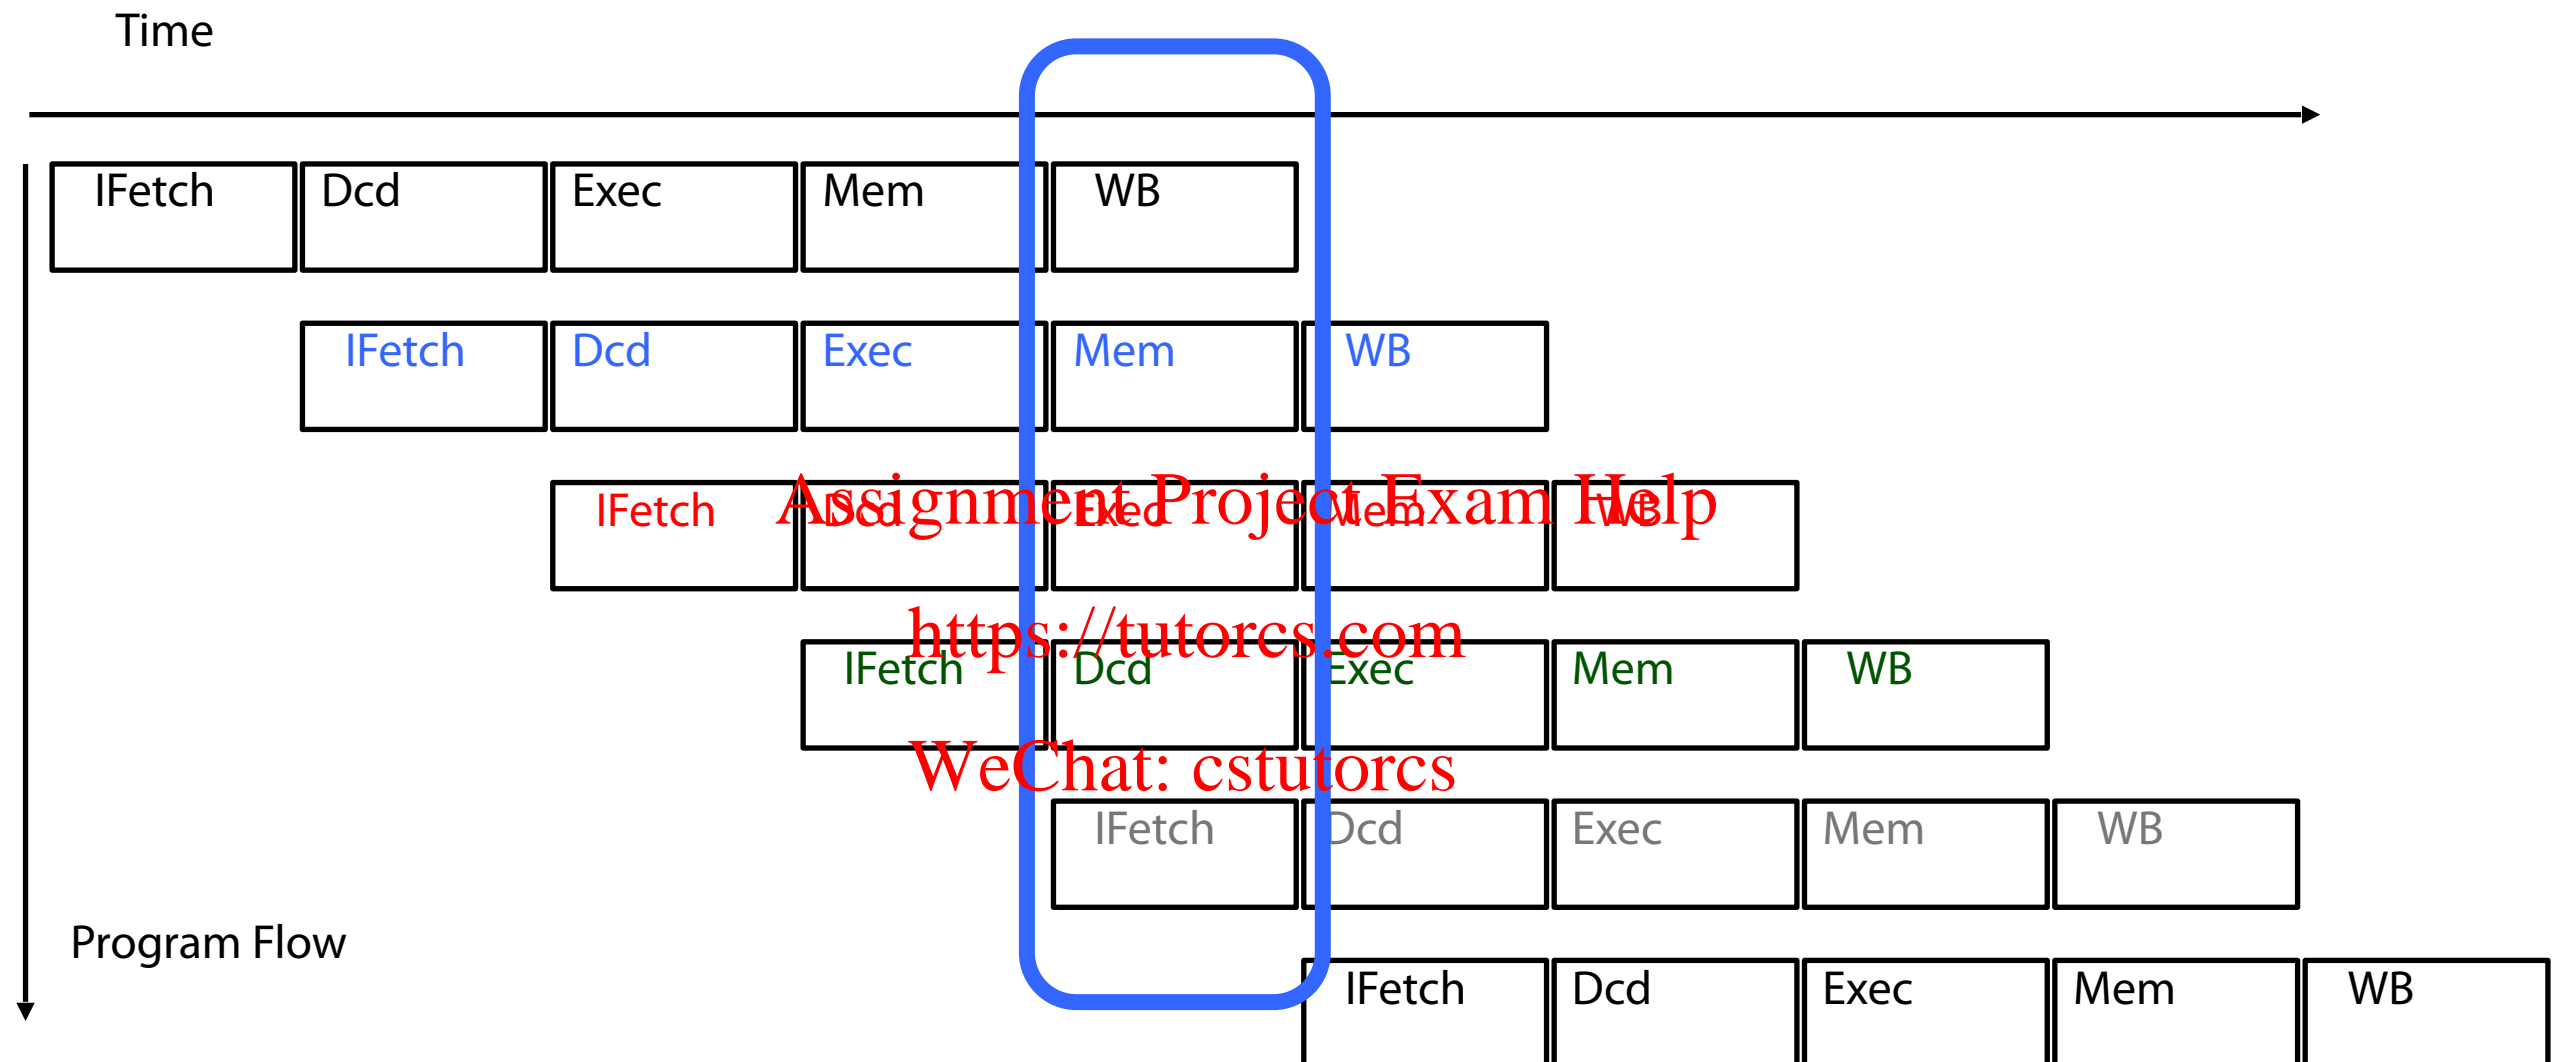
- Five stages, one step per stage
 1. IF: Instruction fetch from memory
 2. ID: Instruction decode & register read
 3. EX: Execute operation or calculate address
 4. MEM: Access memory operand
 5. WB: Write result back to register

Assignment Project Exam Help

<https://tutores.com>

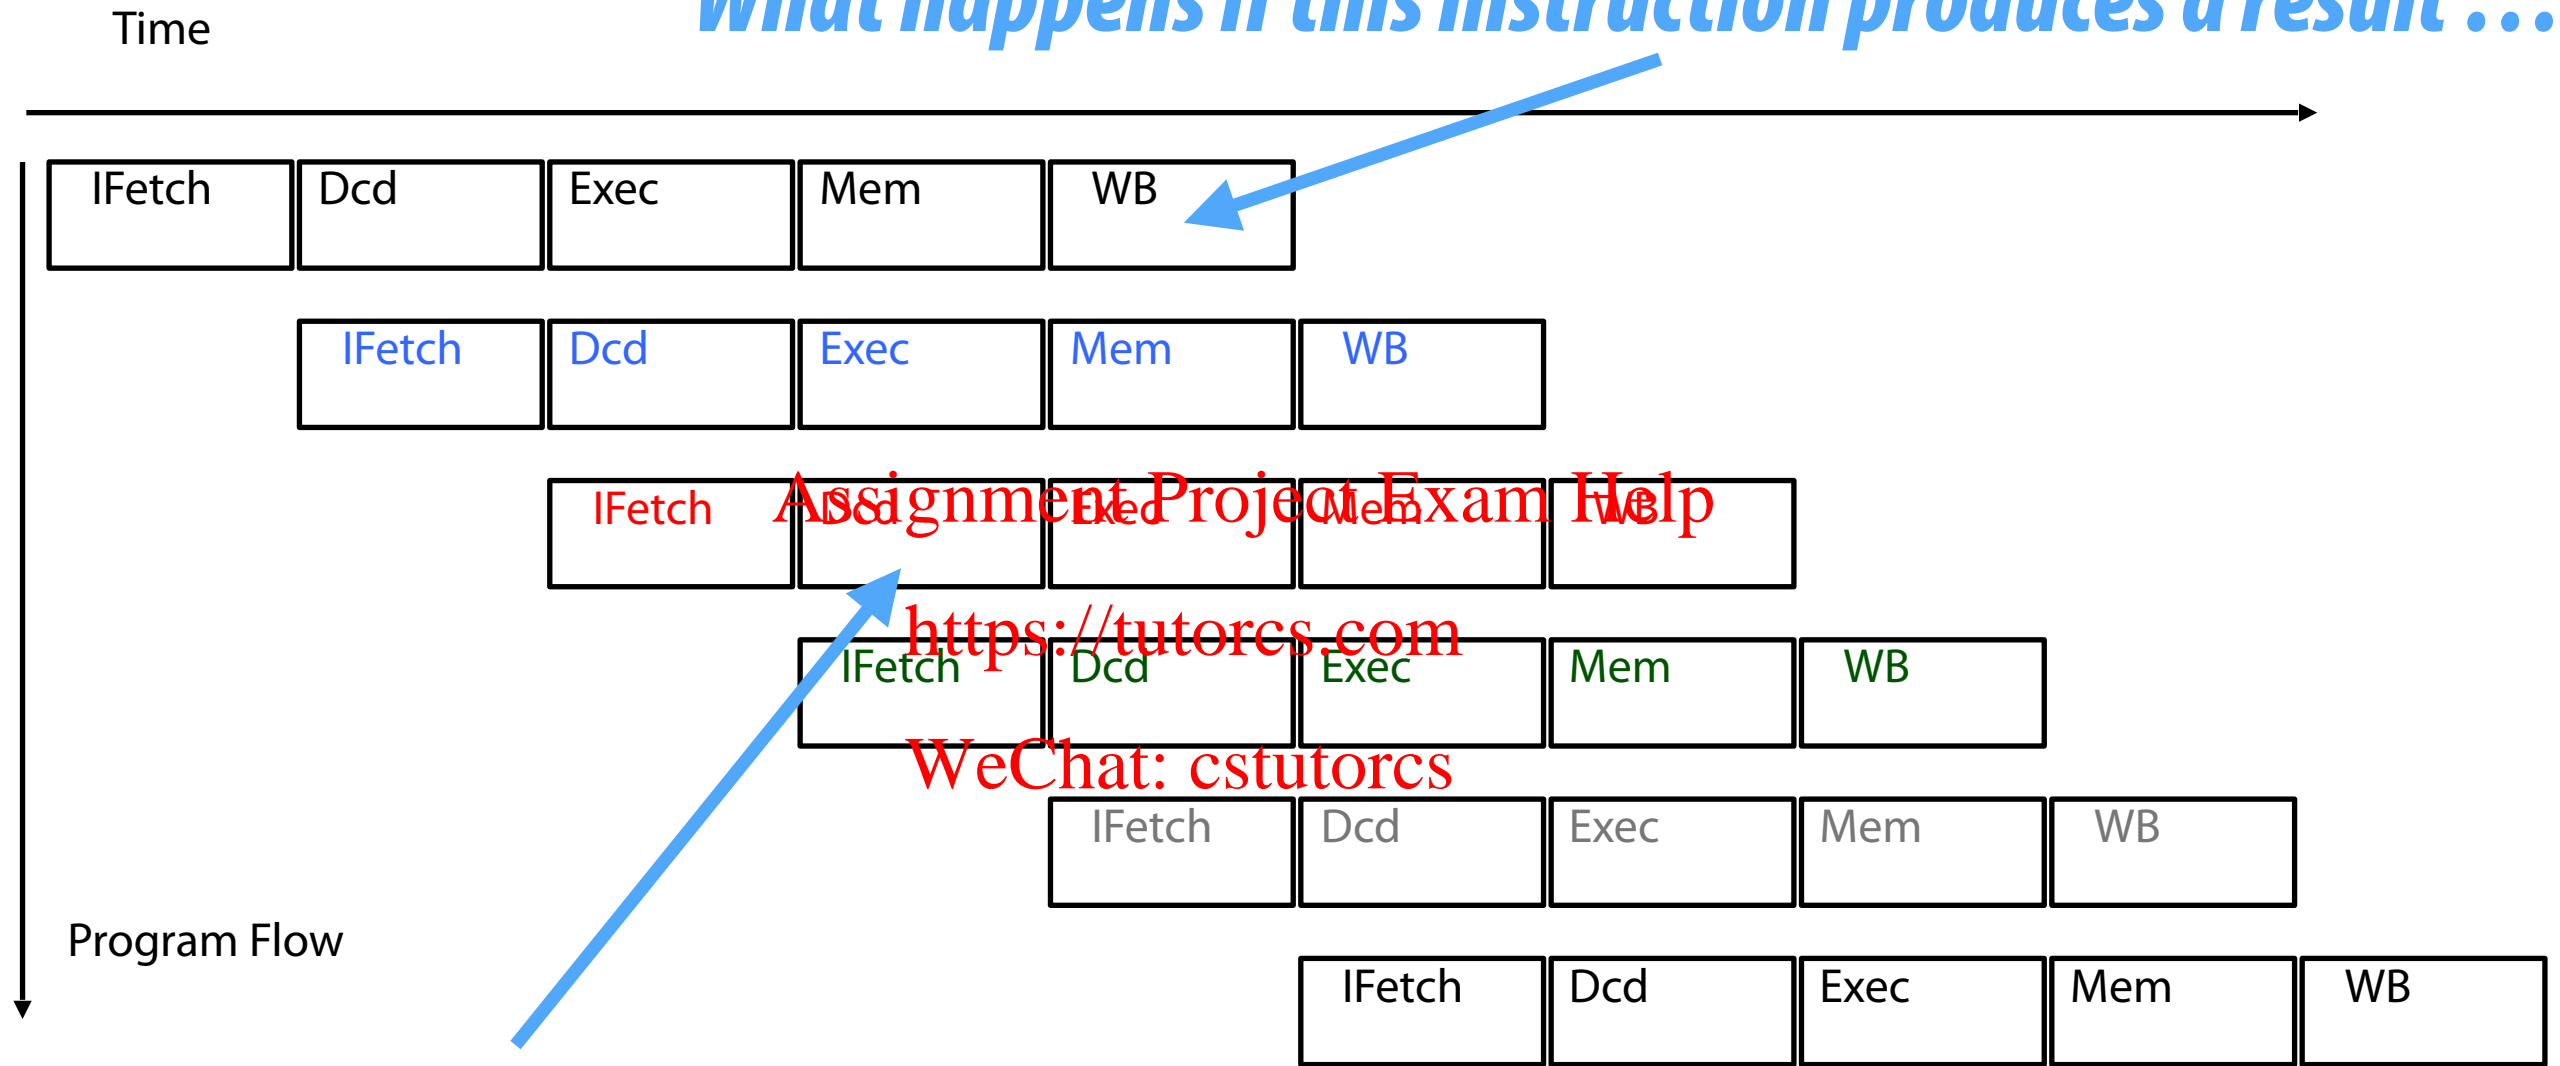
WeChat: cstutores

Conventional Pipelined Execution Representation



We Have No Time Machine

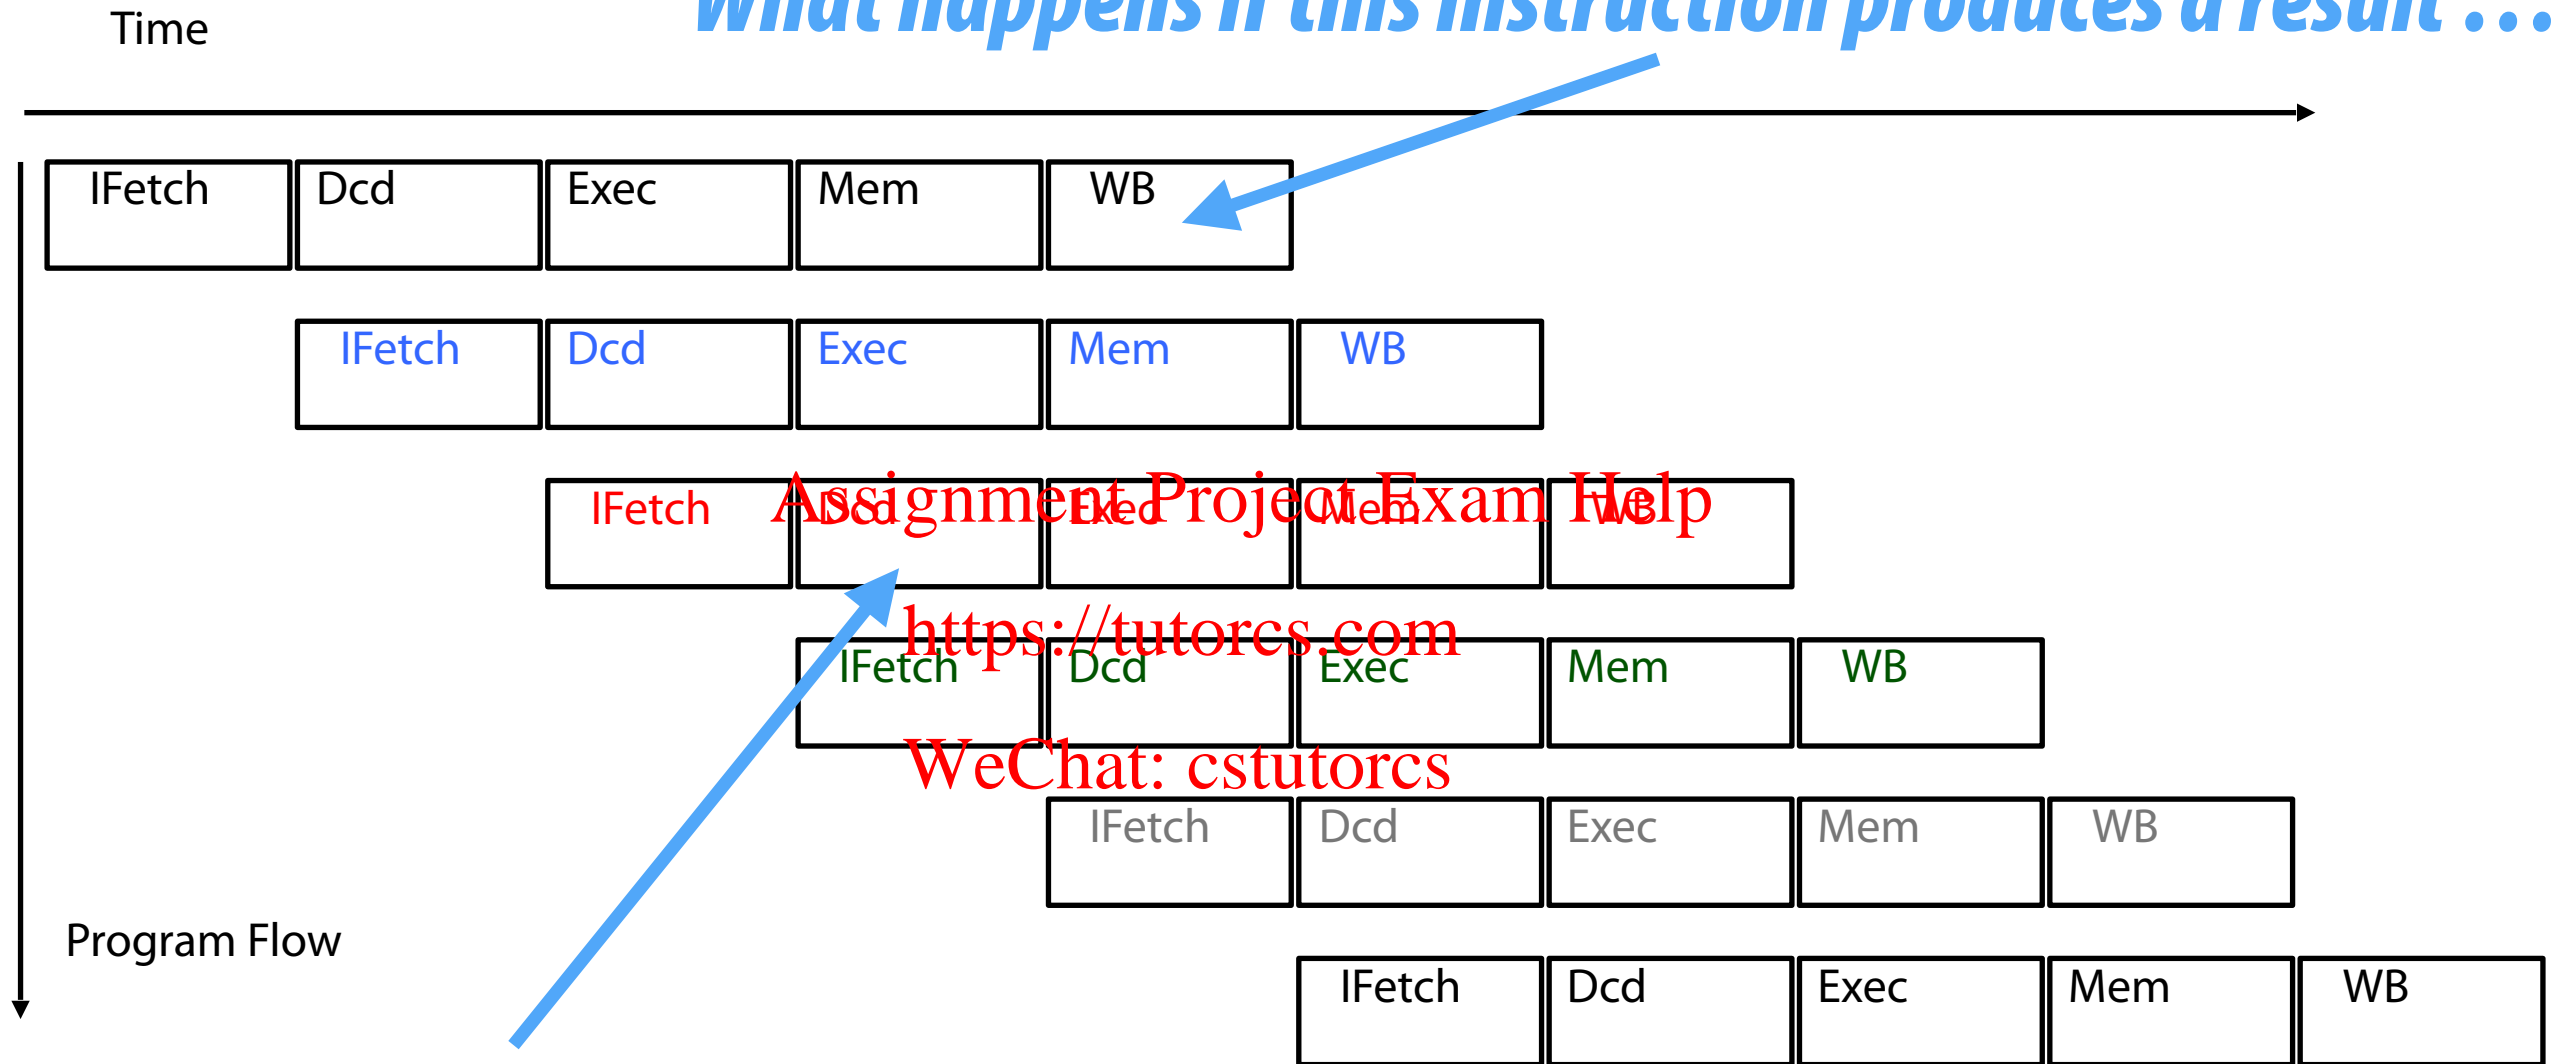
What happens if this instruction produces a result ...



... *that this instruction needs?*

We Have No Time Machine

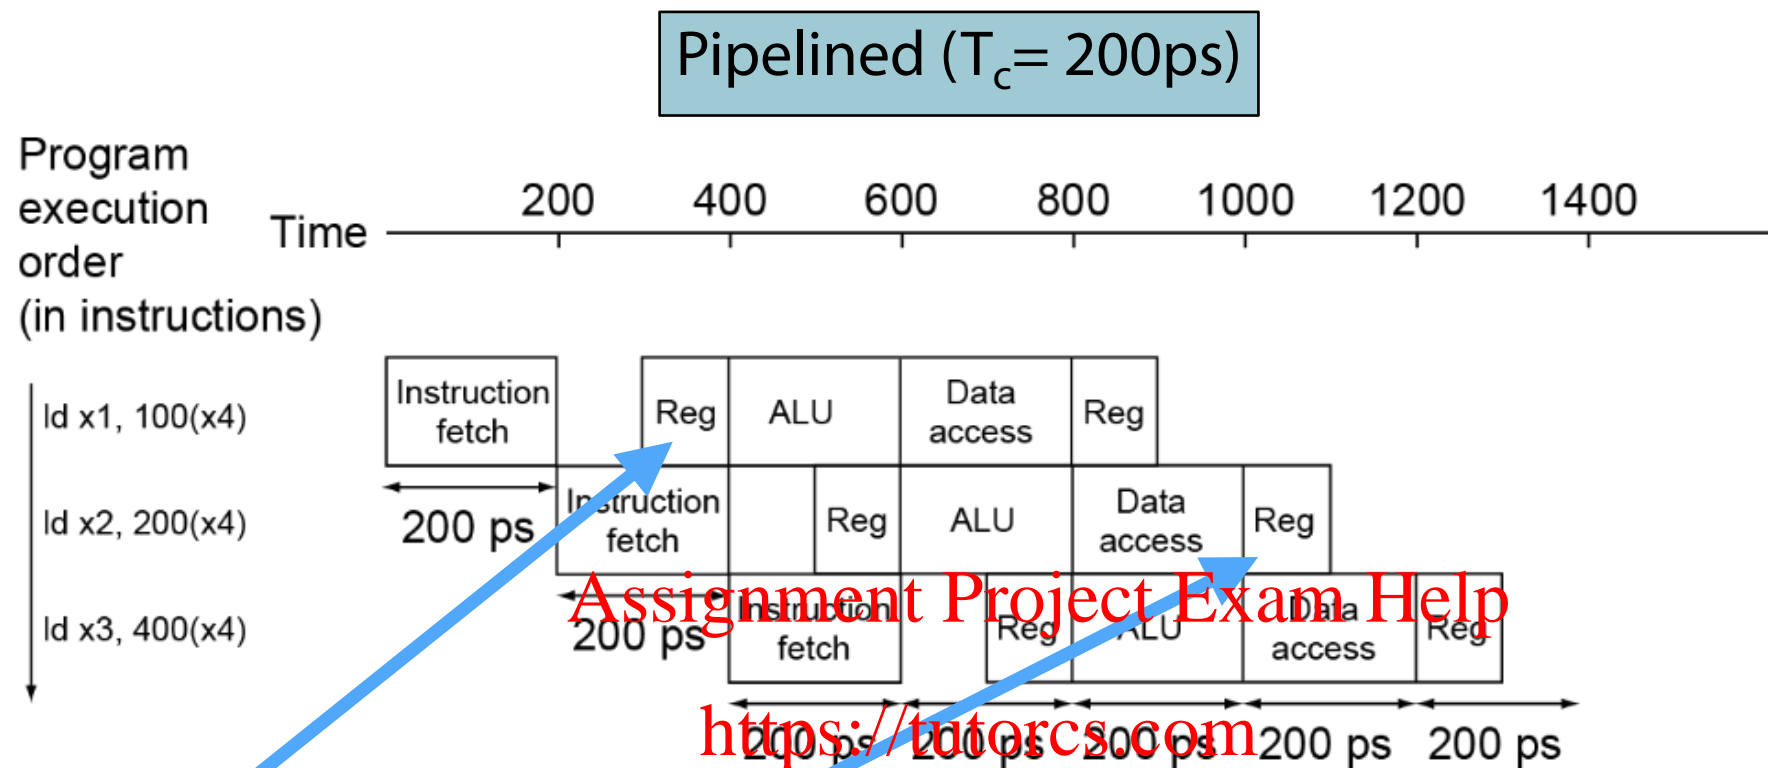
What happens if this instruction produces a result ...



... that this instruction needs?

- **For correctness, we can delay issuing the “red” instruction until we know we can get the result from the “black” instruction.**

RISC-V write/read order



- The “read” phase of the register file reads in the second half of the cycle
- The “write” phase of the register file writes in the first half of the cycle
- Thus an instruction can write a result then read the same result in the same clock cycle. (This is good!)

Hazards

- Situations that prevent starting the next instruction in the next cycle
- Structural hazards
 - A required resource is busy
- Data hazard
 - Need to wait for previous instruction to complete its data read/write
- Control hazard
 - Deciding on control action depends on previous instruction
- We can always resolve a hazard by waiting

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Structural Hazards

- Conflict for use of a resource
- Example: In RISC-V pipeline with a single memory
 - Load/store requires data access
 - Instruction fetch would have to *stall* for that cycle
 - Would cause a pipeline “bubble”
- Hence, pipelined datapaths require separate instruction/data memories
 - Or separate instruction/data caches
- General strategy for addressing structural hazard: replicate

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Data Hazards

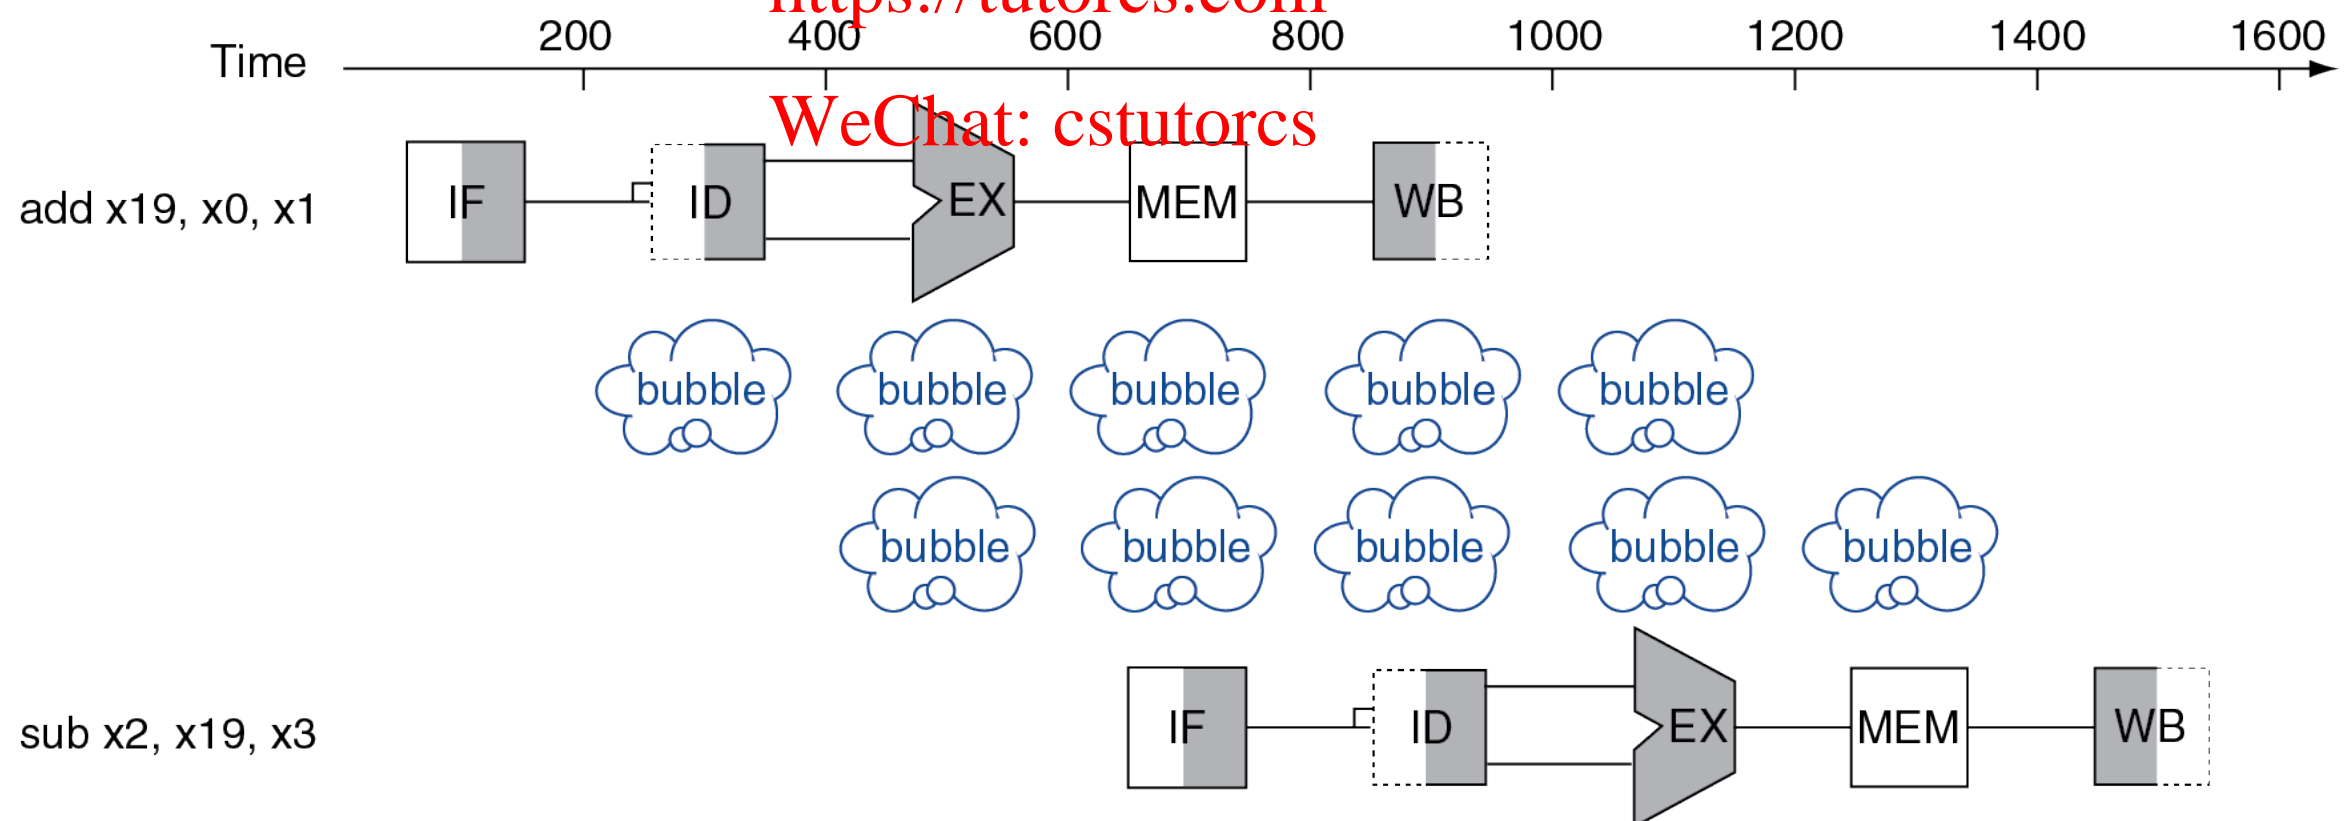
- An instruction depends on completion of data access by a previous instruction (“data dependency”)

— add x19, x0, x1
sub x2, x19, x3

Assignment Project Exam Help

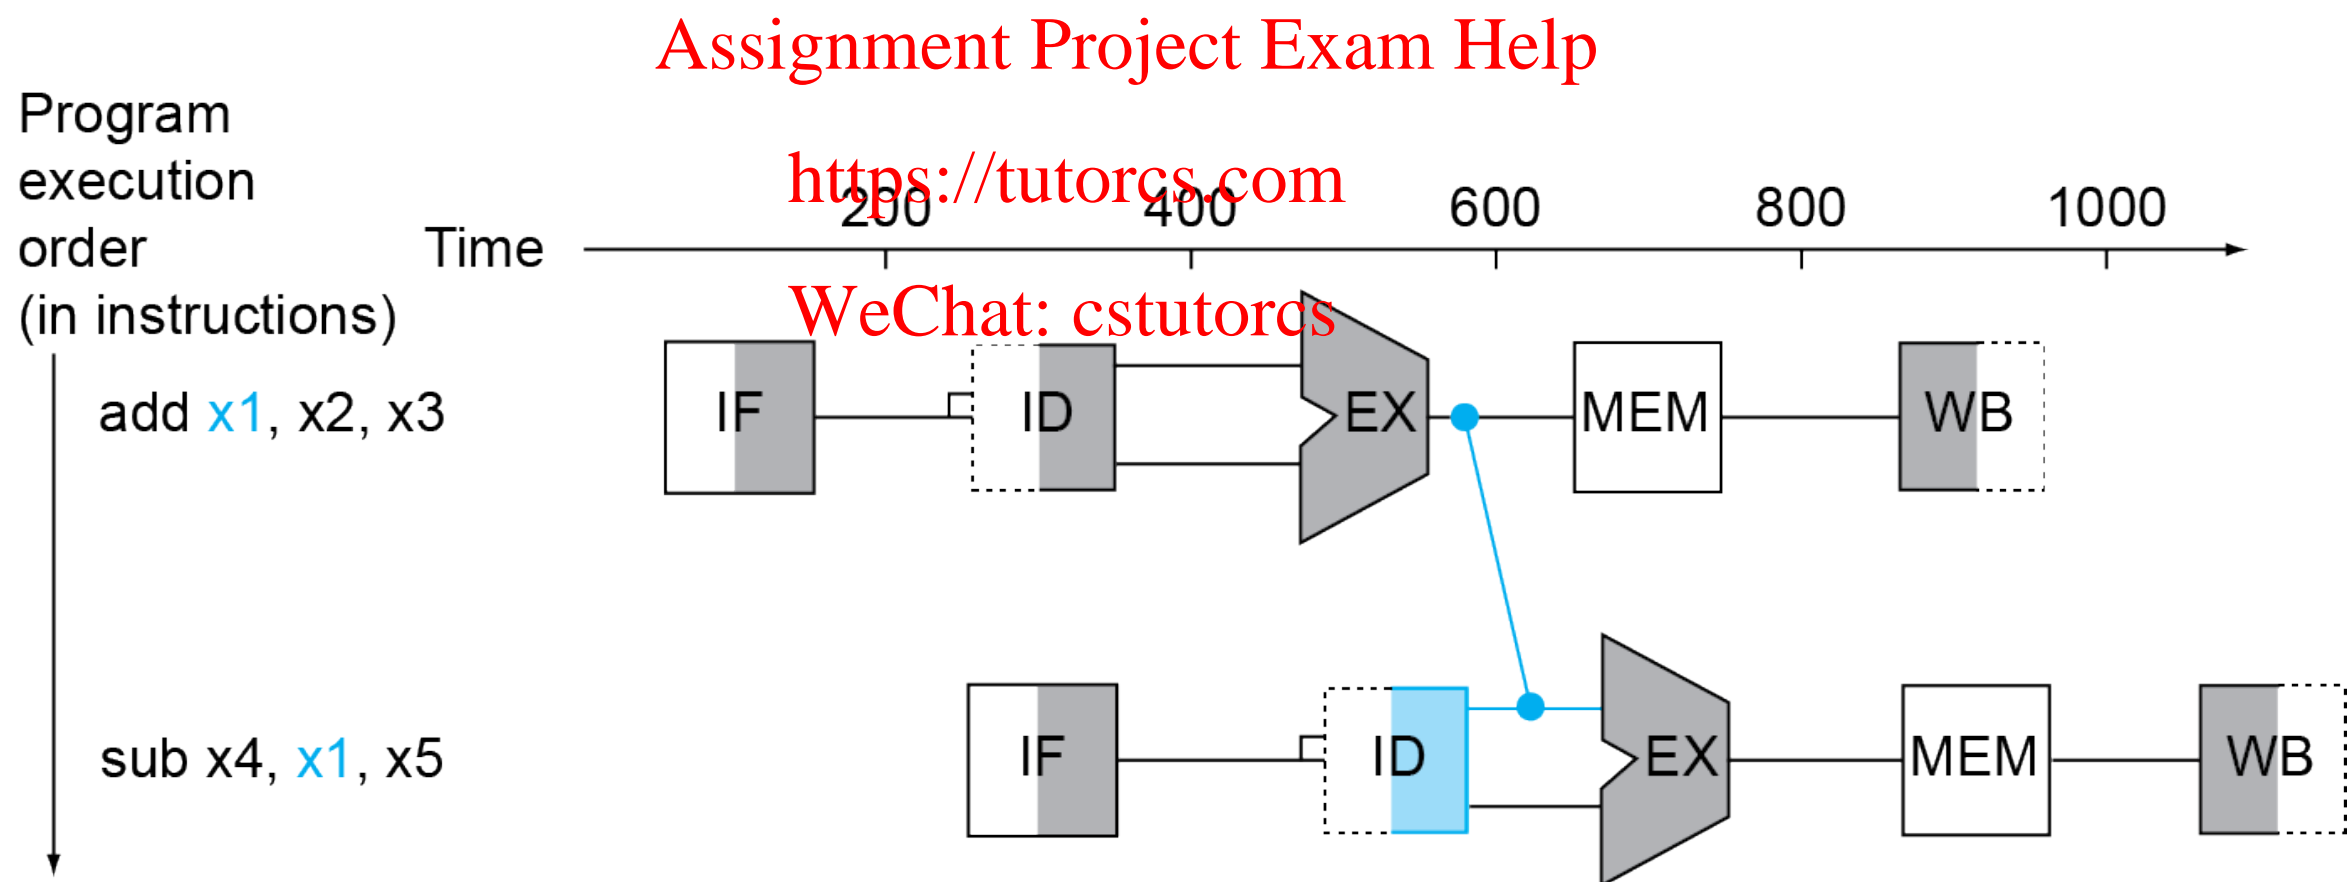
<https://tutorcs.com>

WeChat: cstutorcs



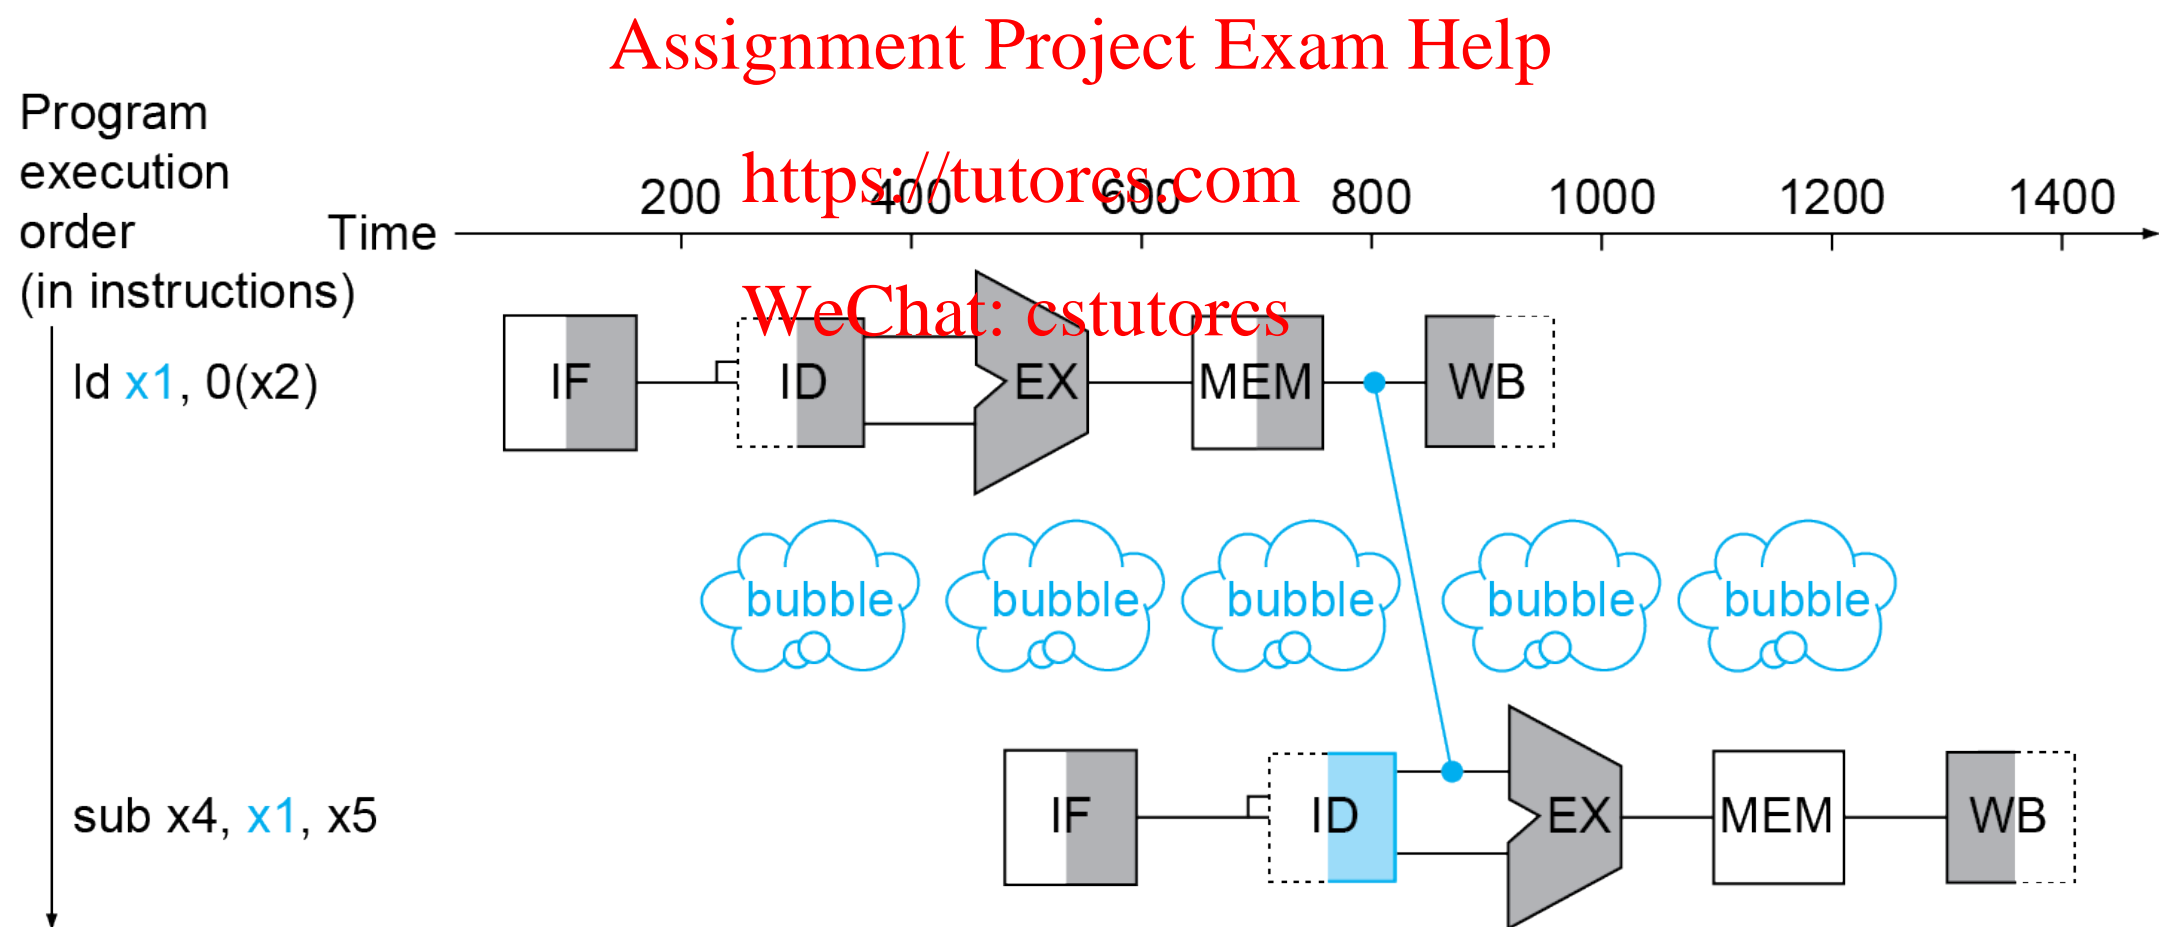
Forwarding (aka Bypassing)

- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath



Load-Use Data Hazard

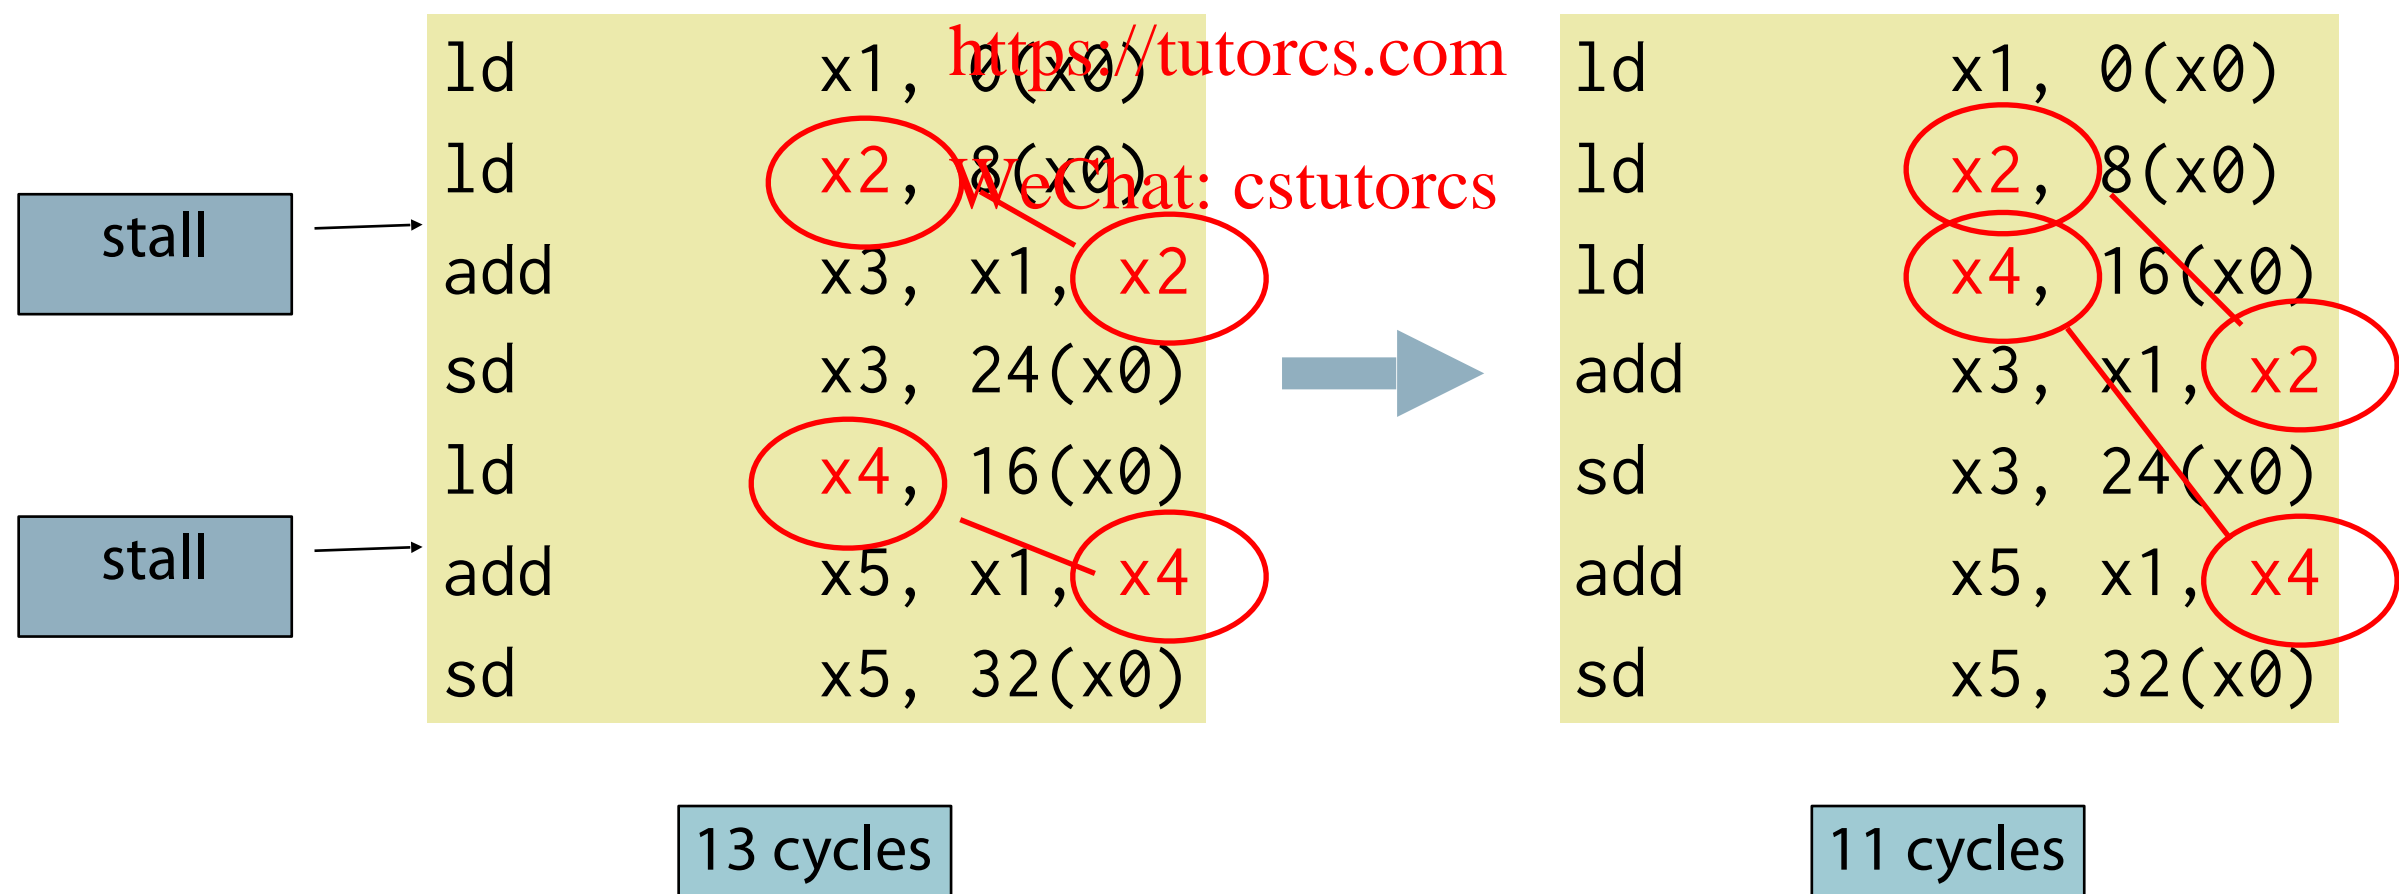
- Can't always avoid stalls by forwarding
 - If value not computed when needed
 - Can't forward backward in time!



Code Scheduling to Avoid Stalls

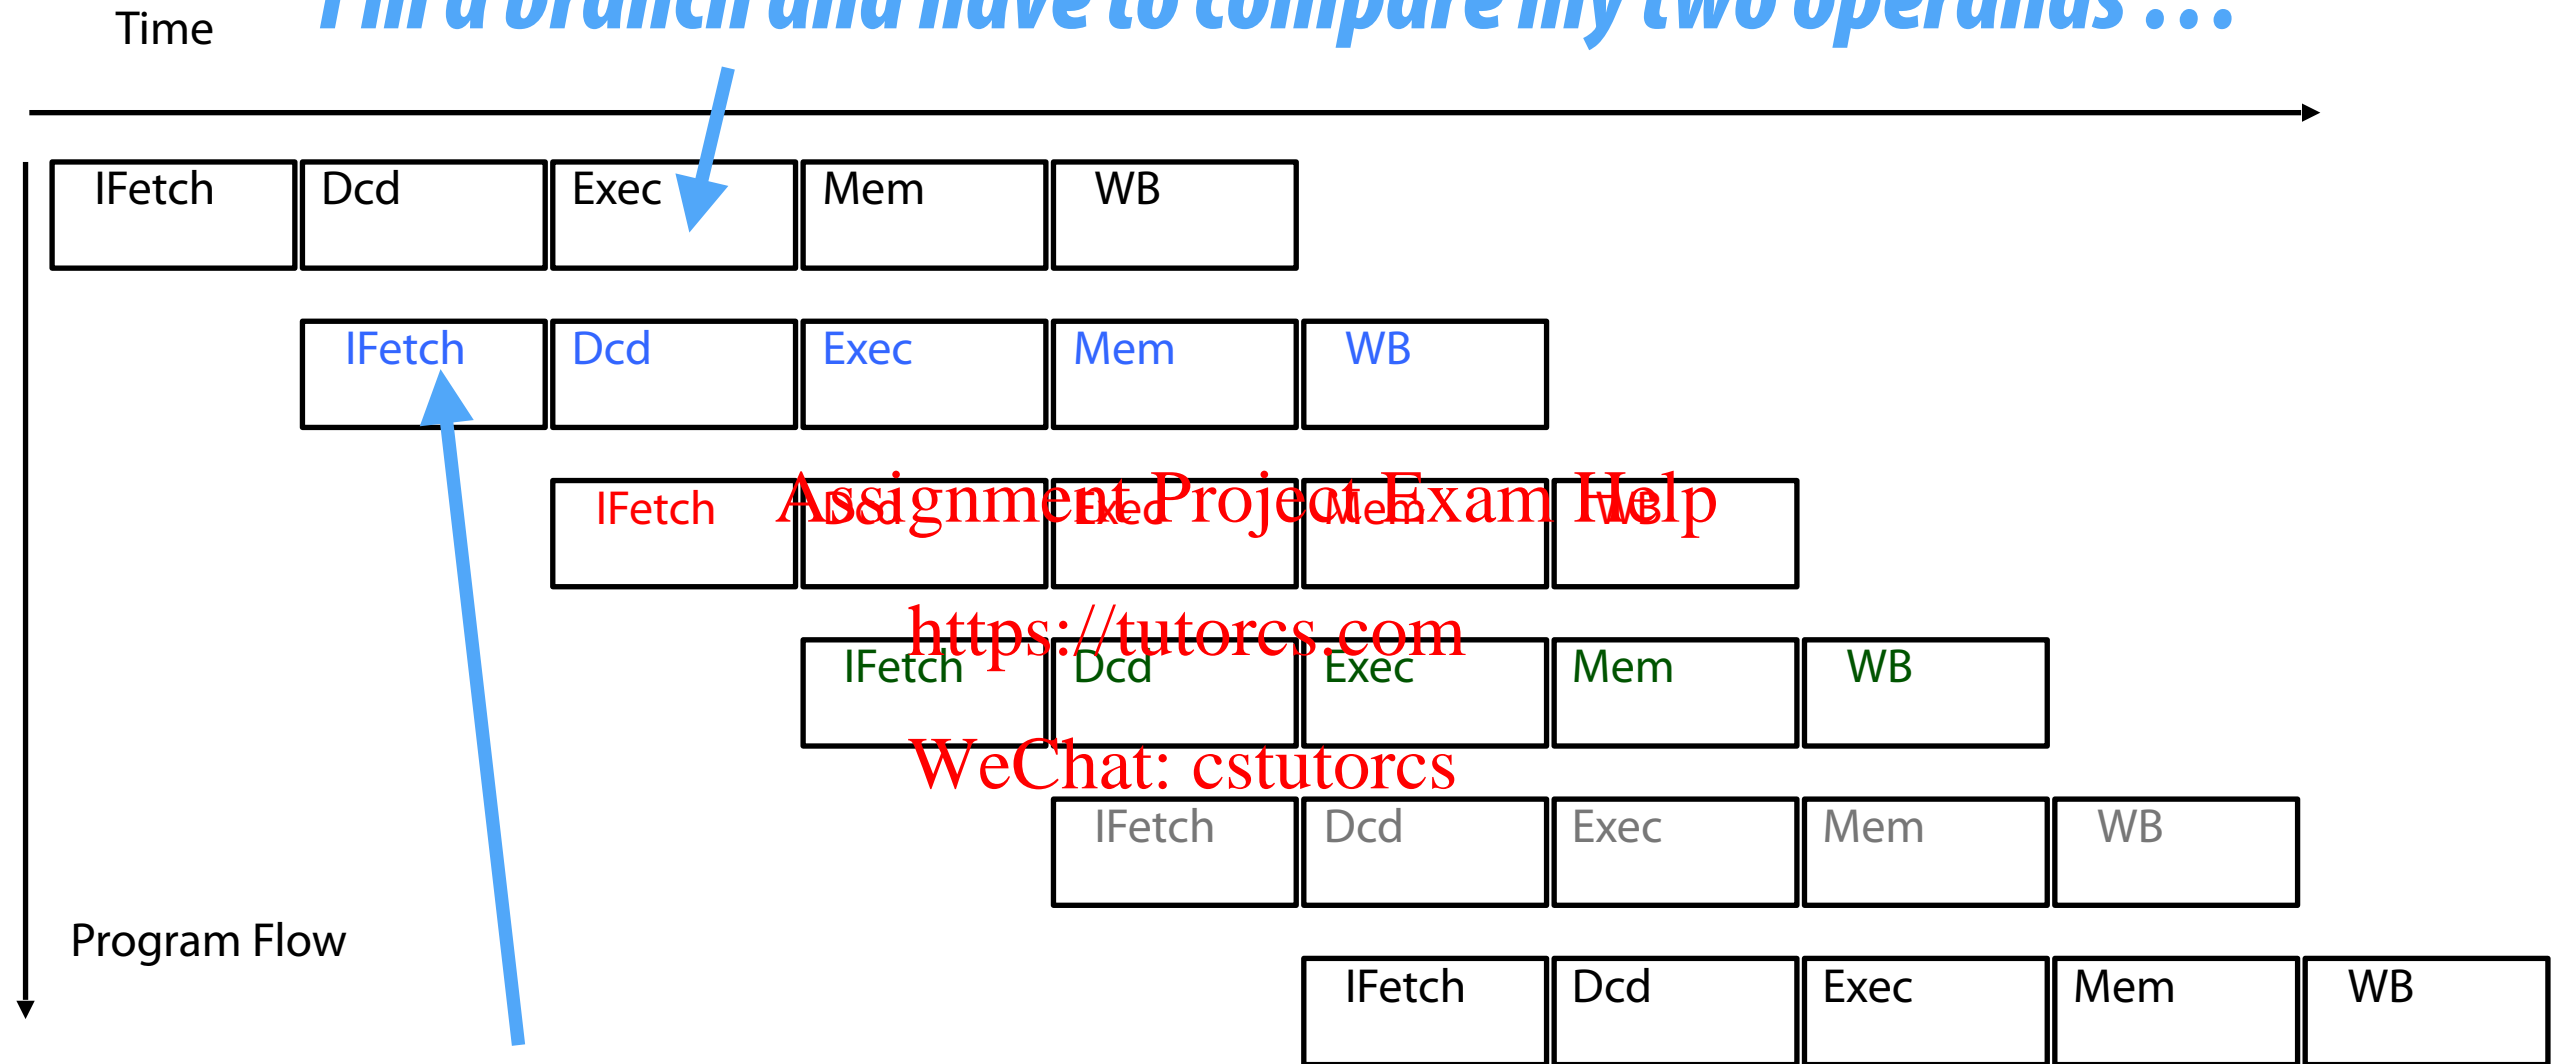
- Reorder code to avoid use of load result in the next instruction
- C code for $a = b + e; c = b + f;$

Assignment Project Exam Help



We Have No Time Machine

I'm a branch and have to compare my two operands ...



*... but I'm the next instruction;
what do I fetch?*

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Control Hazards

- **Branch determines flow of control**

- **Fetching next instruction depends on branch outcome**
- **Pipeline can't always fetch correct instruction**
 - **Still working on ID stage of branch**

Assignment Project Exam Help

<https://tutorcs.com>

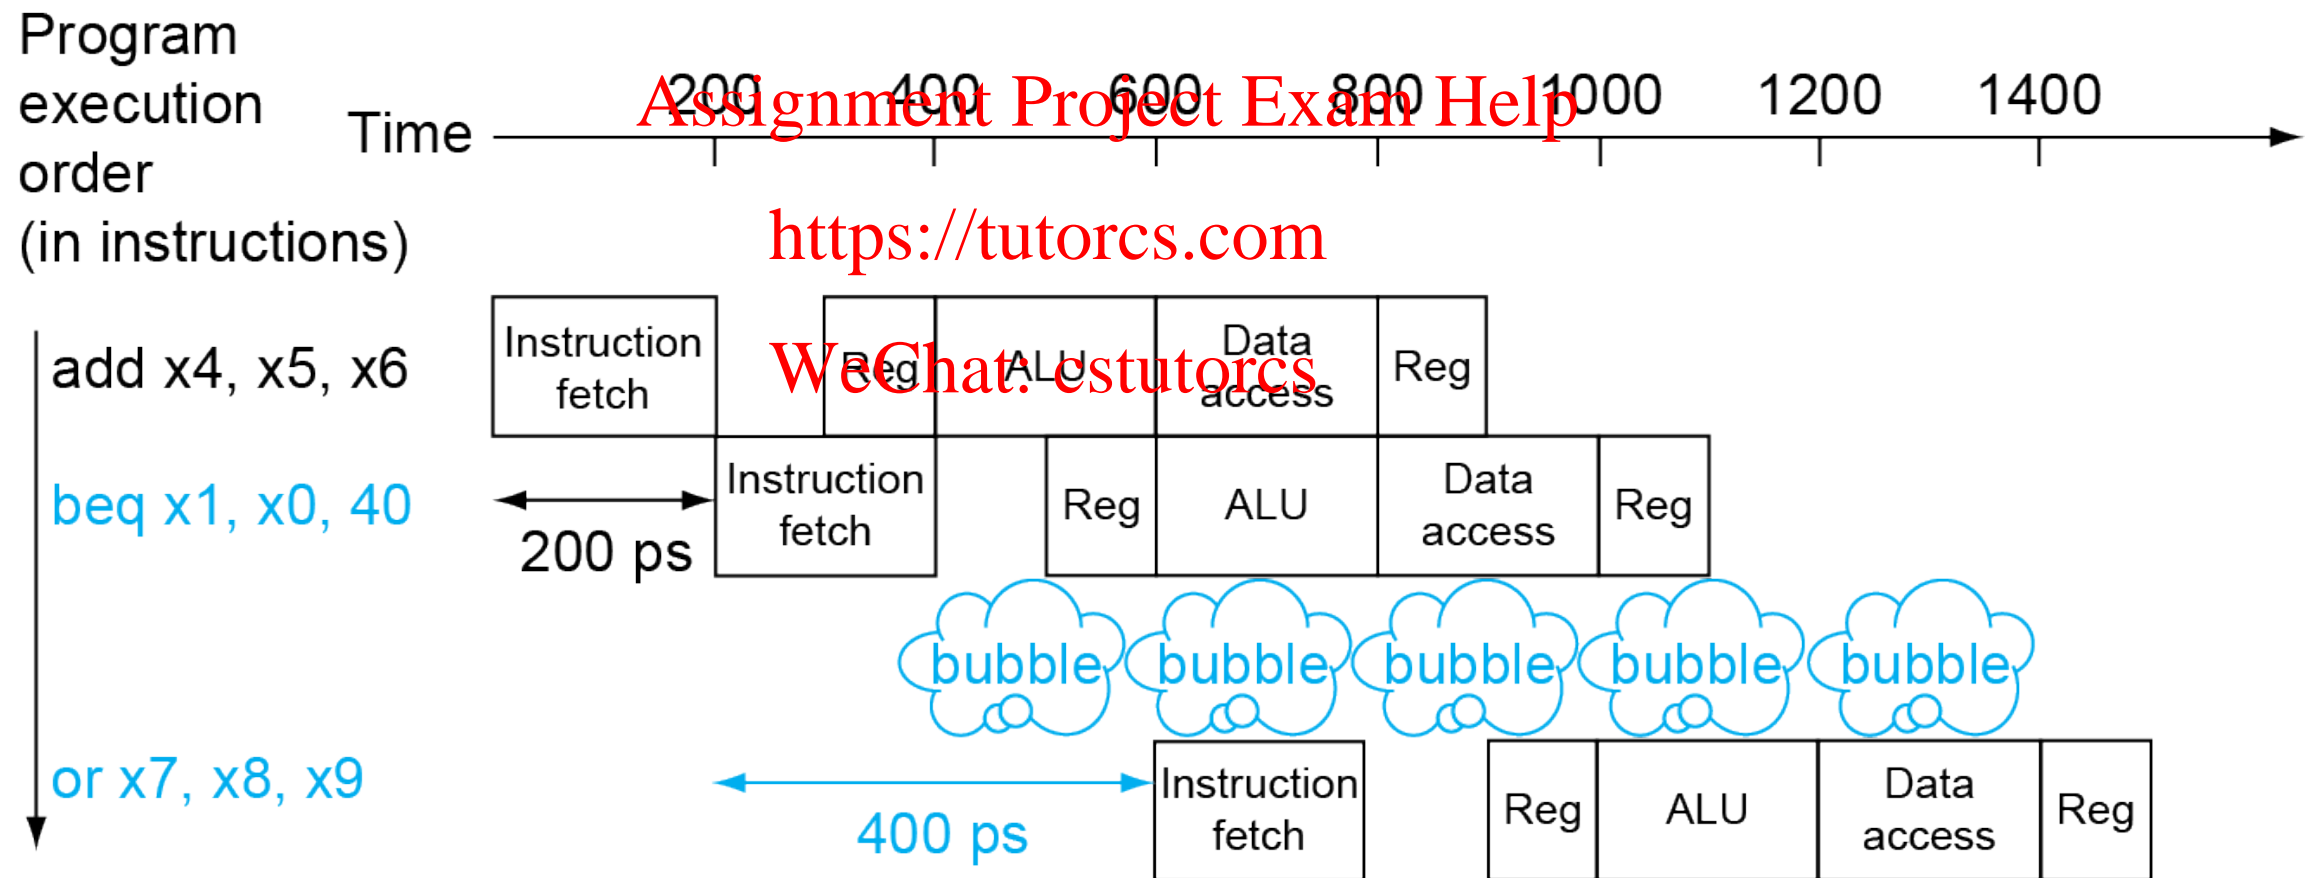
- **In RISC-V pipeline**

WeChat: cstutorcs

- **Need to compare registers and compute target early in the pipeline**
- **Add hardware to do it in ID stage**

Stall on Branch

- Wait until branch outcome determined before fetching next instruction



Branch Prediction

- There's a lot of work to evaluate a branch!
- High-performance pipelines are usually > 5 cycles
- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- *Predict outcome of branch*
 - Only stall if prediction is wrong
- In RISC-V pipeline
 - Can predict branches not taken
 - Fetch instruction after branch, with no delay

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cs_tutorcs

More-Realistic Branch Prediction

■ Static branch prediction

- Based on typical branch behavior
- Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken

Assignment Project Exam Help

<https://tutorcs.com>

■ Dynamic branch prediction

WeChat: cstutorcs

- Hardware measures actual branch behavior
 - e.g., record recent history of each branch
- Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and update history

■ If we have time later in the course, we'll spend more time on this

Pipeline Summary

The Big Picture

- improves performance by increasing instruction throughput
 - Executes multiple instructions in parallel
 - Each instruction has the same latency
- Subject to hazards
 - Structure, data, control
- Instruction set design affects complexity of pipeline implementation

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

