

Lecture 18a:

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Virtualization, Memory}

Introduction to Computer Architecture
UC Davis EEC 170, Fall 2019

Virtual Machines

- **Host computer emulates guest operating system and machine resources**
 - Improved isolation of multiple guests
 - Avoids security and reliability problems
 - Aids sharing of resources
- **Virtualization has some performance impact**
 - Feasible with modern high-performance computers
- **Examples**
 - IBM VM/370 (1970s technology!)
 - VMWare
 - Microsoft Virtual PC

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Virtual Machine Monitor

- **Maps virtual resources to physical resources**
 - **Memory, I/O devices, CPUs**
- **Guest code runs on native machine in user mode**
 - **Traps to VMM on privileged instructions and access to protected resources**
- **Guest OS may be different from host OS**
- **VMM handles real I/O devices**
 - **Emulates generic virtual I/O devices for guest**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Example: Timer Virtualization

- In native machine, on timer interrupt
 - OS suspends current process, handles interrupt, selects and resumes next process
- With Virtual Machine Monitor
 - VMM suspends current VM, handles interrupt, selects and resumes next VM <https://tutorcs.com>
- If a VM requires timer interrupts
 - VMM emulates a virtual timer
 - Emulates interrupt for VM when physical timer interrupt occurs
- Guest VM doesn't get any access to the raw machine. You can't trust it.

Instruction Set Support

- **User and System modes**
- **Privileged instructions only available in system mode**
 - **Trap to system if executed in user mode**
- **All physical resources only accessible using privileged instructions**
 - **Including page tables, interrupt controls, I/O registers**
- **Renaissance of virtualization support**
 - **Current ISAs (e.g., x86) adapting**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Virtualization & Instruction Sets

- **Goal of classical virtualization: Guest OS runs in user mode; any privileged instruction run by the guest OS is trapped and handled by the hypervisor**
- **“ARMv7 is not classically virtualizable because, among other reasons, the return-from-exception instruction, RFE, is not defined to trap when executed in user mode.”**
- **“The [x86] ISA is not classically virtualizable, since some privileged instructions silently fail in user mode rather than trapping. VMware’s engineers famously worked around this deficiency with intricate dynamic binary translation software.”**
 - **“Indeed, engineers from Intel Corporation were convinced their processors could not be virtualized in any practical sense. ... Unfortunately, the description of the x86 architecture, publicly available as the Intel Architecture Manual [Intel Corporation 2010], was at once baroquely detailed and woefully imprecise for our purpose. For example, the formal specification of a single instruction could easily exceed 8 pages of pseudocode while omitting crucial details necessary for correct virtualization.”**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Virtual Memory

- Use main memory as a “cache” for secondary (disk) storage
 - Managed jointly by CPU hardware and the operating system (OS)
- Programs share main memory
 - Each gets a private virtual address space holding its frequently used code and data
 - Protected from other programs
- CPU and OS translate virtual addresses to physical addresses
 - VM “block” is called a page
 - VM translation “miss” is called a page fault

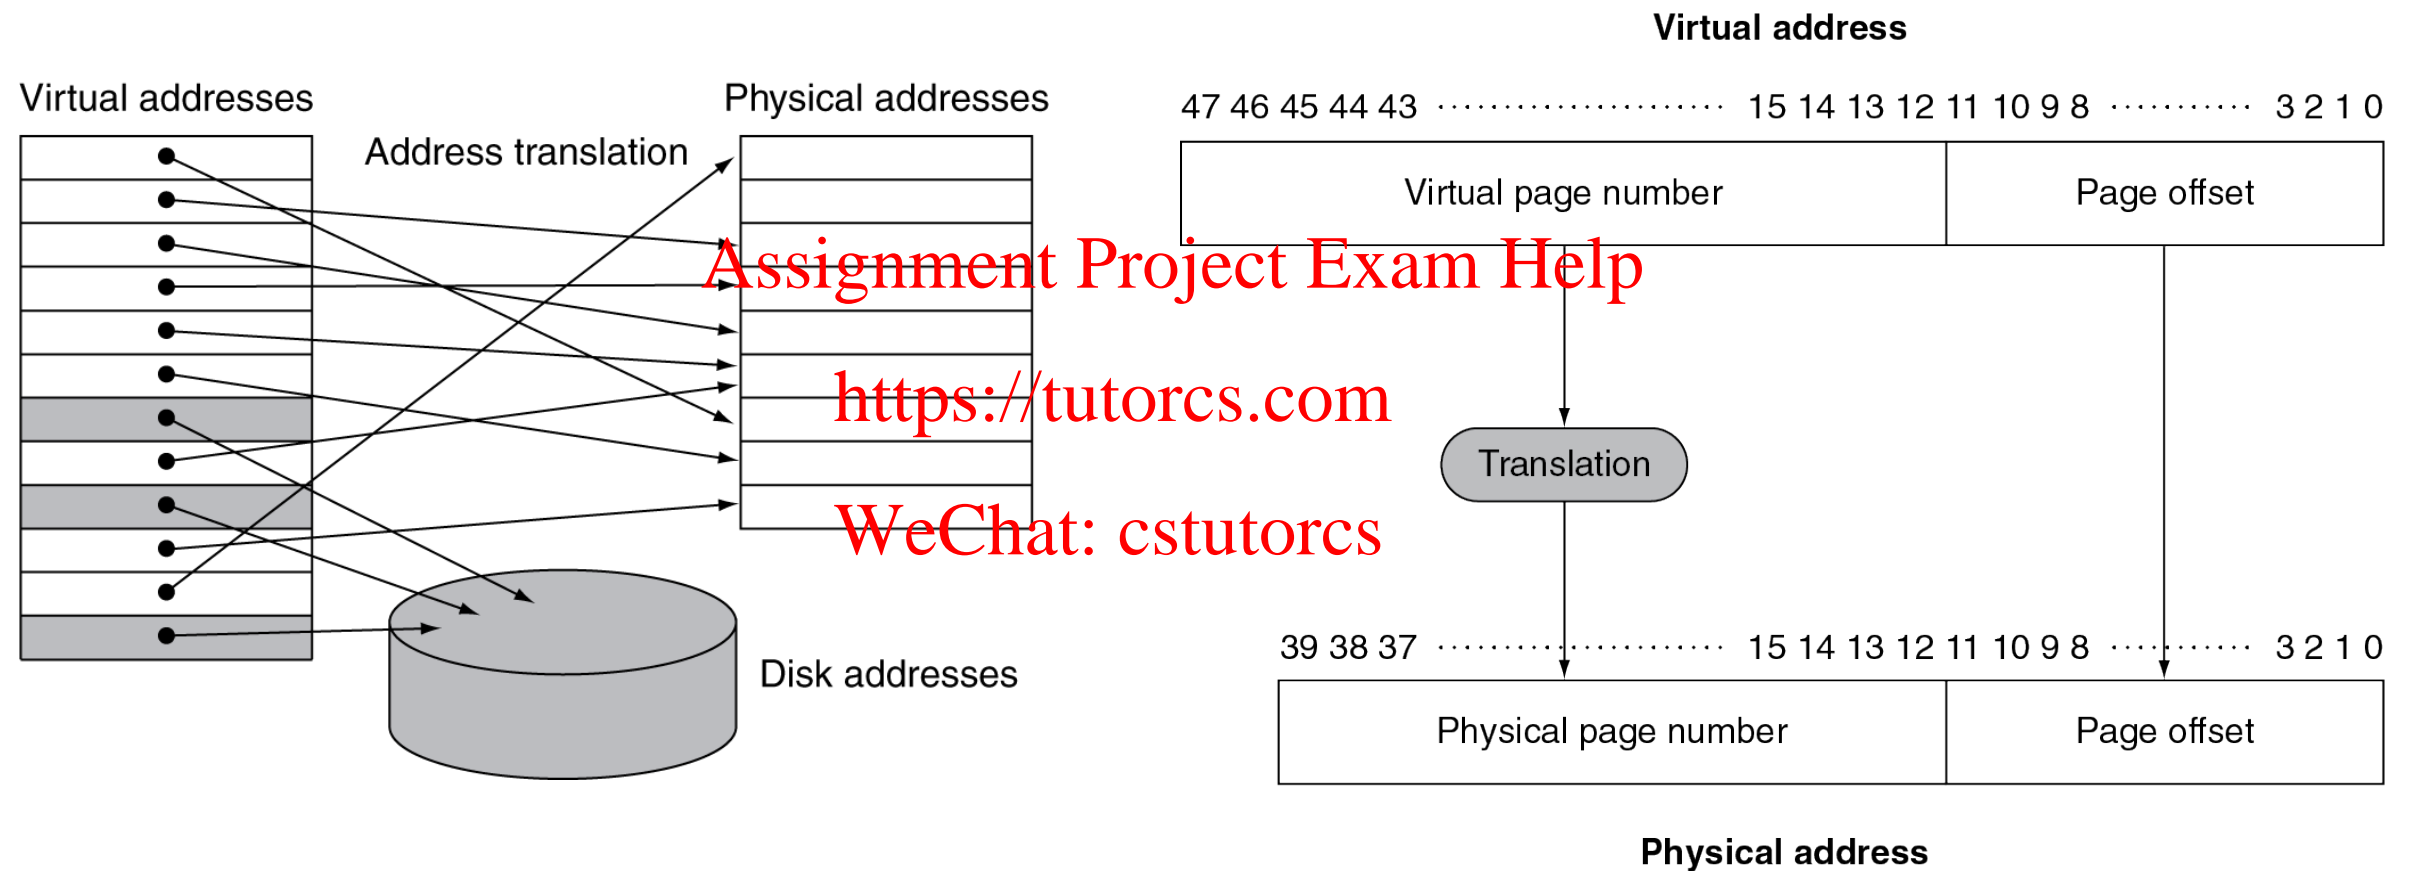
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Address Translation

■ Fixed-size pages (e.g., 4K)



Page Fault Penalty

- On page fault, the page must be fetched from disk
 - Takes millions of clock cycles
 - Handled by OS code
- Try to minimize page fault rate
 - Fully associative placement
 - Smart replacement algorithms

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Page Tables

- **Stores placement information**

- **Array of page table entries, indexed by virtual page number**
- **Page table register in CPU points to page table in physical memory**

Assignment Project Exam Help

- **Protected register, not user-accessible**

<https://tutorcs.com>

- **If page is present in memory**

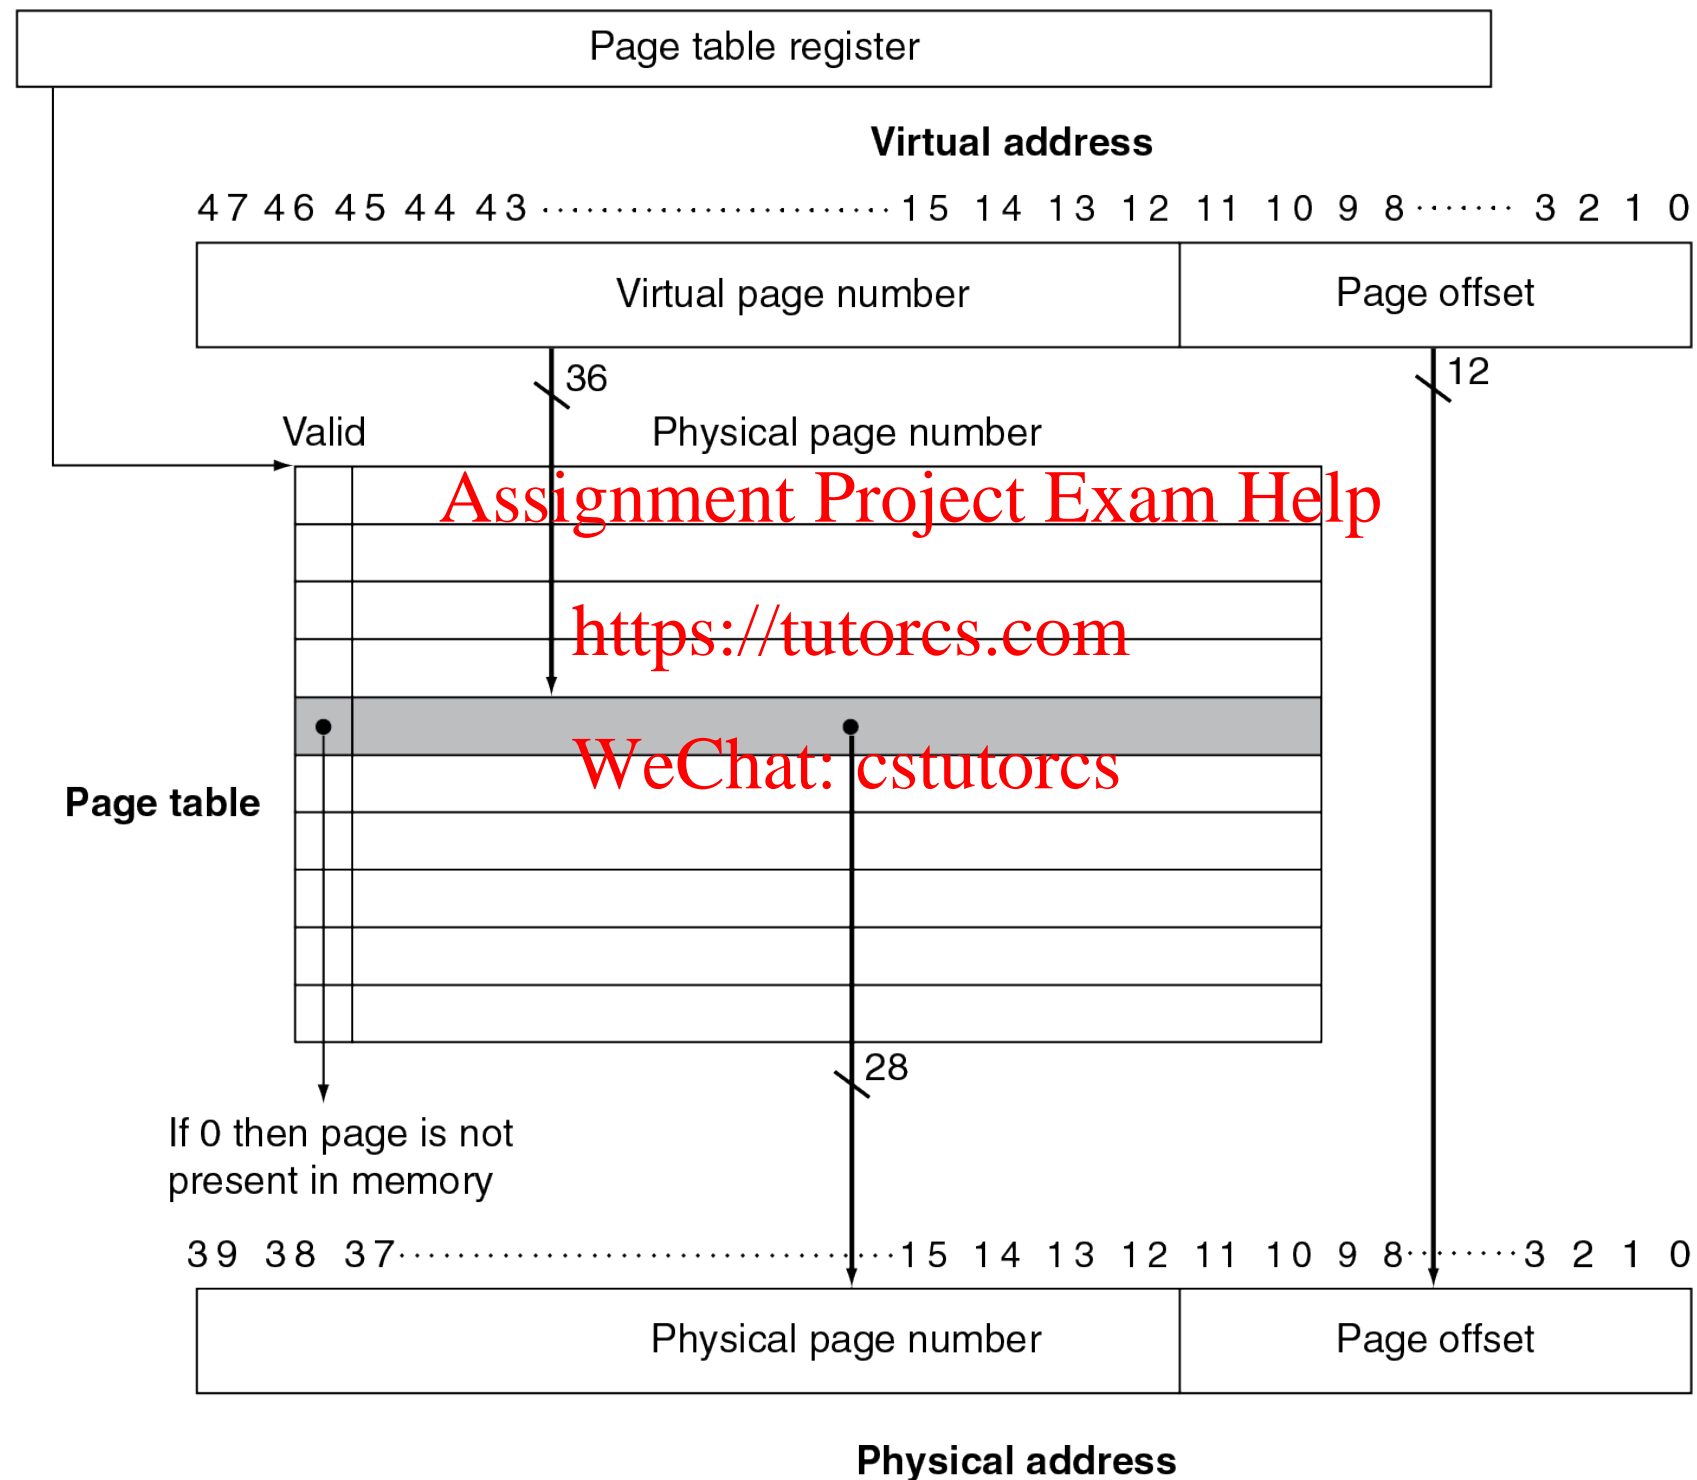
WeChat: cstutorcs

- **PTE stores the physical page number**
- **Plus other status bits (referenced, dirty, ...)**

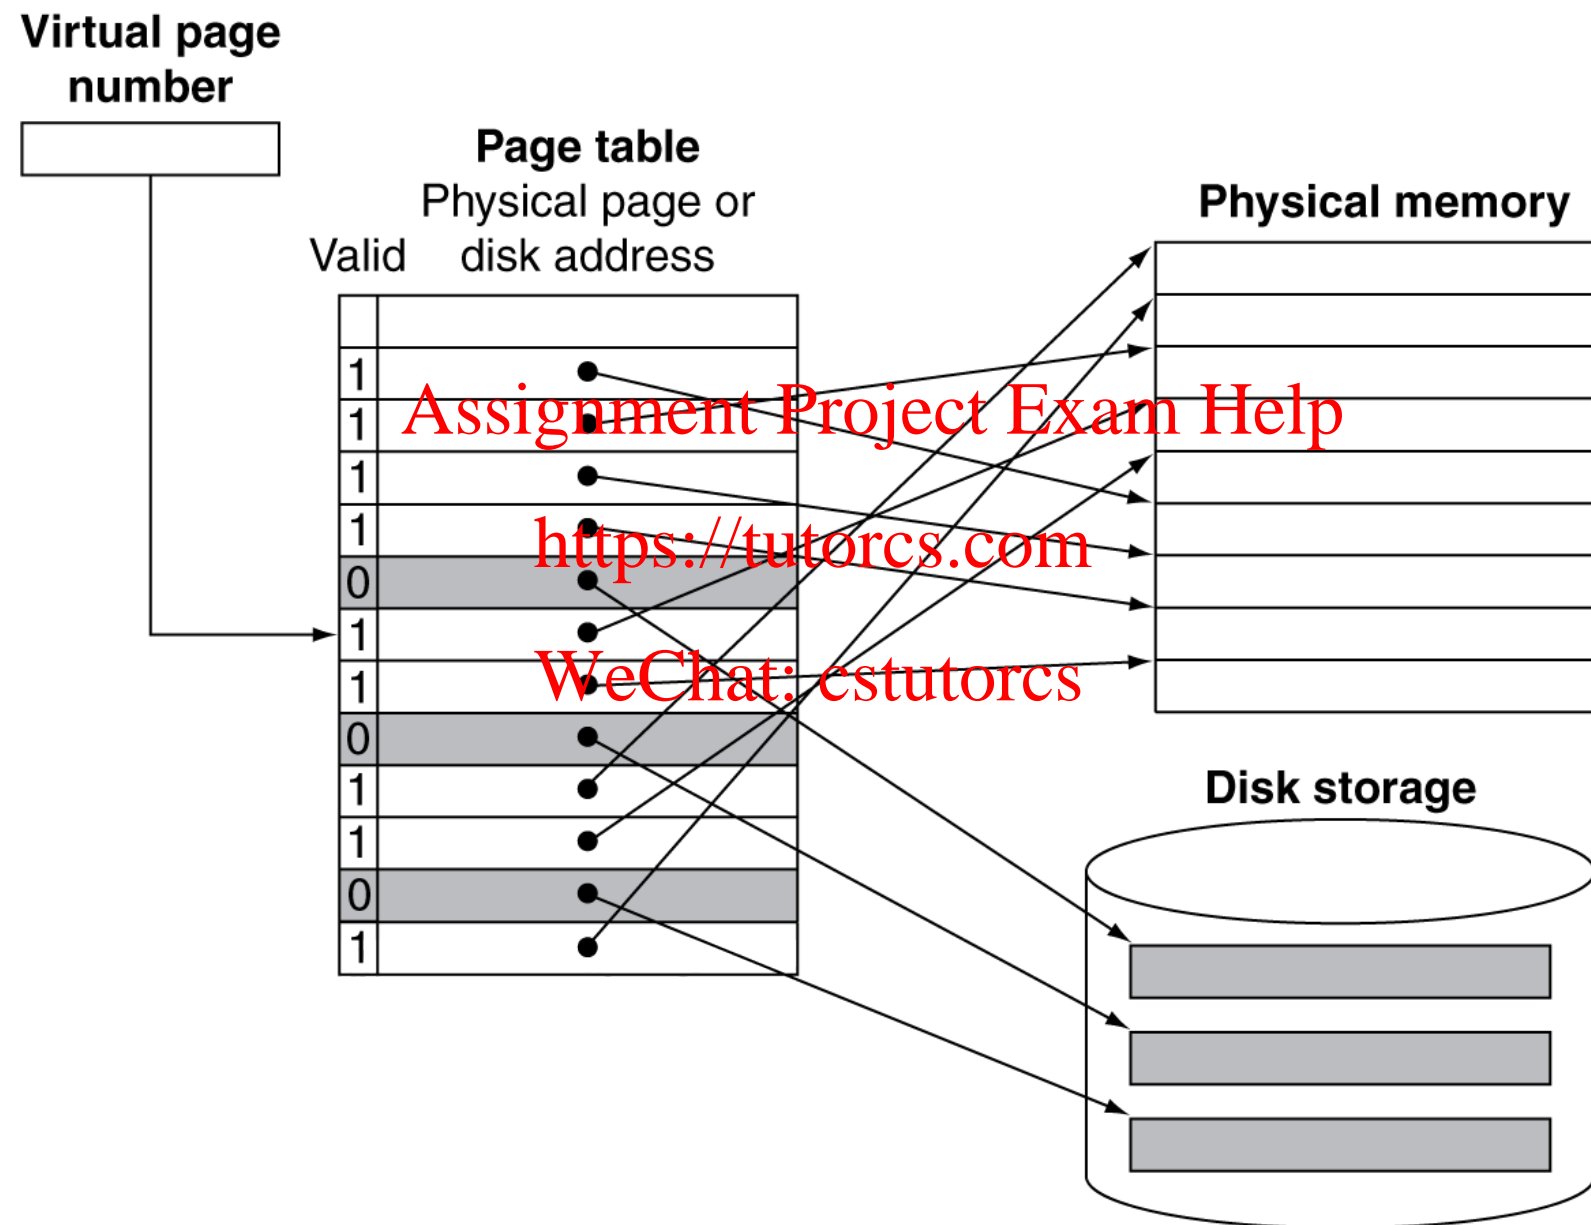
- **If page is not present**

- **PTE can refer to location in swap space on disk**

Translation Using a Page Table



Mapping Pages to Storage



Replacement and Writes

- To reduce page fault rate, prefer least-recently used (LRU) replacement
 - Reference bit (aka use bit) in PTE set to 1 on access to page
 - Periodically cleared to 0 by OS
 - A page with reference bit = 0 has not been used recently
- Disk writes take millions of cycles
 - Block at once, not individual locations
 - Write through is impractical
 - Use write-back
 - Dirty bit in PTE set when page is written

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Fast Translation Using a TLB

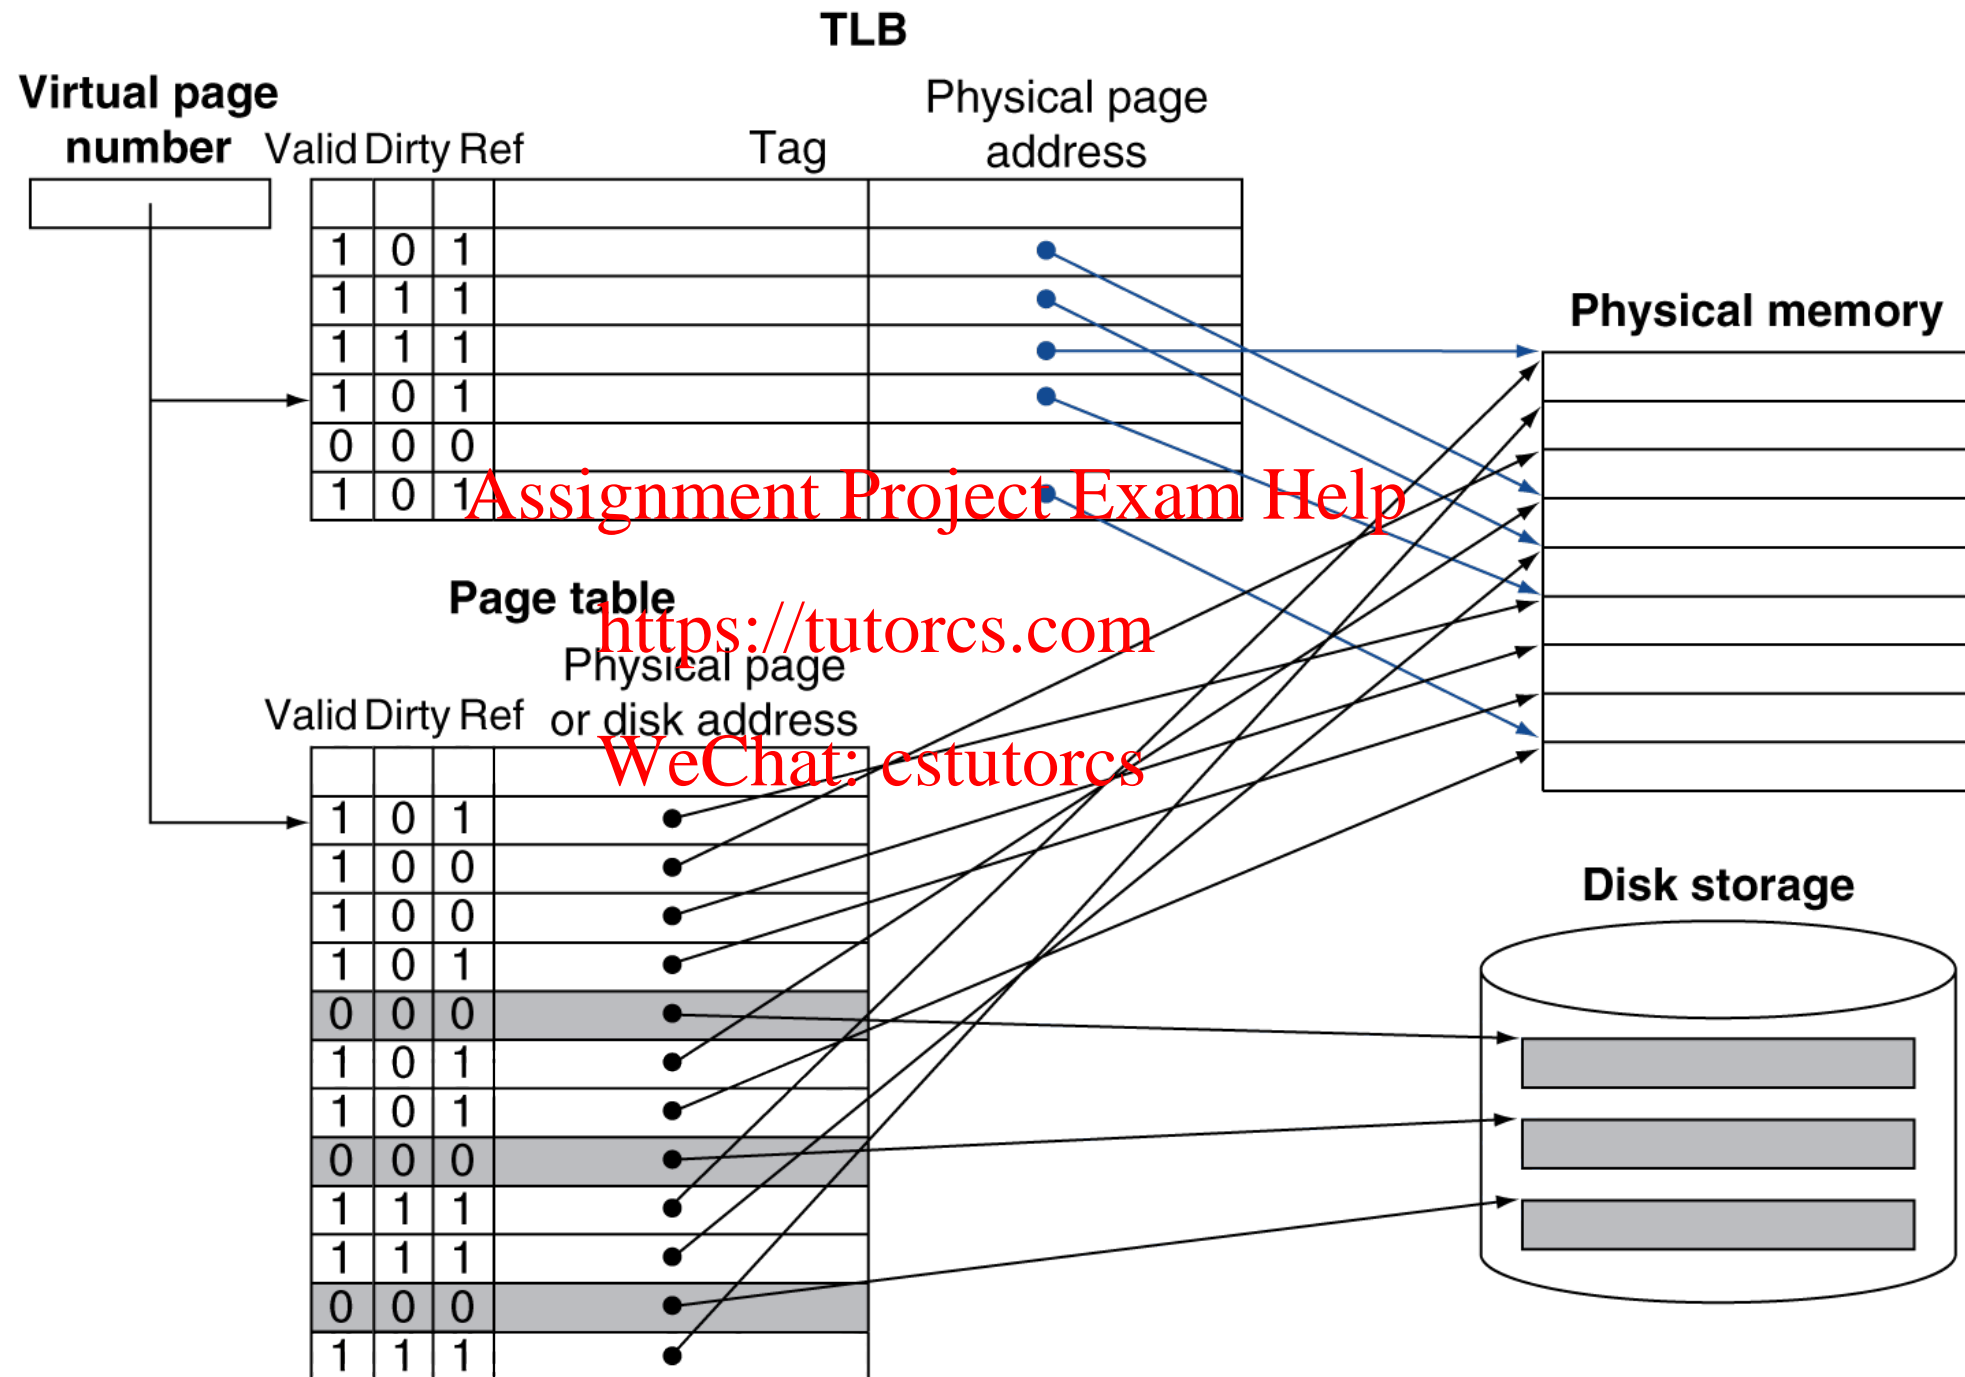
- **Address translation would appear to require extra memory references**
 - **One to access the PTE (virtual->physical translation)**
 - **Then the actual memory access**
- **But access to page tables has good locality**
 - **So use a fast cache of PTEs within the CPU**
 - **Called a Translation Look-aside Buffer (TLB)**
 - **Typical: 16–512 PTEs, 0.5–1 cycle for hit, 10–100 cycles for miss, 0.01%–1% miss rate**
 - **Misses could be handled by hardware or software**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Fast Translation Using a TLB



2-Level TLB Organization

Characteristic	ARM Cortex-A53	Intel Core i7
Virtual address	48 bits	48 bits
Physical address	40 bits	44 bits
Page size	Variable: 4, 16, 64 KiB, 1, 2 MiB, 1 GiB	Variable: 4 KiB, 2/4 MiB
TLB organization	<p>1 TLB for instructions and 1 TLB for data per core</p> <p>Both micro TLBs are fully associative with 10 entries, round robin replacement</p> <p>64-entry, four-way set associative TLBs</p> <p>TLB misses handled in hardware</p>	<p>1 TLB for instructions and 1 TLB for data per core</p> <p>Both L1 TLBs are four-way set associative, LRU replacement</p> <p>L1 I-TLB has 128 entries for small pages, seven per thread for large pages</p> <p>L1 D-TLB has 64 entries for small pages, 32 for large pages</p> <p>The L2 TLB is four-way set associative, LRU replacement</p> <p>The L2 TLB has 512 entries</p> <p>TLB misses handled in hardware</p>

TLB Misses

- If page is in memory
 - Load the PTE from memory and retry
 - Could be handled in hardware
 - Can get complex for more complicated page table structures
 - Note Intel and ARM (previous slide) do this
 - Or in software
 - Raise a special exception, with optimized handler
- If page is not in memory (page fault)
 - OS handles fetching the page and updating the page table
 - Then restart the faulting instruction

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

TLB Miss Handler

- **TLB miss indicates**
 - **Page present, but PTE not in TLB**
 - **Page not present**
- **Pipeline must recognize TLB miss before destination register overwritten**
 - **Raise exception**
- **Handler copies PTE from memory to TLB**
 - **Then restarts instruction**
 - **If page not present, page fault will occur**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Page Fault Handler

- Use faulting virtual address to find PTE
- Locate page on disk
- Choose page to replace
 - If dirty, write to disk first
- Read page into memory and update page table
- Make process runnable again
 - Restart from faulting instruction

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

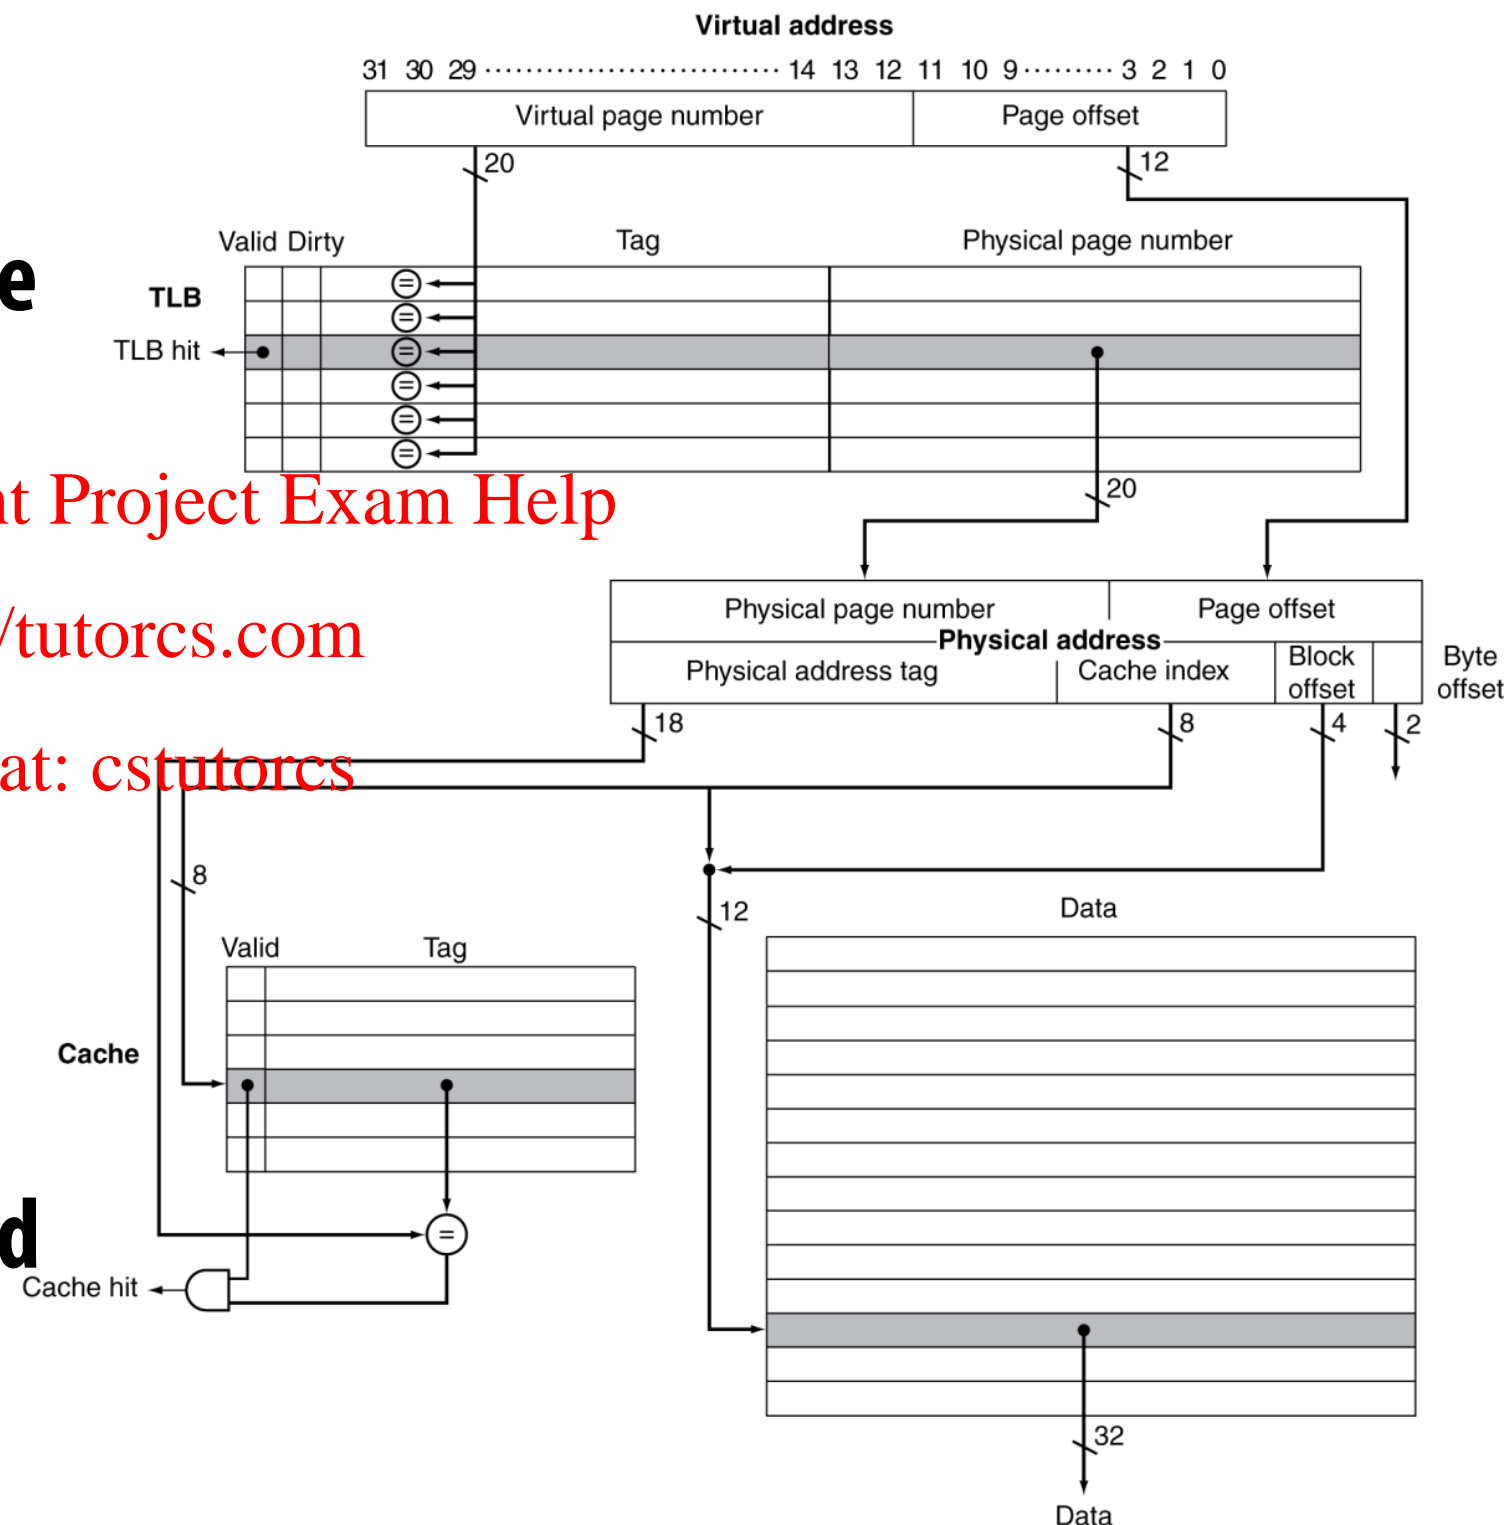
TLB and Cache Interaction

■ If cache tag uses physical address

- Need to translate before cache lookup

■ Alternative: use virtual address tag

- Complications due to aliasing
 - Different virtual addresses for shared physical address



Memory Protection

- **Different tasks can share parts of their virtual address spaces**
 - **But need to protect against errant access**
 - **Requires OS assistance**
- **Hardware support for OS protection**
 - **Privileged supervisor mode (aka kernel mode)**
 - **Privileged instructions**
 - **Page tables and other state information only accessible in supervisor mode**
 - **System call exception (e.g., ecall in RISC-V)**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Memory Hierarchy Summary

- **The following slides are for your review only. We will not cover them in class. I hope that you find all this material familiar at this point.**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

The Memory Hierarchy

The BIG Picture

- Common principles apply at all levels of the memory hierarchy
 - Based on notions of caching
- At each level in the hierarchy
 - Block placement
 - Finding a block
 - Replacement on a miss
 - Write policy

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Block Placement

- **Determined by associativity**
 - **Direct mapped (1-way associative)**
 - **One choice for placement**
 - **n-way set associative**
 - **n choices within a set**
 - **Fully associative**
 - **Any location**
- **Higher associativity reduces miss rate**
 - **Increases complexity, cost, and access time**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Finding a Block

- **Hardware caches**
 - **Reduce comparisons to reduce cost**
- **Virtual memory**
 - **Full table lookup makes full associativity feasible**
 - **Benefit in reduced miss rate**

Associativity	Location method	Tag comparisons
Direct mapped	Index	1
n-way set associative	Set index, then search entries within the set	n
Fully associative	Search all entries	#entries
	Full lookup table	0

Replacement

- **Choice of entry to replace on a miss**
 - **Least recently used (LRU)**
 - **Complex and costly hardware for high associativity**
 - **Random**
 - **Close to LRU, easier to implement**
- **Virtual memory**
 - **LRU approximation with hardware support**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Write Policy

■ Write-through

- Update both upper and lower levels
- Simplifies replacement, but may require write buffer

■ Write-back

- Update upper level only
- Update lower level when block is replaced
- Need to keep more state

■ Virtual memory

- Only write-back is feasible, given disk write latency

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Sources of Misses

- **Compulsory misses (aka cold start misses)**
 - **First access to a block**
- **Capacity misses**
 - **Due to finite cache size**
 - **A replaced block is later accessed again**
- **Conflict misses (aka collision misses)**
 - **In a non-fully associative cache**
 - **Due to competition for entries in a set**
 - **Would not occur in a fully associative cache of the same total size**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Cache Design Trade-offs

Design change	Effect on miss rate	Negative performance effect
Increase cache size	Decrease capacity misses	May increase access time
Increase associativity	Decrease conflict misses	May increase access time
Increase block size	Decrease compulsory misses	Increases miss penalty. For very large block size, may increase miss rate due to pollution.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Pitfalls

■ Byte vs. word addressing

- **Example: 32-byte direct-mapped cache, 4-byte blocks**

- **Byte 36 maps to block 1**

Assignment Project Exam Help

- **Word 36 maps to block 4**

<https://tutorcs.com>

■ Ignoring memory system effects when writing or generating code

WeChat: cstutorcs

- **Example: iterating over rows vs. columns of arrays**
- **Large strides result in poor locality**

Pitfalls

- In multiprocessor with shared L2 or L3 cache
 - Less associativity than cores results in conflict misses
 - More cores \Rightarrow need to increase associativity
- Using AMAT to evaluate performance of out-of-order processors
 - Ignores effect of non-blocked accesses
 - Instead, evaluate performance by simulation

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Pitfalls

- **Extending address range using segments**
 - **E.g., Intel 80286**
 - **But a segment is not always big enough**
 - **Makes address arithmetic complicated**
- **Implementing a VMM on an ISA not designed for virtualization**
 - **E.g., non-privileged instructions accessing hardware resources**
 - **Either extend ISA, or require guest OS not to use problematic instructions**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Concluding Remarks

- **Fast memories are small, large memories are slow**
 - We really want fast, large memories ☹️
 - Caching gives this illusion 😊
- **Principle of locality**
 - Programs use a small part of their memory space frequently
- **Memory hierarchy**
 - L1 cache ↔ L2 cache ↔ ... ↔ DRAM memory
↔ disk
- **Memory system design is critical for multiprocessors**

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs